



KaribuTech AI

INSURANCE · LEADERSHIP SERIES · SEPTEMBER 2026

The AI Blueprint Series for Insurance Leaders

Twelve blueprints, one plan: how every leader in insurance, from the CEO to the regulator, can put AI to work safely and profitably

Insurance CEO

Chief Information Officer

Head of HR

Head of New Business

Head of Claims

Head of Medical Claims

Head of Premium Collections

Head of Actuarial

Head of Risk

Insurance Broker

Reinsurance Executive

Insurance Regulator

By Kennedy Chengeta, PhD, and Justice Muguda
with the KaribuTech AI Engineering Team

kaributech.co.za



Contents

This volume brings together twelve AI blueprints for insurance leaders. Each can be read on its own, and each follows the same structure, so leaders across the

insurer share one language and one plan. Page numbers below refer to this volume.

	01	The AI Blueprint for the Insurance CEO Growth, profit and trust: the seven moves, maturity model and 36-month roadmap	4
	02	The AI Blueprint for the Chief Information Officer Core systems, data, connections and security so AI can scale safely	24
	03	The AI Blueprint for the Head of HR Reskilling, job redesign and fair use of AI inside HR	44
	04	The AI Blueprint for the Head of New Business Faster quotes, paperless onboarding and policies that stay on the books	64
	05	The AI Blueprint for the Head of Claims Faster, fairer claims with less leakage and fraud	84
	06	The AI Blueprint for the Head of Medical Claims Real-time pre-authorisation, adjudication and fraud, waste and abuse	104
	07	The AI Blueprint for the Head of Premium Collections Collect more, lose fewer policies and treat customers fairly	124
	08	The AI Blueprint for the Head of Actuarial Sharper prices, faster reserves and models you can explain	144
	09	The AI Blueprint for the Head of Risk Seeing risk earlier and governing AI as a risk in its own right	164
	10	The AI Blueprint for the Insurance Broker Better advice, faster placement and stronger client relationships	184
	11	The AI Blueprint for the Reinsurance Executive Cleaner cedant data, sharper catastrophe models, faster recoveries	204
	12	The AI Blueprint for the Insurance Regulator Supervising insurers' AI, and using AI to supervise	224


How to use this series

AI succeeds when every leader moves together. The CEO sets the direction, the CIO and HR build the foundations, business leaders

put AI to work, and risk, actuarial and regulators keep it safe.


EXHIBIT 1

Twelve blueprints, four roles in one plan



Source: KaribuTech AI analysis.

Every blueprint follows the same structure



ONE PRINCIPLE ACROSS ALL TWELVE

AI is a tool for professionals. Technology teams build and run it; every leader and professional must learn to use it, judge its output and stay accountable for the result.

INSURANCE · CEO SERIES · SEPTEMBER 2026



The AI Blueprint for the Insurance CEO

From pilots to an AI-first insurer in 36 months: seven moves, a 100-day quick-wins approach and a roadmap for scale

By Kennedy Chengeta, PhD, and Justice Muguda
with the KaribuTech AI Engineering Team

kaributech.co.za





About this report

This report is written for chief executives and boards of insurers, reinsurers and pension funds in developing markets.

It sets out how to move from scattered AI pilots to an AI-first insurer: why the economics have changed, how to assess maturity, seven moves that separate scaled AI insurers from the rest, a 100-day quick-wins approach, a reference architecture, an operating model, a governance agenda and a 36-month roadmap with a CEO scorecard.

External data are sourced beneath each exhibit, and a full list of references appears at the end. Frameworks without an external source are the authors' own.

WRITTEN FOR

CEOs

Boards

Executive teams

Regulators

Contents

01	Introduction	6
02	Why AI, why now	7
03	Lessons from history	8
04	The five-level maturity model	9
05	Seven moves for the insurance CEO	10
06	Why AI programmes stall	12
07	Where AI pays off across the value chain	13
08	Generative AI and AI agents: what is real now	14
09	AI for inclusion	15
10	The 100-day quick-wins approach	16
11	Building the business case	17
12	A reference architecture for the AI insurer	18
13	Organising for AI	19
14	Governing AI: the board and the regulator	20
15	New risks: cyber, deepfakes, privacy, models	21
16	The CEO scorecard and the 36-month roadmap	22
17	About the authors and KaribuTech AI	23

Introduction

Artificial intelligence has moved from the innovation lab to the insurer's income statement. The question for chief executives is no longer whether to use AI, but how to scale it, safely and profitably, faster than competitors.

Almost every insurer now uses AI somewhere. Very few have scaled it. Most portfolios are a collection of pilots: a chatbot here, a fraud model there, and a generative AI assistant that never reached the claims floor. The value sits in the gap between using AI and running the business on it.

In our markets the stakes are higher. Protection gaps are wide, incomes are informal and data is thin, yet customers already live on mobile money and messaging apps. AI is the only way to serve millions of small policies at a cost that makes them viable.

This blueprint distils what KaribuTech AI has learned building automation, integration and AI for insurers and banks in Southern Africa, and combines it with the latest public evidence. It sets out a maturity model to locate where you are today, **seven moves** that separate scaled AI insurers from the rest, a **100-day quick-wins approach** that proves value before large investments, a reference architecture, an operating model, a governance agenda and a **36-month roadmap**.

Our guiding principle is simple: start where your data is best and your losses are largest, prove value in 100 days, then scale what works with the discipline of a regulated financial institution.

AT A GLANCE: THE SEVEN MOVES

- 1 Anchor AI in business value, not technology
- 2 Start with claims and onboarding: win in 100 days
- 3 Build the data foundation once, for AI and IFRS 17
- 4 Automate the whole process, not just the decision
- 5 Organise for AI: hub, spokes and new skills
- 6 Govern AI from the board, not the IT department
- 7 Compete as an ecosystem: partner, share and embed

Why AI, why now

Three forces have changed the economics of AI for insurers at the same time: the cost of capable AI has collapsed, customers now live

on digital rails, and regulators have started to set clear expectations.

The cost of AI has collapsed

Running a GPT-3.5-level model became about **280 times cheaper** between November 2022 and October 2024. Capable AI that once needed a global balance sheet is now within reach of a mid-sized insurer in Harare, Lagos or Kathmandu.

Customers live on digital rails


More than **\$2 trillion** moved through mobile money in 2025, across 2.3 billion accounts. Premiums can be collected and claims paid in minutes, if the processes behind them are automated.

Regulators are setting the rules

The IAIS published its Application Paper on the supervision of AI in July 2025, the EU AI Act classes life and health pricing as high-risk, and African countries are adopting national AI strategies. Governance is now a licence to operate.


EXHIBIT 1

Almost every organisation uses AI; very few have scaled it



Source: McKinsey, The State of AI (2025); Gartner (2024); Stanford HAI, AI Index 2025; GSMA, State of the Industry Report on Mobile Money (2026); IAIS (2025); PwC, Stopping the leaks; Accenture underwriting survey.

Where the money is



WHAT THIS MEANS FOR THE CEO

The scarce resource is no longer the model. It is clean data, redesigned processes, people who can run AI safely, and a board that owns the risk. Insurers that build those four foundations now will compound their advantage; those that keep buying pilots will pay more each year to stand still.

Lessons from history: when technology rewrites an industry

Every generation of insurance leaders believes its disruption is unprecedented. It rarely is. History shows the same pattern again and again: incumbents rarely lose because they lack the new technology; they lose

because they define their business by how they deliver value today rather than by the value customers actually buy. Insurance itself was born from exactly this kind of shift, twice, in the 17th and 18th centuries.

19TH CENTURY · THE ICE KINGS

The ice harvesters who never became refrigeration companies

In 1806 Boston merchant Frederic Tudor shipped a cargo of pond ice to Martinique. Within decades the “Ice King” and his rivals were cutting ice from New England lakes and shipping it as far as Calcutta. Then mechanical refrigeration arrived: first ice factories, then the home refrigerator. By the 1930s the natural ice trade had all but disappeared.

Lesson: customers bought cold, not ice.

Harvested ice

Cut from frozen lakes, shipped in sawdust



Ice factories

Ice made by machine, anywhere, all year



Home refrigerator

Cold delivered inside every kitchen

Three more stories every insurance CEO should know

c. 1688

Lloyd's coffee house

Edward Lloyd's London coffee house became the place to hear shipping news. Underwriters gathered where the information was best; from 1734 the house published Lloyd's List. Marine insurance grew around a data advantage.

Insurance markets form around whoever has the best information.

1762

The Equitable's mortality tables

The Equitable Society in London was among the first to price life cover by age using mortality data, and in 1775 appointed William Morgan as its actuary. Pricing moved from guesswork to evidence, and the actuarial profession was born.

Better data and models create whole new ways to price risk.

1975–2012

Kodak's digital camera

A Kodak engineer, Steven Sasson, built one of the first digital cameras in 1975. Kodak kept protecting its film business; digital photography grew elsewhere, and the company filed for bankruptcy protection in 2012.

Inventing the future is not enough if you protect the past.

THE LESSON FOR INSURERS

Customers do not buy policies, paperwork or branches; they buy protection and peace of mind. Insurers that define themselves by the value they deliver, not the process they run today, will be the ones that build the AI-era equivalent of the refrigerator.

Source: G. Weightman, *The Frozen-Water Trade* (2003); M. E. Ogborn, *Equitable Assurances* (1962); Lloyd's of London; Eastman Kodak Company (2012); KaribuTech AI analysis.


Where are you today? The five-level maturity model

Before choosing where to invest, a CEO needs an honest view of where the organisation stands. The BluFountain AI Insurance Maturity Model assesses six dimensions across five levels. In our experience, most

insurers in developing markets sit at level 2 or 3; the jump to level 4 is about data and governance, not about buying more AI.

EXHIBIT 2

From paper to AI-first: five levels, six dimensions



Source: KaribuTech AI, BluFountain AI Insurance Maturity Model (proprietary).

How to use the model

01

Score each dimension separately: most insurers are uneven, with technology ahead of data and governance behind both.

02

Set a target level per dimension for 36 months. Level 4 across all six is a realistic ambition for most mid-sized insurers.

03

Fix the weakest dimension first. A level-4 model on level-2 data will not survive contact with production.

Seven moves for the insurance CEO

Insurers that scale AI do seven things differently. None of them is about choosing a

model; all of them are about how the business is run.

MOVE 1

01

Anchor AI in business value, not technology

AI programmes justified on technical grounds rarely survive their first budget review. Value comes from redesigning how a claim is settled or a policy is sold, not from bolting a model onto an old process.

What leaders do

- Pick three domains where data is best and losses are largest, typically claims, onboarding and fraud.
- Give each a business owner, a baseline and one number to move, such as days to pay or cost per policy.
- Stop pilots that cannot show a path to scale within 12 months.

ASK YOUR TEAM

Which three AI use cases will move our combined ratio this year, and who owns each number?

MOVE 2

02

Start with claims and onboarding: win in 100 days

Claims is where trust is won and leakage hides; onboarding is where cost and drop-off are highest. Both are process-heavy, document-heavy and measurable, which makes them ideal first wins.

What leaders do

- Automate first notice of loss, document reading and triage before fully automated decisions.
- Digitise KYC, signatures and lead-to-offer tracking on one product, then template it.
- Keep a person on every decline, and publish results within 100 days.

ASK YOUR TEAM

What would it take to pay a simple claim in hours, not weeks, by the end of this quarter?

MOVE 3

03

Build the data foundation once, for AI and IFRS 17

The IFRS 17 accounting standard forced insurers to keep detailed data for every policy and year of sale. The same foundation feeds AI. Insurers that treat reporting data and AI data as two projects pay twice and still lack a single customer view.

What leaders do

- Create one trusted source of customer, policy and claims data.
- Keep sensitive data in-country; use the cloud where it is cheaper and faster.
- Measure data quality as a board KPI, not an IT metric.

ASK YOUR TEAM

Do we have one version of the customer, the policy and the claim, and who signs it off?

MOVE 4

04

Automate the whole process, not just the decision

AI creates value only when its decisions flow into work that gets done. Digital workflows, business rules, software robots and AI together turn each prediction into an action with a deadline and an audit trail.

What leaders do

- Run core processes on digital workflows, with rules the business can change.
- Connect core systems securely, so AI tools and partners can use them without copying data.
- Measure where work gets stuck; let AI assistants help inside the process, not run it.

ASK YOUR TEAM

Can we see, in real time, where every claim and application is stuck, and why?

MOVE 5

05

Organise for AI: hub, spokes and new skills

Separate AI divisions drift away from the risk; AI buried inside IT moves too slowly. The model that works is a central hub for platforms and standards, with AI squads embedded in claims, underwriting and actuarial teams.

What leaders do

- Appoint a Chief Data and AI Officer with a mandate across the business.
- Pair actuaries with AI engineers: engineers build models, actuaries own the risk and sign-off.
- Reskill before you automate: claims clerks become claims analysts reviewing AI decisions.

ASK YOUR TEAM

Who in our organisation is accountable for every model in production?

MOVE 6

06

Govern AI from the board, not the IT department

Regulators are converging on the same expectations: accountability, robustness and security, transparency, fairness and redress. Boards that own AI risk move faster, because they can approve use cases with confidence.

What leaders do

- Keep a register of every AI model, with an owner, monitoring and a retirement date.
- Require explainable decisions, human review of declines and bias testing before use.
- Report AI value and AI incidents to the board every quarter.

ASK YOUR TEAM

If the regulator asked tomorrow, could we explain every automated decline we made last month?

MOVE 7

07

Compete as an ecosystem: partner, share and embed

Banks, telcos, ride-hailing apps and retailers already own the customer relationship and the payment rails. Small insurers cannot build every capability alone, but

together they can share platforms, fraud data and models, as carmakers share components.

What leaders do

- Embed cover in partners' journeys: mobile money, bank accounts, trips and checkouts.
- Agree data rights, customer ownership and AI governance with every partner before launch.

FROM OUR WORK

Underwriting from weeks to days

Using a 16-module insurance suite built on workflow, rules and document AI, the KaribuTech AI engineering team, working with client teams, cut underwriting turnaround from weeks to days. The same flow runs on USSD, WhatsApp and mobile money, so customers without smartphones are served by the same automated process.

Weeks → days

underwriting turnaround

One flow

across USSD, WhatsApp, app and branch

Human in the loop

on every referral and decline

Why AI programmes stall, and how to avoid it

Most failed AI programmes in insurance do not fail on technology. They stall on ownership, data, process and

people. The exhibit below lists the seven pitfalls we see most often and the fix for each.

EXHIBIT 3

Seven pitfalls that stall AI in insurance, and the fix

PITFALL	WHAT IT LOOKS LIKE	THE FIX
1 Pilot purgatory	Dozens of pilots, none in production, no one accountable for scale	Fund only use cases with a path to scale; stop the rest at 12 months
2 Technology-led AI	AI owned by IT, measured in models built, not money made	A business owner, a baseline and one KPI for every use case
3 Data last	Models built on siloed, inconsistent data that no one trusts	One trusted data source and data-quality targets before scaling
4 One model for everything	Expecting one LLM or agent to run the whole claims journey	Use several specialist AI tools inside one controlled process
5 Automating broken processes	AI bolted onto a paper-era process with the same hand-offs	Redesign the process and measure where it gets stuck first
6 Governance as an afterthought	No model inventory, no explanations, no bias testing	Inventory, human review of declines and bias tests from day one
7 Forgetting people	Staff resist, skills are missing, adoption stalls	Reskill before automating; involve front-line staff in design

Source: KaribuTech AI analysis.

Build, buy or partner: source AI by what differentiates you

Buy

COMMODITY CAPABILITIES

Document reading, chatbots, speech tools, cloud AI services, software robots

Speed and cost; avoid rebuilding what the market sells

Build

WHAT DIFFERENTIATES YOU

Pricing, underwriting and fraud models trained on your own data; your workflows and rules

Advantage competitors cannot copy

Partner

WHAT NO ONE CAN DO ALONE

Telcos, banks and platforms for reach; shared industry data and platforms; insurtechs

Scale and reach at a fraction of the cost

BEFORE FUNDING ANY AI USE CASE, ASK FOUR QUESTIONS

Value

Which number moves, and by how much?

Data

Is the data good enough today?

Process

Which workflow will act on the output?

Owner

Who is accountable for the result?

Where AI pays off across the value chain

AI value is not spread evenly. It concentrates in the steps that are document-heavy, decision-heavy and repeated at volume. The exhibit below maps the five

places where insurers in our markets see the fastest payback, what AI changes and the metric a CEO should track.

EXHIBIT 4

Five places AI pays off, and the metric that proves it

STAGE	TODAY'S PROBLEM	WHAT AI DOES	KPI TO TRACK
Product design	Products that don't fit informal incomes	Designs micro, usage-based and parametric covers from customer and alternative data	Take-up and first-year persistency
Distribution	High cost to reach rural and young customers	Next-best offers and local-language chat inside mobile-money, bank and telco apps	Cost per policy; digital share of sales
Underwriting	Thin data, broad risk pools, manual referrals	Scores risk with mobile, satellite and vehicle-tracking data; reads submissions automatically	Time to issue; loss ratio by segment
Claims	Weeks of paperwork and phone calls	Photo-based assessment, document AI, automated triage and payment by mobile money	Days to pay; claims leakage
Fraud	Leakage hidden in manual files	Spots suspicious patterns on every claim, with a person on every decline	Fraud detected; false-positive rate

Source: KaribuTech AI analysis.

The evidence from leading insurers

23 days

faster liability assessment on complex claims at Aviva after rewiring claims with AI

65%

fewer customer complaints at Aviva following the same programme

80%

faster settlement of simple claims with Allianz's AI agents

5–10%

claims cost savings typical from reducing leakage; up to 25% leakage seen in life

Source: McKinsey, Rewired in action: Aviva; Allianz (2025); PwC, Stopping the leaks.

In insurance, the claim is the product. Start where your data is best and your losses are largest, and let the first win pay for the next.


Generative AI and AI agents: what is real now

Generative AI reads and writes; AI agents plan and act. Together they can take on the document-heavy, conversation-heavy work that still consumes most insurance operations. Insurers are moving quickly:

publicised AI deployments by insurers grew 87% in the fourth quarter of 2025, and 21% of them used AI agents, more than half of those in claims. The question for the CEO is how much autonomy to give, and where.

EXHIBIT 5

Four levels of autonomy: raise the level only as trust is earned



MORE AUTONOMY →

Source: Evident (2026); KaribuTech AI analysis.

Where it pays off today

- **Claims:** first notice of loss by WhatsApp or voice, claim summaries, document checks and customer updates.
- **Underwriting:** reading broker submissions, schedules and medical reports; drafting quotes for review.
- **Service:** answering policy questions in local languages, around the clock.
- **Operations:** comparing policy wordings, drafting letters, matching reinsurance statements.
- **IT:** documenting and updating old systems before they are replaced.

Guardrails that make it safe

- **Grounding:** answer only from the insurer's approved documents and data.
- **Human review:** a person on every decline, payment above a limit and any complaint.
- **Security:** protection against manipulation and data leaks; no customer data in public AI tools.
- **Logging:** every prompt, answer and action recorded for audit.
- **Testing:** accuracy and bias tested before launch and monitored after.

AI AGENTS HELP THE PROCESS; THEY DO NOT RUN IT

The safest pattern puts AI agents inside a governed workflow: the workflow decides what happens next, rules set the limits, and agents do the reading, writing and checking in between.

AI for inclusion: reaching the uninsured profitably

The largest growth opportunity in our markets is not the insured customer switching providers; it is the customer who has never been insured. Traditional distribution cannot reach them at a price they can afford. AI, mobile

money and parametric design change that arithmetic, because they cut the cost of selling, servicing and paying a small policy close to zero.

EXHIBIT 6

Three levers that make small policies viable

01

Embedded distribution

Cover sold inside the apps and channels customers already use: mobile money, airtime, bank accounts, ride-hailing and retail checkouts.

- Next-best offer at the moment of need
- Premiums from airtime or wallets
- No agent commission on every sale

02

Parametric design

Pay automatically when an index is crossed, such as rainfall, temperature or a flood level, instead of assessing each loss.

- Weather and satellite data as triggers
- Payouts in days by mobile money
- No loss adjuster for small claims

03

Alternative data

Price and underwrite customers with no credit or claims history, using signals they already generate.

- Mobile usage and payment history
- Satellite and mapping data
- Consent-based, explainable models

Source: KaribuTech AI analysis.

The shift is already happening

2.3bn

registered mobile money accounts worldwide, moving more than \$2 trillion in 2025

15%

of African microinsurance products used mobile money in 2023, up from 12% in 2022

17%

relied on cash in 2023, down from 22% a year earlier

Source: GSMA, State of the Industry Report on Mobile Money (2026); Microinsurance Network, Landscape of Microinsurance (2024).

EXAMPLES FROM OUR MARKETS

Zimbabwe

Funeral cover sold and paid through airtime reached the mass market at low cost: the phone became the channel and the payment rail.

Kenya

Weather-index cover for smallholders pays out on rainfall data, with no field loss assessment, through mobile money.

Examples: Econet EcoSure (Zimbabwe); ACRE Africa (Kenya), published programme information.

The 100-day quick-wins approach

Large AI programmes fail when they try to do everything at once. Our approach delivers value in 100-day cycles, run by small cross-functional teams focused on one customer journey. Each cycle builds the

base architecture, delivers one production use case end to end, and leaves templates that make the next use case cheaper and faster.

EXHIBIT 7

Every 100 days: discover, build, deploy, then repeat on the next journey

DAYS 1-30

Discover and design

- As-is assessment of the chosen journey
- Redesigned process and rules agreed
- Technology, security and data plan

DAYS 31-60

Build and set the scene

- Workflow, system links and AI tools built
- Document reading and ID checks connected
- Team delivers in two-week steps

DAYS 61-100

Test and deploy

- User testing with real cases
- Go live on one product, one channel
- Measure against the baseline

Run it as: one small cross-functional team · two-week delivery steps · a customer-journey lens · templates reused on every later journey

Source: KaribuTech AI quick-wins methodology.

Three quick wins we recommend for insurers

QUICK WIN 1

Automated onboarding and KYC

- Digital KYC and identity checks
- Documents submitted on any channel
- Digital signatures and contract templates
- Referrals to underwriters with deadlines

MEASURE

Time to issue, drop-off rate, cost per policy

QUICK WIN 2

Claims intake and triage

- First notice of loss on WhatsApp, USSD and app
- Document and photo AI for assessment
- Automatic routing to the right handler
- Fraud screening on every claim

MEASURE

Days to pay, leakage, straight-through rate

QUICK WIN 3

Queries, complaints and process analytics

- Log queries on any channel, track status
- Bots resolve simple requests instantly
- Rules route the rest with deadlines and escalation
- Analytics show where work is stuck

MEASURE

Resolution time, complaints, missed deadlines

Building the business case: what AI costs and what it returns

Boards approve what they can measure. A credible AI business case is built bottom-up from a few use cases, with a baseline, a cost line and a named owner for

every benefit. The example below is illustrative, for a mid-sized insurer, to show the shape of the numbers rather than a forecast.

EXHIBIT 8

Illustrative example: a mid-sized insurer with \$200m of gross written premium

Benefit lever	Assumption	Annual value
Claims leakage reduced	5% of \$110m claims	\$5.5m
Admin work automated	15% of \$50m expenses	\$7.5m
Fraud detected earlier	1% of claims	\$1.1m
Lapses prevented	2% of premium retained	\$4.0m
Total annual benefit at scale		\$18.1m


WHERE THE MONEY GOES

People and change	35–45%
Data platform	20–25%
Cloud and compute	15–20%
Software and licences	10–15%
Governance and risk	5–10%

Typical split of a scaled programme

Source: Illustrative KaribuTech AI example; benefit ranges informed by PwC, Stopping the leaks. Not a forecast.

The J-curve: why year one looks expensive



- **Year 1:** invest in data, platforms and the first 100-day wins; value starts but costs lead.
- **Year 2:** templates and reuse make each new use case cheaper; the programme breaks even.
- **Year 3:** value compounds across journeys; AI funds its own expansion.

THREE RULES FOR A CREDIBLE CASE

- 1 Baseline every KPI before you start
- 2 Count only benefits a business owner signs
- 3 Track value and cost every quarter

A reference architecture for the AI insurer

Scaling AI is an architecture problem as much as a data-science problem. The pattern below separates customer channels, workflows, AI tools and data, so

that each can change without breaking the others, and keeps security and governance running through every layer.

EXHIBIT 9

Six layers, with security and governance running through all of them



Source: KaribuTech AI reference architecture.

Five design principles

01

Automation first

Every AI decision lands in a process with an owner, a deadline and an audit trail.

02

Connect, don't copy

Core systems are reached through secure, controlled links, not copied spreadsheets or shared passwords.

03

No lock-in

Open standards let you switch cloud or vendor without rebuilding.

04

Secure by design

Strong identity checks for every user and system; data kept in-country where the law requires.

05

Business in control

Business teams change rules and forms without waiting for IT.

IN PRACTICE

BluFountain Insure, KaribuTech AI's insurance AI marketplace, follows this pattern: core platforms for life, short-term, medical and funeral business share customer-checking, anti-money-laundering, document, actuarial and valuation tools, with AI agents that detect fraud and help run processes.


Organising for AI: operating model and talent

AI changes who does what. The hub-and-spoke model keeps platforms and standards central while putting AI capability next to the people who own the risk.

Accountability for insurance risk stays with actuaries and underwriters; accountability for the platform sits with the AI lead.

EXHIBIT 10

Hub and spoke: one platform, AI squads beside the business



Source: KaribuTech AI analysis.

Who owns what

Actuaries

own pricing, reserving, IFRS 17 and model sign-off

AI practitioners

own the data, the models and keeping them running

Model risk (second line)

validates independently; builders never sign off alone

Business owners

own the process, the KPI and the customer outcome

Reskill before you automate

Claims clerk	→	Claims analyst reviewing AI decisions
Underwriting assistant	→	Underwriter using AI triage
Call-centre agent	→	Customer advocate for complex cases
Data entry clerk	→	Data quality steward
Actuarial analyst	→	Actuarial data scientist
Branch sales	→	Digital and embedded-channel adviser

The World Economic Forum expects 39% of core skills to change by 2030. HR, not IT, decides whether AI replaces people or upgrades them.

Five roles to hire or build

Chief Data and AI Officer

Owens the AI agenda, platform and value

Data and AI engineers

Build and run data flows and models

AI model validators

Test accuracy and fairness independently

Process designers

Redesign workflows and business rules

AI product owners

Turn use cases into adopted products

Source: World Economic Forum, Future of Jobs Report 2025; KaribuTech AI analysis.

Governing AI: the board and the regulator

The rules for AI in insurance are converging around the same principles: govern, explain, be fair and stay secure. Insurers that build to the International

Association of Insurance Supervisors (IAIS) principles now will be ready for whichever national rules follow.


EXHIBIT 11

The AI rulebook for insurers is taking shape




Source: NAIC (2023); African Union (2024); EU AI Act and Digital Omnibus (2024, 2026); Government of Kenya (2025); IAIS, Application Paper on the supervision of AI (2025).

Ten questions every board should ask about AI



Governing AI inside the insurer: five controls



Responsible AI is not a brake. It is what lets customers, regulators and reinsurers trust AI-driven decisions, and what lets the board approve the next use case with confidence.

New risks: cyber, deepfakes, privacy and model risk

AI creates new attack surfaces and new ways for decisions to go wrong. Insurers face these risks twice: as users of AI in their own operations, and as

underwriters of customers who use it. Governance protects the first; new products serve the second.

\$4.44m

global average cost of a data breach (2025)

13%

of organisations reported breaches of AI models or applications

97%

of those lacked proper AI access controls

1 in 5

organisations had a breach involving unapproved "shadow" AI

Source: IBM, Cost of a Data Breach Report 2025.

EXHIBIT 12

Five AI risks for insurers, and the control for each

RISK	WHAT CAN GO WRONG	THE CONTROL
Cyber and deepfakes	Synthetic voices and documents used to open policies or claim fraudulently	Liveness checks, document forensics, call-back verification
Data privacy and sovereignty	Customer data sent to external AI tools or stored outside the country	Data classification, in-country hosting, approved tools only
Model risk	Models that get worse over time, fail silently or are confidently wrong	A model register, independent testing, monitoring and an off switch
Conduct and fairness	Pricing or claims decisions that disadvantage groups of customers	Bias testing, explainable decisions, human review of declines
Third-party concentration	Dependence on one cloud, model provider or vendor	Exit plans, portable architecture, multi-vendor options

Source: KaribuTech AI analysis.

THE OPPORTUNITY ON THE OTHER SIDE

The same risks create demand. Cyber cover, AI liability, cover for technology errors and cover for model failure are among the fastest-growing needs of insurers' own business customers. Insurers that manage AI risk well internally will be best placed to price it for others.

PRODUCTS TO BUILD

Cyber cover

AI liability

Technology errors cover

Model failure cover

Deepfake fraud cover

The CEO scorecard and the 36-month roadmap

EXHIBIT 13

Three horizons from first win to AI-first insurer

	0-100 DAYS Foundation	100 DAYS-18 MONTHS Scale	18-36 MONTHS Transform
Business	Three domains chosen, baselines set	Five to eight use cases live	AI embedded in every core journey
Data	Data audit; one journey's data cleaned	One trusted view of every customer	Shared industry data and models
Technology	Digital workflow and system links built	Cloud in place; AI assistants piloted	AI assistants working across processes
People	AI literacy for leaders	Hub and squads staffed; reskilling at scale	AI-literate whole company
Governance	AI policy and model inventory	Board AI committee; bias testing	Audited, explainable AI across the firm

Source: KaribuTech AI analysis.

EXHIBIT 14

The CEO scorecard: six numbers, plus four for the AI era

Combined ratio $\text{Claims plus expenses} \div \text{earned premium}$	GWP growth Premium growth, year on year	Days to pay a claim First notice to settlement
Cost per policy Cost to acquire and serve a policy	Retention and NPS Renewals and customer advocacy	Solvency coverage $\text{Own funds} \div \text{capital required}$
AI ERA		
AI value delivered (\$) Processes AI-enabled (%) AI incidents and complaints Staff reskilled (%)		

Where CEOs should start

You do not need to commit to a multi-year programme on day one. Start with a focused, time-boxed

diagnostic, let the first win fund the second, and hold the board accountable for the scorecard every quarter.

Weeks 1-2

Score your maturity across the six dimensions

Weeks 3-4

Choose three domains; name owners and baselines

Month 2

Appoint the AI lead; set up the board AI committee

Month 3

Launch the first 100-day cycle on one journey

Day 100

Report results to the board; fund the next cycle

About the authors



Kennedy Chengeta, PhD

Co-founder, KaribuTech AI

Enterprise Architect and AI Leader · Camunda Chapter Lead, Africa

Leads AI, automation and integration programmes for insurers and banks across Southern Africa.



Justice Muguda

Chief Technology Officer, KaribuTech AI


Cloud, integration and cybersecurity

Leads KaribuTech AI's technology and engineering, from cloud platforms to secure connections between systems.

With contributions from the KaribuTech AI Engineering Team. To discuss this report, or to run a 100-day AI quick-wins diagnostic, contact KaribuTech AI at kaributech.co.za.

References

1. Accenture, underwriting survey, as reported by Insurance Business.
2. African Union (2024), Continental Artificial Intelligence Strategy, July 2024.
3. Allianz (2025), announcements on AI agents in claims handling.
4. European Union (2024), Regulation (EU) 2024/1689, the Artificial Intelligence Act.
5. European Union (2026), Digital Omnibus on AI, amending the AI Act timeline.
6. Gartner (2024), research on AI adoption in organisations.
7. Government of Kenya (2025), Kenya Artificial Intelligence Strategy 2025–2030.
8. GSMA (2026), State of the Industry Report on Mobile Money.
9. IAIS (2025), Application Paper on the supervision of artificial intelligence, July 2025.
10. IASB (2017, amended 2020), IFRS 17 Insurance Contracts.
11. McKinsey & Company (2025), The State of AI.
12. McKinsey & Company, Rewired in action: Aviva.
13. McKinsey & Company, research on the value of AI in insurance.
14. NAIC (2023), Model Bulletin: Use of Artificial Intelligence Systems by Insurers.
15. PwC, Stopping the leaks: claims leakage in insurance.
16. Stanford HAI (2025), Artificial Intelligence Index Report 2025.
17. World Economic Forum (2025), Future of Jobs Report 2025.
18. Weightman, G. (2003), The Frozen-Water Trade.
19. Ogborn, M. E. (1962), Equitable Assurances.
20. Lloyd's of London, history of Lloyd's and Lloyd's List.
21. Eastman Kodak Company (2012), Chapter 11 filing.
22. Evident (2026), insurance AI deployment tracking.
23. IBM (2025), Cost of a Data Breach Report 2025.
24. Microinsurance Network (2024), Landscape of Microinsurance.
25. KaribuTech AI, BluFountain AI Insurance Maturity Model (proprietary).



About KaribuTech AI

KaribuTech AI is a South African technology company specialising in artificial intelligence, machine learning, intelligent automation and integration for insurers and banks across Africa. Our teams design and build intelligent workflows, decision engines, document and computer-vision AI, and cloud-native core platforms, including the **BluFountain Insure** platform: an insurance AI

marketplace with core systems for life, short-term, medical, funeral and investment business. We combine deep process expertise (Camunda, IBM BPM, MuleSoft, webMethods and document automation platforms) with practical AI engineering, so that AI moves out of the lab and into the operations that customers feel.

Disclaimer: this report provides general information and does not constitute legal, actuarial or investment advice. Third-party figures are cited beneath each exhibit and have not been independently verified by KaribuTech AI.



KaribuTech AI

INSURANCE · LEADERSHIP SERIES · SEPTEMBER 2026

The AI Blueprint for the Chief Information Officer

Modernising core systems, data and security so AI can scale safely: the technology agenda for insurers


By Kennedy Chengeta, PhD, and Justice Muguda
with the KaribuTech AI Engineering Team

kaributech.co.za



Contents

01	Introduction	26
02	Why the technology agenda must change	27
03	Lessons from history	28
04	Where IT delivery breaks	29
05	Six moves for the CIO	30
06	Modernising the core: three paths	32
07	The architecture for AI	33
08	Generative AI and AI agents in IT	34
09	Cloud, cost and data sovereignty	35
10	Security for the AI era	36
11	The 100-day quick-wins plan	37
12	The business case	38
13	What you need underneath	39
14	The technology team	40
15	Risks and guardrails	41
16	The scorecard and 12-month roadmap	42
17	About the authors and KaribuTech AI	43



About this report

This report is written for chief information officers, chief technology officers, chief data officers and technology leaders at insurers in developing markets.

AI only scales on solid foundations. This report shows the CIO how to build them without a risky big-bang programme: the pressures, six moves, three paths to modernise core systems, the architecture for AI, cloud and sovereignty, security, a 100-day plan, the business case, the team, the risks and a scorecard.

External data are sourced beneath each exhibit, and a full list of references appears at the end. Frameworks and illustrative figures without an external source are the authors' own.

WRITTEN FOR

CIOs and CTOs

Chief data officers

Architects

Security leaders

Introduction

Every AI ambition in insurance lands on the CIO's desk. AI can only scale as fast as the core systems, data and security underneath it allow.

Many insurers run on core systems that are decades old, with data spread across many databases and spreadsheets. New AI tools get bolted on at the edges, staff use public AI tools without permission, and security teams struggle to keep up.

The CIO's job is shifting from running systems to enabling the business: connecting systems so AI can reach them, building one trusted data platform, securing AI, and modernising the core step by step while the business keeps running.

This report sets out **six moves** for the CIO, **three paths to modernise** the core, the **architecture for AI, cloud and sovereignty, security**, a **100-day plan** and a **scorecard**, in plain language.

THREE QUESTIONS THIS REPORT ANSWERS

01

Which core systems should we keep, move or replace?

02

What foundations must be in place before AI can scale?

03

How do we secure AI and keep its costs under control?

AT A GLANCE: SIX MOVES FOR THE CIO

- 1 Modernise the core in steps, not a big bang
- 2 Build one trusted data platform
- 3 Connect everything securely
- 4 Secure AI by design
- 5 Run AI like a product
- 6 Govern vendors, cloud and cost

Why the technology agenda must change

Five pressures are converging on insurers' technology teams. Each is also a chance for

the CIO to move from cost centre to growth partner.

Legacy cores hold AI back

Old policy and claims systems are hard to connect, slow to change and expensive to run.

Data is scattered

Customer, policy and claims data sits in many systems and spreadsheets, with no single trusted view.

Cyber threats are growing

AI makes attacks cheaper, and AI systems themselves are new targets.

Staff are using AI anyway

Employees use public AI tools with company data, creating unseen risk.

Costs and choices are multiplying

Cloud, AI services and vendors bring flexibility, and new costs that are easy to lose control of.

EXHIBIT 1

The numbers behind the pressure

280x

cheaper to run capable AI in 2024 than in 2022

\$4.44m

global average cost of a data breach (2025)

13%

of organisations reported breaches of AI models or applications

97%

of those lacked proper AI access controls

Source: Stanford HAI, AI Index 2025; IBM, Cost of a Data Breach Report 2025.

AI models are becoming a commodity. The foundations underneath them, data, connections and security, are what set one insurer apart from another.

Lessons from history: insurers have modernised before

Insurers were among the earliest business users of computers, and many still run systems built decades

ago. Each wave of technology rewarded insurers that modernised steadily rather than all at once.

1959

COBOL

COBOL, a programming language designed for business, appeared in 1959. Many insurers' policy and claims systems are still written in it, more than 60 years later.

Today: AI explains old code, so it can be connected and replaced safely.

THE LESSON

Systems outlive the people who built them. Document and connect before you replace.

1964

IBM System/360

IBM announced the System/360 in 1964, a family of compatible computers that let businesses upgrade hardware without rewriting all their software.

Today: open standards and APIs keep change cheap.

THE LESSON

Standards and compatibility make change cheaper.

2006

The public cloud

Amazon Web Services launched its first cloud storage and computing services in 2006, letting companies rent computing power instead of buying it.

Today: rent AI capacity; keep sensitive data in-country.

THE LESSON

Pay for what you use; scale when you need to.

Today

AI on every system

AI now needs access to data and systems across the insurer, securely and in real time.

Today: one platform so every AI use case builds on the last.

THE LESSON

Connect, secure and modernise step by step.

Source: Histories of COBOL (1959), the IBM System/360 (1964) and Amazon Web Services (2006).

The insurers that did best in each wave did not replace everything at once. They connected the old to the new, then modernised piece by piece.

Where IT delivery breaks, and how AI helps

Technology delivery follows a cycle, from strategy to retiring old systems. At each step, projects slow down,

costs rise or risks are missed. The exhibit shows where it typically breaks and what AI does.

EXHIBIT 2

The technology delivery cycle: eight steps from strategy to retirement

	WHERE IT BREAKS TODAY	WHAT AI DOES
Strategy Deciding what to build	IT plans separate from business goals	Portfolio linked to business value
Architecture How systems fit together	Point-to-point links everywhere	Maps systems and finds dependencies
Data Collecting and cleaning data	Many copies, no single truth	Automated data checks and cataloguing
Build Writing and changing software	Slow delivery; old code nobody understands	Explains old code; helps write and test new
Test and release Getting changes live safely	Manual testing; risky releases	Automated tests; safer releases
Security Protecting systems and data	Alerts overwhelm the team	Prioritises alerts; spots unusual activity
Operations Keeping systems running	Problems found by customers first	Predicts failures; faster fixes
Retire and modernise Replacing old systems	Old systems never switched off	Maps what can move and what can go

Source: KaribuTech AI analysis.

The fastest way to deliver AI is to remove the friction in how technology is built and run, not to buy more AI tools.

Six moves for the Chief Information Officer

Technology leaders who scale AI well do six things differently. Each is about foundations,

connections and discipline, not about choosing the latest model.

MOVE 1

01

Modernise the core in steps, not a big bang

Replacing a core insurance system in one go is one of the riskiest projects an insurer can take on. A safer path connects the old core, moves products or processes across in stages, and switches off the old system piece by piece.

What leaders do

- Put a secure connection layer in front of every core system.
- Move one product or process at a time to the new platform.
- Set a date to retire each old component, and keep to it.

ASK YOUR TEAM

Which core system would hurt us most if it failed, and what is our plan to reduce that risk?

MOVE 2

02

Build one trusted data platform

AI, IFRS 17 reporting and customer service all need the same thing: one trusted view of customers, policies and claims. Building it once saves paying for it several times over.

What leaders do

- Create one governed data platform for customer, policy and claims data.
- Give every data set an owner and a quality measure.
- Make data available to AI tools securely, not by copying it around.

ASK YOUR TEAM

How many different answers do we have to the question "how many customers do we have"?

MOVE 3

03

Connect everything securely

AI and partners need to reach core systems safely. Secure connections, digital workflows and business rules turn separate systems into one joined-up operation, and let AI decisions flow into real work.

What leaders do

- Expose core systems through secure, managed connections (APIs).
- Run key processes on digital workflows with clear rules.
- Connect partners, such as banks and mobile money, through the same layer.

ASK YOUR TEAM

How long does it take today to connect a new partner to our policy system?

MOVE 4

04

Secure AI by design

AI brings new ways to attack and new ways to leak data. Security must be built into every AI tool from the start: who can use it, what data it can see and what it can do.

What leaders do

- Approve AI tools centrally and block unapproved ones for sensitive data.
- Control access to AI systems as tightly as to core systems.
- Test AI tools against misuse and manipulation before launch.

ASK YOUR TEAM

Which AI tools are our staff using today, and what data are they putting into them?

MOVE 5

Run AI like a product

Many AI pilots never reach production, and many that do are left unmonitored. Treat each AI tool like a product, with an owner, users, a roadmap and a plan to keep it working.

What leaders do

- Give every AI tool a business owner and a technical owner.
- Monitor accuracy, cost and use after launch, not just before.
- Reuse components, so each new AI tool is faster to build.

ASK YOUR TEAM

Which of our AI pilots has a clear owner and a date to go live, or be stopped?

05

MOVE 6

Govern vendors, cloud and cost

Cloud and AI services make it easy to start and easy to overspend. The CIO needs a clear view of every vendor, what it costs, where data goes and how to leave if needed.

What leaders do

- Keep one register of vendors, cloud services and AI tools, with owners and costs.
- Track cloud and AI spending monthly, by team and use case.
- Have an exit plan for every critical vendor.

ASK YOUR TEAM

Which vendor could we not replace within a year if we had to?

06

FROM OUR WORK

Integration and automation at scale

KaribuTech AI's team has built integration and automation for insurers and banks using platforms such as Camunda, IBM BPM, MuleSoft and webMethods, including hundreds of automated workflows for a large African bank. BluFountain Insure, its cloud-native insurance platform, connects life, short-term, medical and funeral cores through shared engines and self-service portals, so new partners and AI tools plug into one layer rather than many.

Hundreds

of automated workflows for a large bank

Cloud-native

BluFountain Insure platform

One layer

for partners, portals and AI


Modernising the core: three paths

There is no single right answer for every core system. The CIO should choose, system by system, whether to keep and connect it, move it to modern infrastructure,

or replace it. The choice depends on risk, cost and how much the business needs to change.

EXHIBIT 3

Three paths for each core system



LOWER RISK →

← BIGGER CHANGE

Source: KaribuTech AI analysis.

How to choose the path

- Business need:** how much must this product change?
- Technical risk:** is the technology still supported?
- Cost:** what does it cost to run and change?
- Skills:** can we still find people who know it?

Rules for any path

- No big bang:** move in stages.
- Connect first:** APIs before replacement.
- Data moves with care:** test every migration.
- Retire on time:** switch off what is replaced.

The architecture for AI, in plain language

An architecture is simply how systems fit together. The pattern below separates the parts that change often, such as customer channels and AI tools, from the parts

that must be stable, such as core systems and data, so each can change without breaking the others.

EXHIBIT 4

Six layers of the AI-ready insurer

LAYER	WHAT IT INCLUDES	WHY IT MATTERS FOR AI
Channels	Apps, WhatsApp, USSD, web, portals for brokers and partners	Where customers and staff meet AI
Workflows and rules	Digital workflows, business rules, software robots	Turns AI outputs into actions with deadlines
AI services	Prediction, generative AI, document and image reading, AI agents	Reusable AI tools, not one-off pilots
Connections	Secure links (APIs), real-time data feeds, partner links	Lets AI reach core systems safely
Core and data	Policy, claims and billing systems; one trusted data platform	The data AI learns from and acts on
Infrastructure	Cloud and own data centres; security and monitoring	Scale, resilience and data sovereignty

Source: KaribuTech AI reference architecture.

Three design principles

01

Connect, don't copy

Reach data through secure links rather than copies.

02

No lock-in

Open standards let you switch vendors.

03

Secure by design

Security built into every layer.


Generative AI and AI agents in IT

Generative AI reads and writes; AI agents plan and carry out tasks. In technology teams, that means explaining and documenting old code, helping write and

test new code, answering IT helpdesk questions and sorting security alerts. Engineers stay responsible for what goes live.

EXHIBIT 5

Six jobs AI can take on in the technology team today



Source: KaribuTech AI analysis. Labels show the level of AI autonomy we recommend.

What stays with engineers

- Approving every change that goes live.
- Security decisions and access to sensitive data.
- Architecture and vendor choices.
- Responding to major incidents.

Guardrails that make it safe

- Approved AI tools only for code and data.
- No secrets or customer data in public AI tools.
- Every AI-written change reviewed and tested.
- Licences checked for AI-generated code.

WHERE TO START

Start with documenting legacy code and an IT helpdesk assistant. Both deliver quickly and reduce the risk of the next modernisation step.

Cloud, cost and data sovereignty

Cloud computing makes AI affordable and scalable, but insurers in many markets must keep some data

in-country, and cloud costs can grow quickly. Most insurers will use a mix. The table compares the options.

EXHIBIT 6

Four hosting options, and when each fits

OPTION	BEST FOR	WHAT TO WATCH
Own data centre	Sensitive data that must stay in-country; stable workloads	Up-front cost; skills; slower to scale
Local or sovereign cloud	Regulated data with cloud flexibility	Fewer services; provider strength
Global public cloud	AI services, analytics and variable workloads	Where data is stored; cost control; lock-in
Hybrid	Keeping sensitive data local and using cloud for the rest	More complex to manage and secure

Source: KaribuTech AI analysis.

Keep cloud costs under control

01

Tag everything

Every service tagged to a team and use case.

02

Review monthly

Spending reviewed with business owners.

03

Switch off waste

Unused services and oversized machines removed.

Choose hosting by the data, not by fashion: sensitive data where the law requires it, everything else where it runs best.

Security for the AI era

AI changes both attack and defence. Attackers use AI to write convincing phishing and clone voices; AI systems themselves can be misused or leak data.

Defenders can use AI to spot threats faster. The table shows the main threats and controls.

EXHIBIT 7

Five AI-era security threats, and the controls that work

THREAT	WHAT HAPPENS	THE CONTROL
AI-written phishing	Convincing messages trick staff into giving access	Strong sign-in checks; phishing tests; AI filtering
Cloned voices and deepfakes	Fake calls authorise payments or changes	Call-back rules; code words; liveness checks
Shadow AI	Staff put company data into unapproved AI tools	Approved tools; data-loss monitoring; training
Attacks on AI systems	AI tools manipulated or data extracted through them	Access controls; testing; monitoring of AI use
Supply chain	A vendor or cloud provider is breached	Vendor reviews; least access; exit plans

Source: IBM, Cost of a Data Breach Report 2025; KaribuTech AI analysis.

What good looks like

01

Detect in hours

Know about an intrusion in hours, not months.

02

Recover in days

Restore critical systems within agreed times.

03

Zero trust

Every user and system verified every time.

The 100-day quick-wins plan for the CIO

Do not start with a large programme. Pick one system or process, deliver a working improvement in 100 days, measure it, then reuse the approach. Each cycle is run

by a small team from technology, the business and security.

EXHIBIT 8

One technology improvement in 100 days: design, build, run, measure

DAYS 1-30

Discover and design

- Map the system, its data and links
- Agree the target design and security
- Baseline cost, speed and incidents

DAYS 31-60

Build

- Build the connection layer or data flow
- Automate testing
- Set up monitoring

DAYS 61-100

Run and measure

- Go live on one product or process
- Compare with the baseline
- Report to the executive team

Source: KaribuTech AI quick-wins approach.

Three quick wins we recommend

QUICK WIN 1

APIs on one core system

- Secure connections in front of the core
- One partner or channel connected
- Access controlled and logged
- Reusable for the next partner

MEASURE

Time to connect a partner, incidents

QUICK WIN 2

Data platform for one domain

- Customer or claims data first
- Owner and quality measures
- Available securely to AI tools
- Feeds reporting and IFRS 17

MEASURE

Data quality score, reports automated

QUICK WIN 3

AI for IT itself

- Legacy code documented
- IT helpdesk assistant
- Security alerts ranked
- Approved AI tools for developers

MEASURE

Tickets resolved, alert response time

The business case for the CIO's AI agenda

The value of the CIO's agenda comes from lower run costs, faster delivery, fewer incidents and the AI value it unlocks across the business. The example below is

illustrative, for an insurer with a \$20m annual technology budget.

EXHIBIT 9

Illustrative example: a \$20m technology budget, 70% spent on running existing systems

Measure	Before	After AI	Annual value
Share of budget on running systems	70%	60%	≈ \$2.0m moved to change
Time to connect a new partner	6 months	6 weeks	Faster growth
Major incidents a year	12	6	Fewer outages
Time to answer IT questions	Days	Minutes	Staff time saved
Total measurable value			≈ \$2.0m + faster delivery

Source: Illustrative KaribuTech AI example, not a forecast. Run costs: ten points of a \$20m budget moved from running old systems to new work.

Where the value comes from

- **Lower run costs:** automation and retiring old systems.
- **Faster delivery:** reusable connections and data.
- **Fewer incidents:** better monitoring and testing.
- **AI value unlocked:** every business AI use case builds on the foundations.

Where the money goes

Connections and workflows	25–30%
Data platform	25–30%
Security	15–20%
Cloud and AI services	10–15%
Skills and change	10%

Typical split of a first-year budget (illustrative)

THREE RULES FOR A CREDIBLE CASE

- 1 Baseline run costs, delivery times and incidents first
- 2 Tie every investment to a business use case
- 3 Report value and cost to the board quarterly


What you need underneath: the technology platform

The CIO owns the engine that every other blueprint in this series relies on. The key idea is one platform with shared layers, so each business area's AI and

automation builds on the same foundations instead of starting again.


EXHIBIT 10

One technology platform for every business area



Source: KaribuTech AI analysis.

Build, buy or partner



Ask one question of every technology decision: will the next AI use case be faster and safer to deliver because of it?

The technology team in the AI era

AI changes how technology teams work. Routine coding, testing and support become faster; architecture, data, security and working with the

business become more important. The CIO leads that shift.

EXHIBIT 11

How the main roles change

ROLE	TODAY	WITH AI
Developer	Writes most code by hand	→ Works with AI assistants; reviews and tests
Architect	Designs individual systems	→ Designs shared platforms and connections
Data engineer	Builds one-off data feeds	→ Builds and runs the trusted data platform
Security analyst	Works through alert queues	→ Investigates the threats AI ranks highest
Service desk	Answers routine IT questions	→ Handles complex issues; runs the IT assistant
CIO	Runs systems within budget	→ Enables business growth and AI at scale

Source: KaribuTech AI analysis.

Who you work with

- **Business heads:** own use cases and value.
- **Chief data officer:** data ownership and quality.
- **Risk and security:** AI and cyber risk.
- **Finance:** cloud and technology costs.
- **Vendors and partners:** platforms and integration.

Skills to build in the team

Architecture and integration

Data engineering

AI security

Cloud cost management

Working with the business

Risks and guardrails for the technology agenda

Technology programmes fail more often from risk than in the AI era, and the controls that manage them. from lack of ambition. These are the risks the CIO owns

EXHIBIT 12

Five technology risks, and the control for each

RISK	WHAT CAN GO WRONG	THE CONTROL
Big-bang failure	A large core replacement runs late and over budget	Move in stages; connect first; retire on time
Shadow AI	Staff use unapproved AI tools with company data	Approved tools; data-loss controls; training
Vendor lock-in	One provider becomes impossible to replace	Open standards; exit plans; own your data
Cyber attack	Breach of core systems or AI tools	Zero trust; monitoring; tested recovery plans
Cost overruns	Cloud and AI spending grows unchecked	Tagging; monthly reviews; business ownership

Source: KaribuTech AI analysis; IAS, Application Paper on the supervision of AI (2025).

Eight questions to ask your team

01

Which core system poses our biggest technology risk?

02

How long does it take to connect a new partner?

03

Which AI tools are staff using without approval?

04

What share of our budget runs old systems?

05

Could we leave our main cloud provider if needed?

06

How quickly would we detect an intruder?

07

Which AI pilots have no owner or go-live date?


08

Is our data platform ready for AI and IFRS 17?

The technology scorecard and 12-month roadmap

EXHIBIT 13


Six numbers to run technology for the AI era, plus four for AI



Source: KaribuTech AI analysis. Gold: keep within a target band; green: higher is better; red: lower is better.

EXHIBIT 14

Four quarters to AI-ready technology foundations



Source: KaribuTech AI analysis.

WHERE TO START ON MONDAY

List every AI tool in use across the insurer, approved or not. It is the fastest way to see both the risk and the opportunity.

About the authors



Kennedy Chengeta, PhD

Co-founder, KaribuTech AI

Enterprise Architect and AI Leader · Camunda Chapter Lead, Africa

Leads AI, automation and integration programmes for insurers and banks across Southern Africa.



Justice Muguda

Chief Technology Officer, KaribuTech AI


Cloud, integration and cybersecurity

Leads KaribuTech AI's technology and engineering, from cloud platforms to secure connections between systems.

With contributions from the KaribuTech AI Engineering Team. To discuss this report, or to run a 100-day technology and AI diagnostic, contact KaribuTech AI at kaributech.co.za.

References

1. Stanford HAI (2025), Artificial Intelligence Index Report 2025.
2. IBM (2025), Cost of a Data Breach Report 2025.
3. Histories of COBOL (1959).
4. IBM, System/360 announcement (1964).
5. Amazon Web Services, launch of cloud services (2006).
6. IAIS (2025), Application Paper on the supervision of artificial intelligence, July 2025.
7. European Union (2024), Regulation (EU) 2024/1689, the Artificial Intelligence Act.
8. KaribuTech AI, BluFountain Insure platform.
9. KaribuTech AI, The AI Blueprint for the Insurance CEO (2026).
10. KaribuTech AI, reference architecture and quick-wins approach (proprietary).



About KaribuTech AI

KaribuTech AI is a South African technology company specialising in artificial intelligence, machine learning, intelligent automation and integration for insurers and banks across Africa. Our teams design and build intelligent workflows, decision engines, document and computer-vision AI, and cloud-native core platforms, including the **BluFountain Insure** platform: an insurance AI

marketplace with core systems for life, short-term, medical, funeral and investment business. We combine deep process expertise (Camunda, IBM BPM, MuleSoft, webMethods and document automation platforms) with practical AI engineering, so that AI moves out of the lab and into the operations that customers feel.

Disclaimer: this report provides general information and does not constitute legal, actuarial or investment advice. Third-party figures are cited beneath each exhibit and have not been independently verified by KaribuTech AI. Illustrative figures are not forecasts.



KaribuTech AI

INSURANCE · LEADERSHIP SERIES · SEPTEMBER 2026

The AI Blueprint for the Head of HR

Reskilling people, redesigning jobs and using AI fairly inside HR: the people side of AI in insurance


By Kennedy Chengeta, PhD, and Justice Muguda
with the KaribuTech AI Engineering Team

kaributech.co.za



Contents

01	Introduction	46
02	Why HR must lead on AI	47
03	Lessons from history	48
04	Where the people cycle breaks	49
05	Six moves for the Head of HR	50
06	Grow, change or redeploy	52
07	Reskilling pathways	53
08	Generative AI and AI agents in HR	54
09	Using AI in HR fairly	55
10	Skills for the AI era	56
11	The 100-day quick-wins plan	57
12	The business case	58
13	What you need underneath	59
14	The HR team	60
15	Risks and guardrails	61
16	The scorecard and 12-month roadmap	62
17	About the authors and KaribuTech AI	63



About this report

This report is written for chief people officers, heads of HR and learning leaders at insurers, reinsurers and pension funds in developing markets.

AI will change almost every job in insurance. This report shows HR how to lead that change and how to use AI inside HR fairly: the pressures, six moves, which roles grow, change or need redeployment, reskilling pathways, fair use of AI in HR, skills for the AI era, a 100-day plan, the business case, the team, the risks and a scorecard.

External data are sourced beneath each exhibit, and a full list of references appears at the end. Frameworks and illustrative figures without an external source are the authors' own.

WRITTEN FOR

Chief people officers

HR business partners

Learning leaders

Executives

Introduction

AI is not only a technology change. It is a people change. HR has two jobs: leading the workforce through the shift, and using AI inside HR itself, fairly.

In most insurers, AI projects are run by IT and the business, while HR is asked to deal with the consequences afterwards: new skills nobody planned for, anxious staff and roles that no longer fit. That is the wrong order.

When HR leads, the insurer plans which roles will grow, change or shrink, reskills people before automating their work, redesigns jobs around people and AI working together, and uses AI in hiring and development without bias.

This report sets out **six moves** for the head of HR, how to decide which roles **grow, change or need redeployment, reskilling pathways, fair use of AI in HR, skills for the AI era, a 100-day plan** and a **scorecard**.

THREE QUESTIONS THIS REPORT ANSWERS

01

Which roles will AI grow, change or shrink?

02

How do we reskill people before we automate their work?

03

How can HR itself use AI without bias?

AT A GLANCE: SIX MOVES FOR HR

1 Plan the workforce for AI

2 Reskill before you automate

3 Redesign jobs, not just tasks

4 Use AI in HR fairly

5 Lead the change and bring people with you

6 Build an AI-literate culture

Why HR must lead on AI

Five pressures make AI an HR priority, not only an IT project. Each is also a chance to build a

stronger, more capable workforce.

Skills are changing fast

Many of the skills insurance jobs need today will change within a few years, especially in claims, underwriting and operations.

Staff are anxious

Uncertainty about AI and jobs hurts engagement, trust and retention if nobody explains the plan.

Talent is scarce

Data, AI and actuarial skills are in short supply and in demand from banks, telcos and global firms.

AI in HR is high-risk

Using AI in hiring, promotion or dismissal is classed as high-risk in the EU AI Act, and bias in HR AI damages trust and brand.

Staff already use AI

Employees use public AI tools with or without permission, often with company or customer data.

EXHIBIT 1

The scale of the shift

39%

of workers' core skills expected to change by 2030

170m

new jobs created globally by 2030

92m

jobs displaced globally by 2030

1 in 5

organisations had a breach involving unapproved "shadow" AI

Source: World Economic Forum, Future of Jobs Report 2025; IBM, Cost of a Data Breach Report 2025; EU AI Act (2024).

AI will not replace most insurance professionals. Insurance professionals who use AI will replace those who do not. HR decides which of those our people become.

Lessons from history: technology changes jobs, and people adapt when helped

Each wave of workplace technology has destroyed some tasks and created new ones. The outcome for

workers depended on how quickly they were helped to adapt.

1811

The Luddites

As textile machines spread across England in the late 18th and early 19th centuries, skilled weavers who lost their livelihoods protested and broke machines. Change without support created lasting resistance.

Today: explain the AI plan early and fund reskilling first.

THE LESSON

Change without a plan for people creates fear and resistance.

1913

The moving assembly line

Ford's moving assembly line transformed car-making and the skills factory workers needed. New roles in maintenance, planning and supervision grew.

Today: new roles grow around AI in data, oversight and care.

THE LESSON

Technology reshapes jobs; new roles appear around it.

1979

The spreadsheet

VisiCalc, the first spreadsheet program for personal computers, automated hours of manual calculation. Bookkeeping clerks declined, while analysts and accountants who used spreadsheets grew.

Today: professionals who use AI well will move ahead.

THE LESSON

People who learn the new tool move up; those who do not are left behind.

Today

AI at work

AI now reads, writes and predicts, touching office work in claims, underwriting, finance and HR itself.

Today: AI literacy for everyone; deep skills for some.

THE LESSON

Reskill early, redesign jobs, and bring people with you.

Source: Histories of the Luddite movement (1811–1816), Ford's moving assembly line (1913) and VisiCalc (1979).

The spreadsheet did not end accounting; it changed who did well in it. AI will do the same across insurance, and HR can decide how many of our people come out ahead.

Where the people cycle breaks, and how AI helps

Every employee passes through the same cycle, from workforce planning to leaving or moving roles. At each

step, HR can lose time, fairness or talent. The exhibit shows where it typically breaks and what AI does.

EXHIBIT 2

The people cycle: eight steps from planning to redeployment

	WHERE IT BREAKS TODAY	WHAT AI DOES
Workforce planning How many people, with which skills	Plans based on headcount, not skills	Skills mapped against future needs
Hiring Finding and selecting people	Slow screening; hidden bias	Faster matching, checked for fairness
Onboarding Getting new staff productive	Paper forms; slow access	Digital onboarding and a help assistant
Learning Building skills	Generic courses, low take-up	Personal learning paths for AI skills
Performance Goals and feedback	Annual reviews, little data	Regular feedback; fair, explained insights
Pay and benefits Rewarding people fairly	Pay gaps unseen	Pay equity checked across groups
Wellbeing and relations Engagement and concerns	Surveys read too late	Themes found early across feedback
Moves and exits Internal moves, redeployment, leaving	Redundancy before redeployment	Internal roles matched to skills first

Source: KaribuTech AI analysis.

The goal is not an automated HR function. It is a workforce that is ready for AI, and treated fairly along the way.

Six moves for the Head of HR

HR leaders who handle AI well do six things differently. Each is about people, skills and

trust, not about choosing a technology.

MOVE 1

01

Plan the workforce for AI

HR should know which roles AI will grow, change or shrink before the technology arrives, not after. That starts with a map of today's skills and an honest view of what each role will need in two to three years.

What leaders do

- Build a skills inventory for every role, not only a headcount plan.
- Agree with business leaders which roles grow, change or shrink.
- Update the plan every six months as AI use cases go live.

ASK YOUR TEAM

Which of our roles will change most in the next two years, and have we told the people in them?

MOVE 2

02

Reskill before you automate

Redundancy followed by hiring is slow, costly and damages trust. Reskilling people into new roles, such as claims clerks becoming claims analysts, keeps knowledge in the business and shows staff that AI is a path, not a threat.

What leaders do

- Fund reskilling before, not after, each automation project.
- Create clear pathways from shrinking roles to growing ones.
- Give people time and support to learn, not only online courses.

ASK YOUR TEAM

For every role AI will shrink, do we have a pathway into a role it will grow?

MOVE 3

03

Redesign jobs, not just tasks

AI takes over tasks, not whole jobs. The value comes from redesigning roles around what people do best, such as judgement, empathy and complex cases, with AI handling the routine.

What leaders do

- Redesign roles with the people who do them.
- Write new job descriptions that include working with AI.
- Update pay and grading so new skills are recognised.

ASK YOUR TEAM

Have we redesigned any role around AI yet, or only removed tasks from it?

MOVE 4

04

Use AI in HR fairly

AI can help HR screen candidates, plan learning and answer questions. But AI in hiring and promotion can repeat past bias and is treated as high-risk by regulators. Every HR use of AI needs testing, explanation and human decisions.

What leaders do

- Test hiring and promotion tools for bias before and during use.
- Keep a person responsible for every hiring, promotion and exit decision.
- Tell candidates and employees when AI is used, and how.

ASK YOUR TEAM

Could we explain to a rejected candidate how AI was used in their application?

MOVE 5

05

Lead the change and bring people with you

People fear what they do not understand. Clear, honest communication, early involvement of staff and worker representatives, and visible support turn anxiety into energy.

What leaders do

- Explain the AI plan openly: what changes, when, and how people will be supported.
- Involve staff and worker representatives in redesigning work.
- Celebrate people who learn and use AI well.

ASK YOUR TEAM

What have we told our staff about AI and their jobs, and do they believe it?

MOVE 6

06

Build an AI-literate culture

Every employee, from the board to the front line, needs to understand what AI can and cannot do, how to use approved tools safely and when to question an AI output.

What leaders do

- Run AI basics for everyone, and deeper training for those who use AI daily.
- Provide approved AI tools, so people do not turn to unsafe ones.
- Build a network of AI champions across departments.

ASK YOUR TEAM

What share of our staff could explain what AI is, and is not, to a customer?

FROM OUR WORK

Building AI skills in our markets

KaribuTech AI invests in AI skills beyond its own projects. Its team has mentored Zimbabwean students in AI and process automation since 2021, supports the Deep Learning Zimbabwe community and leads the Camunda Chapter for Africa. Every client project includes hands-on training, so that insurers' own people can run and improve the workflows and AI tools after go-live.

Since 2021

mentoring students in AI and automation

Community

Deep Learning Zimbabwe and Camunda Africa

Every project

includes training for client teams


Grow, change or redeploy: planning roles around AI

AI affects roles differently. Some grow, most change, and some shrink. Sorting roles into three groups lets HR

plan hiring, reskilling and redeployment in good time, instead of reacting after the technology arrives.

EXHIBIT 3

Three groups of roles, and what HR does for each



GROWING ROLES →

← ROLES TO RESKILL FROM

Source: KaribuTech AI analysis.

How to decide the group

- **Tasks AI can do:** how much of the role is routine?
- **Judgement needed:** how much depends on people?
- **Customer contact:** how much empathy is involved?
- **Business plans:** which areas are growing?

Principles HR should hold to

- **No surprises:** people hear about changes early.
- **Reskill first:** redeployment before redundancy.
- **Fair selection:** clear criteria for new roles.
- **Support for all:** including those who leave.

Reskilling pathways: from today's roles to tomorrow's

The most practical way to protect people and keep knowledge in the business is to build clear pathways

from roles that are shrinking to roles that are growing. The table shows six pathways common in insurance.

EXHIBIT 4

Six reskilling pathways in insurance

FROM	TO	SKILLS TO BUILD	TIME TO RESKILL
Claims clerk	Claims analyst	Reviewing AI decisions; data basics; customer care	3–6 months
Underwriting assistant	Underwriter using AI	Risk judgement; reading AI triage; product rules	6–12 months
Call-centre agent	Customer advocate	Complex cases; empathy; digital channels	3–6 months
Data entry clerk	Data quality steward	Data checks; tools; process knowledge	3–6 months
Actuarial analyst	Actuarial data scientist	Programming; machine learning; explaining models	12–18 months
Branch sales	Digital channel adviser	Digital sales; WhatsApp and app journeys; partners	3–6 months

Source: KaribuTech AI analysis; times are indicative.

What makes reskilling work

01

Time to learn

Protected hours in the working week, not only evenings.

02

Learning by doing

Real projects with coaching, not just courses.

03

A job at the end

A named role for people who complete the pathway.

Generative AI and AI agents in HR

Generative AI reads and writes; AI agents plan and carry out tasks. In HR, that means drafting job descriptions, answering employees' policy questions,

creating learning content and finding themes in feedback. Decisions about people stay with people.

EXHIBIT 5

Six jobs AI can take on in HR today

Answer policy questions Answers employees' questions on leave, benefits and policies, any time ACT WITHIN LIMITS	Draft job descriptions Drafts inclusive job descriptions and adverts for HR to review RECOMMEND	Create learning content Turns procedures and policies into short courses and quizzes ASSIST
Find themes in feedback Summarises surveys and exit interviews into clear themes ASSIST	Support onboarding Guides new staff through their first weeks with answers and reminders ACT WITHIN LIMITS	Match internal roles Suggests internal roles that fit an employee's skills RECOMMEND

Source: KaribuTech AI analysis. Labels show the level of AI autonomy we recommend.

What stays with people

- Every hiring, promotion and exit decision.
- Performance ratings and pay decisions.
- Disciplinary and grievance cases.
- Conversations about change and wellbeing.

Guardrails that make it safe

- **Bias tested** before any tool touches hiring or promotion.
- **No employee data in public AI tools.**
- **Transparency:** staff know when AI is used.
- **Every output reviewed** before it affects anyone.

WHERE TO START

Start with an HR policy assistant and learning content. They save time for HR and help every employee, with little risk.

Using AI in HR fairly

AI in HR affects people's livelihoods, so the bar is high. The EU AI Act classes AI used in recruitment, promotion and termination as high-risk, and similar expectations

are spreading. The table shows the main HR uses, their risks and the controls to apply.

EXHIBIT 6

Four HR uses of AI, their risks and the controls

USE	RISK	CONTROL
Hiring and screening	Repeats past bias against gender, age, region or background	Bias testing; human decisions; tell candidates AI is used
Performance and promotion	Judges people on narrow data they cannot see or challenge	Explained insights; manager decides; right to challenge
Monitoring productivity	Feels like surveillance; damages trust	Clear limits; staff consultation; no hidden monitoring
Pay and benefits	Hidden pay gaps grow through automated decisions	Regular pay-equity checks; transparent rules

Source: EU AI Act (2024), Annex III; KaribuTech AI analysis.

Three rules for HR AI

01

A person decides

No hiring, promotion or exit decision by AI alone.

02

Test for fairness

Check outcomes across groups before and after launch.

03

Be open

Tell people when and how AI is used about them.

AI in HR should widen opportunity, not narrow it. If a tool cannot be explained to the person it affects, it should not be used.

Skills for the AI era: who needs what

Not everyone needs to become a data scientist. But everyone needs some AI literacy, and some groups

need much more. The table shows the skills each group needs and how to build them.

EXHIBIT 7

AI skills by group, and how to build them

GROUP	WHAT THEY NEED	HOW TO BUILD IT
Board and executives	What AI can do, its risks and how to govern it	Board briefings; AI in strategy sessions
Managers	Leading teams through change; judging AI use cases	Leadership programmes; peer learning
Professionals	Using AI tools well and questioning their output	Hands-on training on real work
All staff	What AI is and is not; safe use of approved tools	Short AI basics course for everyone
Technology and data teams	Building, testing and running AI safely	Technical training; certifications; mentoring

Source: KaribuTech AI analysis.

Measure what matters

01

Reach

Share of staff who completed AI basics.

02

Use

Share of staff using approved AI tools each week.

03

Impact

Hours saved and quality improved in real work.


The 100-day quick-wins plan for HR

Do not start with a large programme. Pick one area, deliver a working improvement in 100 days, measure it,

then build on it. Each cycle is run by a small team from HR, the business and technology.


EXHIBIT 8

One HR improvement in 100 days: design, build, run, measure



Source: KaribuTech AI quick-wins approach.

Three quick wins we recommend



The business case for AI in HR

The value of HR's AI agenda comes from keeping skills in the business, lowering the cost of hiring and redundancy, freeing HR time and keeping good people.

The example below is illustrative, for an insurer with 1,000 staff.

EXHIBIT 9

Illustrative example: an insurer with 1,000 staff and 60 roles changing through AI

Measure	Before	After AI	Annual value
60 changing roles	Exit and rehire	Reskill and redeploy	≈ \$1.26m saved
HR time on routine queries	30% of HR time	15% of HR time	More time for people
Time to fill roles	60 days	35 days	Faster hiring
Attrition of key talent	15%	11%	Knowledge kept
Total measurable value			≈ \$1.26m + retention

Source: Illustrative KaribuTech AI example, not a forecast. Redundancy plus replacement hiring at about \$25,000 per role, versus reskilling at about \$4,000, for 60 roles.

Where the value comes from

- **Lower change costs:** reskilling instead of redundancy and hiring.
- **Knowledge kept:** experienced people move into new roles.
- **Freed HR time:** routine questions answered automatically.
- **Better retention:** staff see a future with AI.

Where the money goes

Reskilling programmes	35–45%
AI literacy for all staff	15–20%
HR tools and assistant	15–20%
Skills data and planning	10–15%
Change and communication	10%

Typical split of a first-year budget (illustrative)

THREE RULES FOR A CREDIBLE CASE

- 1 Baseline skills, attrition and change costs first
- 2 Count only savings the CFO and CHRO sign
- 3 Report people outcomes, not only savings, every quarter


What you need underneath: the people platform

The head of HR does not need to be a technologist, but should know what the platform looks like. The key idea

is one view of skills across the insurer, connected to hiring, learning and workforce planning.

EXHIBIT 10

One people platform, built around skills



Source: KaribuTech AI analysis.

Build, buy or partner

Buy

HR and learning platforms, recruitment tools, HR assistants

Build

Your skills framework, pathways and workforce plans

Partner

Universities, training providers, industry bodies

Ask one question of every HR technology decision: will it help our people build the skills they need, fairly?

The HR team in the AI era

AI takes over much of HR's routine administration. The team's work moves to workforce planning, reskilling,

change leadership and making sure AI is used fairly. The head of HR leads that shift.

EXHIBIT 11

How the main roles change

ROLE	TODAY	WITH AI
HR business partner	Handles day-to-day people issues	→ Plans skills and change with business leaders
Recruiter	Screens CVs by hand	→ Focuses on candidates, with fair AI support
Learning and development	Runs a course catalogue	→ Designs reskilling pathways and AI literacy
HR administrator	Answers routine questions	→ Manages the HR assistant and exceptions
Employee relations	Handles cases as they arise	→ Leads consultation on AI changes
Chief people officer	Oversees HR services	→ Leads the workforce through the AI shift

Source: KaribuTech AI analysis.

Who you work with

- **Executive team:** agree which roles grow, change or shrink.
- **IT and data teams:** HR systems and approved AI tools.
- **Risk and compliance:** fair use of AI and data privacy.
- **Worker representatives:** consultation on change.
- **Finance:** funding reskilling and measuring value.

Skills to build in the team

Workforce and skills planning

Change leadership

People analytics

AI fairness testing

Designing learning

Risks and guardrails for AI and people

The biggest risks of AI for HR are not technical. They are about trust, fairness and people left behind. These

are the risks the head of HR owns, and the controls that manage them.

EXHIBIT 12

Five people risks of AI, and the control for each

RISK	WHAT CAN GO WRONG	THE CONTROL
Biased HR decisions	AI in hiring or promotion treats groups unfairly	Bias testing; human decisions; clear explanations
Loss of trust	Staff fear AI and stop engaging	Open communication; involvement; visible support
Skills left behind	People in changing roles are not reskilled in time	Reskilling funded before automation
Employee data misuse	Personal data used without consent or leaked	Consent; access controls; approved tools only
Shadow AI	Staff use unapproved AI tools with company data	Approved tools; clear rules; training

Source: KaribuTech AI analysis; IAIS, Application Paper on the supervision of AI (2025).

Eight questions to ask your team

01 Which roles will change most in the next two years?

02 Do we have a pathway for every shrinking role?

03 Have our HR AI tools been tested for bias?

04 What have we told staff about AI and their jobs?

05 What share of staff have completed AI basics?

06 Which AI tools are staff using without approval?


07 How do we support people who choose to leave?

08 Are worker representatives involved in redesigning work?

The HR scorecard and 12-month roadmap

EXHIBIT 13


Six numbers to run HR through the AI shift, plus four for the AI era



Source: KaribuTech AI analysis. Gold: keep within a target band; green: higher is better; red: lower is better.

EXHIBIT 14

Four quarters to an AI-ready workforce



Source: KaribuTech AI analysis.

WHERE TO START ON MONDAY

Ask each business head to name three roles AI will change most in the next two years, and start the first reskilling pathway with the people in those roles.

About the authors



Kennedy Chengeta, PhD

Co-founder, KaribuTech AI

Enterprise Architect and AI Leader · Camunda Chapter Lead, Africa

Leads AI, automation and integration programmes for insurers and banks across Southern Africa.



Justice Muguda

Chief Technology Officer, KaribuTech AI


Cloud, integration and cybersecurity

Leads KaribuTech AI's technology and engineering, from cloud platforms to secure connections between systems.

With contributions from the KaribuTech AI Engineering Team. To discuss this report, or to run a 100-day people and AI diagnostic, contact KaribuTech AI at kaributech.co.za.

References

1. World Economic Forum (2025), Future of Jobs Report 2025.
2. European Union (2024), Regulation (EU) 2024/1689, the Artificial Intelligence Act, Annex III.
3. IBM (2025), Cost of a Data Breach Report 2025.
4. Histories of the Luddite movement (1811–1816).
5. Histories of Ford's moving assembly line (1913).
6. Histories of VisiCalc, the first personal-computer spreadsheet (1979).
7. IAIS (2025), Application Paper on the supervision of artificial intelligence, July 2025.
8. KaribuTech AI, community and mentoring programmes.
9. KaribuTech AI, The AI Blueprint for the Insurance CEO (2026).
10. KaribuTech AI, quick-wins approach (proprietary).



About KaribuTech AI

KaribuTech AI is a South African technology company specialising in artificial intelligence, machine learning, intelligent automation and integration for insurers and banks across Africa. Our teams design and build intelligent workflows, decision engines, document and computer-vision AI, and cloud-native core platforms, including the **BluFountain Insure** platform: an insurance AI

marketplace with core systems for life, short-term, medical, funeral and investment business. We combine deep process expertise (Camunda, IBM BPM, MuleSoft, webMethods and document automation platforms) with practical AI engineering, so that AI moves out of the lab and into the operations that customers feel.

Disclaimer: this report provides general information and does not constitute legal, actuarial or investment advice. Third-party figures are cited beneath each exhibit and have not been independently verified by KaribuTech AI. Illustrative figures are not forecasts.



KaribuTech AI

INSURANCE · LEADERSHIP SERIES · SEPTEMBER 2026


The AI Blueprint for the Head of New Business

Faster quotes, paperless onboarding, smarter underwriting and policies that stay on the books

By Kennedy Chengeta, PhD, and Justice Muguda
with the KaribuTech AI Engineering Team

kaributech.co.za





KaribuTech AI

About this report

This report is written for heads of new business, sales and underwriting at insurers in developing markets, and for the executives they report to.

It shows how AI can turn the new business funnel, from first contact to a policy that is still in force a year later, into a fast, low-cost and well-controlled process: where the leaks are, six moves to fix them, how to underwrite and onboard with AI, how to work with partners, a 100-day quick-wins plan, the team you need, the risks to manage and a scorecard to run it.

External data are sourced beneath each exhibit, and a full list of references appears at the end. Frameworks and illustrative figures without an external source are the authors' own.

WRITTEN FOR

Heads of new business

Sales leaders

Underwriters

Distribution

Contents

01	Introduction	66
02	Why new business must change	67
03	Lessons from history	68
04	Where the funnel leaks	69
05	Six moves for the Head of New Business	70
06	Underwriting with AI: three lanes	72
07	Paperless onboarding	73
08	Generative AI and AI agents in new business	74
09	Distribution: every channel, one engine	75
10	Designing products that sell themselves	76
11	The 100-day quick-wins plan	77
12	The business case	78
13	What you need underneath	79
14	The new business team	80
15	Risks and guardrails at the point of sale	81
16	The scorecard and 12-month roadmap	82
17	About the authors and KaribuTech AI	83

Introduction

New business is where an insurer wins or loses the customer. Today, too many applications leak away between the first quote and the first premium, and too many of the policies that do get sold lapse within a year.

Customers now expect a quote in minutes on the phone they already use. Instead, many still meet paper forms, repeated requests for the same documents and underwriting decisions that take weeks. Underwriters spend much of their day on admin rather than judgement, and distribution costs eat the margin on small policies.

AI changes this. It can pre-fill forms, read documents, check identity, triage risks, recommend the right cover and warn when a new customer is likely to lapse. Used well, it makes the first yes faster, cheaper and fairer, and keeps a person in charge of every difficult case.

This report sets out **six moves** for the head of new business, a practical view of **underwriting and onboarding with AI**, how to work with **distribution partners**, a **100-day quick-wins plan**, the team and guardrails you need, and a **scorecard** to run the function. It is written in plain language, for business leaders rather than technologists.

THREE QUESTIONS THIS REPORT ANSWERS

- 01** Where do we lose applicants between quote and first premium?
- 02** How fast could we quote, onboard and issue a policy?
- 03** How do we keep new policies on the books past year one?

AT A GLANCE: SIX MOVES FOR NEW BUSINESS

- 1** Meet customers where they already are
- 2** Quote in minutes, not days
- 3** Make onboarding paperless
- 4** Underwrite with more data and fewer forms
- 5** Protect the first year
- 6** Measure the whole funnel

Why new business must change

Four pressures are squeezing the traditional new business model at the same time. Each

one is also an opening for the insurers that move first.

Customers expect instant

People who can open a bank account or borrow money on their phone in minutes will not wait weeks for a policy. Every extra form and every extra day loses applicants.

Distribution is too expensive

Commission and branch costs can consume much of the premium on a small policy. Without cheaper channels, the mass market stays uninsured.

Underwriters are buried in admin


Skilled underwriters spend a large share of their time keying data, chasing documents and answering emails instead of judging risk.

Rules are getting tighter

Know-your-customer, anti-money-laundering, data-protection and fair-treatment rules all apply at the point of sale, and AI decisions must be explainable.

EXHIBIT 1

Underwriters spend too much of their time on work AI can do



Source: Accenture underwriting survey, as reported by Insurance Business; Capgemini; McKinsey & Company.

The prize

\$400bn

McKinsey

AI value at stake across pricing, underwriting and promotion globally

\$2tn+

GSMA

moved through mobile money in 2025: a ready-made channel for premiums

Weeks → days

KaribuTech AI

underwriting turnaround achieved with workflow, rules and document AI

Source: McKinsey & Company; GSMA, State of the Industry Report on Mobile Money (2026); KaribuTech AI project experience.

Lessons from history: selling insurance has been reinvented before

Every big leap in insurance growth has come from a new way of pricing or selling cover, not from a new kind of risk. The 18th century brought pricing by data, the 19th brought mass-market selling door to door, and the

21st brought premiums paid by phone. AI is the next of these shifts, and the insurers that led the earlier ones were rarely the incumbents.

1762

Pricing by age

The Equitable Society in London began charging premiums that varied by age, using mortality data, so life cover could be priced fairly and sold at scale.

Today: AI prices customers with no history using mobile and payment data.

THE LESSON

Better data lets you price, and therefore sell, to more people.

1850s

The man at the door

In Britain, Prudential and others built "industrial assurance": agents called on working families every week to collect a few pennies. Millions were insured for the first time.

Today: WhatsApp, USSD and partner apps are the new door.

THE LESSON

Growth comes from going to the customer, on their payment rhythm.

2007

Premiums by phone

Safaricom launched M-Pesa in Kenya in 2007. Mobile money soon became a way to pay small premiums and receive claims with no bank account and no agent visit.

Today: premiums from daily to monthly, with no bank account.

THE LESSON

The channel that customers already use becomes the distribution channel.

Today

AI at the point of sale

AI can price, onboard and underwrite in minutes on WhatsApp, in a bank app or at a checkout, at a cost that makes a very small policy worth selling.

Today: quote, onboard and issue in minutes; people decide the hard cases.

THE LESSON

The door-to-door agent's reach, at the cost of a text message.

Source: M. E. Ogborn, Equitable Assurances (1962); histories of industrial assurance in Britain; Safaricom (2007); KaribuTech AI analysis.

Each shift rewarded the insurers who redesigned how cover was sold, not those who sold the old way a little faster. The same will be true of AI.

Where the funnel leaks, and how AI plugs it

Every stage of new business loses some customers. The exhibit follows one application from first contact to a policy still in force after a year, shows where it

typically leaks, and what AI does about it. The widths are illustrative; your own data will show where your biggest leaks are.

EXHIBIT 2

The new business funnel: eight stages from interest to a policy that stays



Source: KaribuTech AI analysis; funnel widths are illustrative.

The cheapest new customer is the one you already nearly won. Fix the biggest leak first; it is usually identity checks or underwriting, not marketing.

Six moves for the Head of New Business

The leaders who use AI well in new business do six things differently. Each is about the

customer journey and the way the team works, not about choosing a technology.

MOVE 1

01

Meet customers where they already are

Customers will not come to a branch or a long web form. They live on WhatsApp, USSD, mobile money and in their bank, telco and retail apps. New business has to be sold there, in their language, at the moment they need cover.

What leaders do

- Launch quote-and-apply on WhatsApp and USSD, not only on the website.
- Embed simple cover in partners' journeys: loans, bank accounts, trips, checkouts.
- Give agents and brokers an AI assistant that prepares quotes and answers product questions.

ASK YOUR TEAM

Which three channels bring us the most applications today, and which would bring the most tomorrow?

MOVE 2

02

Quote in minutes, not days

Most of the data needed for a quote already exists: in the customer's previous applications, in partner data and in public records. AI can pre-fill it, ask only what is missing and recommend the right level of cover.

What leaders do

- Pre-fill applications from data the customer has already shared, with consent.
- Use next-best-offer models to suggest the right cover, not the most expensive.
- Publish a target: quote in under five minutes for simple products.

ASK YOUR TEAM

How many questions does a customer answer before seeing a price, and how many do we really need?

MOVE 3

03

Make onboarding paperless

Identity checks, document collection and signatures are where most applications stall. AI can read identity documents and payslips, match a selfie to an ID, screen against sanctions lists and generate the policy for digital signature.

What leaders do

- Digital KYC with ID and selfie checks, and automatic anti-money-laundering screening.
- Let customers upload documents on any channel and at any point, never twice.
- Digital signatures and automatically generated policy documents.

ASK YOUR TEAM

Why do we ever ask a customer for the same document twice?

MOVE 4

04

Underwrite with more data and fewer forms

AI can sort applications by complexity, approve simple risks automatically within agreed rules, and prepare complex cases for underwriters with the key facts already extracted. Underwriters spend their time on judgement, not typing.

What leaders do

- Triage every application: straight-through, light-touch or full underwriting.
- Use new data, such as mobile, satellite and health-app data, only with consent and explainable models.
- Keep an underwriter on every decline and every exception.

ASK YOUR TEAM

What share of our applications could be approved safely without a person touching them?

MOVE 5

05

Protect the first year

A sale is not a success until the policy is still in force a year later. Early lapses and not-taken-up policies waste acquisition cost and often signal mis-selling. AI can spot the warning signs at the point of sale and in the first months.

What leaders do

- Predict early lapse at the point of sale, and check the cover fits the customer.
- Collect the first premium on the customer's payday, by mobile money or debit order.
- Run a welcome journey: confirmation, reminders and a check-in after 60 days.

ASK YOUR TEAM

Of the policies we sold last year, how many are still in force, and do we pay commission on the ones that are not?

MOVE 6

06

Measure the whole funnel

Most new business teams measure sales volumes. The best measure every stage, from first contact to month twelve, and use the data to find where work gets stuck and which offers really work.

What leaders do

- Build one dashboard from first contact to 12-month persistency, by channel and product.
- Use process analytics to see where applications wait and why.
- Test offers and journeys against each other, and keep what works.

ASK YOUR TEAM

Where exactly, by stage and by channel, did we lose applicants last month?

FROM OUR WORK

Underwriting from weeks to days

Using an insurance suite built on digital workflows, business rules and document reading, the KaribuTech AI engineering team, working with client teams, cut underwriting turnaround from weeks to days. The same application journey runs on USSD, WhatsApp and mobile money, so customers without smartphones are served by the same automated process, and underwriters see only the cases that need them.

Weeks → days

underwriting turnaround

One journey

across USSD, WhatsApp, app and branch

People on exceptions

underwriters focus on the hard cases


Underwriting with AI: three lanes, one set of rules

The fastest way to cut turnaround without adding risk is to stop treating every application the same. AI sorts applications into three lanes. The rules for each lane are

set by underwriters and actuaries, and every decline or exception still goes to a person.

EXHIBIT 3

Three underwriting lanes: let AI take the simple cases so people can focus on the hard ones



MORE SPEED →

← MORE PERSONAL JUDGEMENT

Source: KaribuTech AI analysis.

New data, used fairly

- **Mobile and payment data:** airtime top-ups and mobile-money history for customers with no credit record.
- **Satellite and mapping data:** flood, fire and crop risk for property and agriculture, without a site visit.
- **Vehicle-tracking data:** how and when a vehicle is driven, for motor cover.
- **Health-app and wellness data:** only with explicit consent, for life and health cover.

Guardrails underwriters should insist on

- **Consent first:** customers know what data is used and agree to it.
- **Explainable decisions:** every decline and loading has a reason a customer can understand.
- **Fairness tests:** outcomes checked by gender, age, region and income before and after launch.
- **Human review:** underwriters own the rules, the exceptions and every decline.

Paperless onboarding: from application to signed policy

Onboarding is where most applications stall and most cost hides. Breaking it into four steps, and automating each one, lets a customer start on one channel, finish

on another, and see where their application stands at every point.


EXHIBIT 4

Four onboarding steps, each with a clear owner, deadline and AI support



Source: KaribuTech AI onboarding approach; targets are illustrative and depend on product and regulation.

What the workflow adds behind the scenes



Start with one retail product and one business product. Once their journeys work end to end, every other product follows the same template, faster and cheaper each time.

Generative AI and AI agents in new business

Generative AI reads and writes; AI agents plan and carry out tasks. In new business, that means reading applications and broker emails, drafting quotes and proposals, answering customer questions in local

languages, and chasing missing documents. The rule is simple: give AI more freedom only where the risk is low and the rules are clear.

EXHIBIT 5

Six jobs AI can take on in new business today

Answer and advise Explains products and answers questions on WhatsApp and web, 24 hours a day, in local languages ASSIST	Read submissions Extracts details from broker emails, schedules, medical and financial reports into the application ASSIST	Draft quotes and proposals Prepares quotes, proposals and policy summaries for an agent or underwriter to check RECOMMEND
Chase missing items Reminds customers and agents about missing documents and follows up until complete ACT WITHIN LIMITS	Summarise for underwriters Turns a complex file into a one-page case summary with the key risks highlighted RECOMMEND	Check before issue Compares the final policy with the application and flags mismatches before it is sent RECOMMEND

Source: KaribuTech AI analysis. Labels show the level of AI autonomy we recommend: assist, recommend or act within limits.

What stays with people

- Every decline and every price loading.
- Advice on complex or high-value cover.
- Any case where the customer seems not to understand the product.
- Complaints and disputes.

Guardrails that make it safe

- **Approved sources only:** AI answers from the insurer's own product documents.
- **No customer data in public AI tools.**
- **Every answer logged** for quality checks and audit.
- **Tested before launch** for accuracy and fairness, and sampled weekly after.

WHERE TO START

Give every agent and underwriter an AI assistant that reads documents and drafts quotes. It is low risk, quick to launch and frees hours every week.

Distribution with AI: every channel, one engine

Most insurers sell through several channels at once, each with its own forms, rules and systems. AI works best when every channel uses the same quote,

onboarding and underwriting engine underneath, so a customer gets the same answer whether they apply through an agent, a bank or a phone.

EXHIBIT 6

What AI does for each channel, and what to watch

CHANNEL	WHAT AI DOES	WHAT CHANGES	WATCH
Agents and brokers	Prepares quotes, compares products, answers questions, drafts proposals	More time advising; fewer errors and resubmissions	Commission on lapsing policies
Bank partners	Offers cover at account opening or loan approval, pre-filled from bank data	Cover sold in the bank's own journey	Data rights; mis-selling
Telcos and mobile money	Simple cover sold and paid through airtime or wallets	Mass-market reach at very low cost	Customer understanding of cover
Platforms and retail	Cover offered at checkout, per trip or per purchase	New customers who never visit an insurer	Revenue share; claims experience
Direct digital	Chat on WhatsApp and web; instant quotes; save and resume	Lowest cost per policy; 24-hour sales	Drop-off; need for human help

Source: KaribuTech AI analysis.

Agreeing terms with partners

1 Data rights

Who owns and can use which data, with customer consent

2 The customer

Whose brand, and who services the policy

3 Revenue share

Linked to persistency and claims quality, not just sales

4 Joint rules

Shared fairness testing, complaints handling and AI governance

Pay partners and agents for policies that stay on the books, not just for policies sold. It is the single most effective protection against mis-selling.

Designing products that sell themselves

The fastest way to cut the cost of new business is to design products that are simple to explain, quick to underwrite and easy to pay for. AI helps both to design

these products, using data on what customers need, and to sell them, by matching the right product to the right customer at the right moment.

EXHIBIT 7

Five product designs that suit AI-led new business

Micro and funeral cover

Small, simple cover with few questions and fast payouts

WHY IT SUITS AI

Straight-through issue; premiums from airtime or wallets

Embedded cover

Cover built into a loan, account, trip or purchase

WHY IT SUITS AI

Offered automatically at the moment of need

Parametric cover

Pays automatically when rainfall, temperature or a flood level crosses a trigger

WHY IT SUITS AI

No loss adjuster; payout in days by mobile money

Pay-as-you-go cover

Premiums by the day, trip, shift or kilometre

WHY IT SUITS AI

Priced from usage data; no annual lump sum

Modular cover

A simple core policy with optional add-ons

WHY IT SUITS AI

Customers build their own cover; AI suggests add-ons

Source: KaribuTech AI analysis.

A product checklist for the AI era

✓ **Five questions or fewer**
to get a price for simple cover

✓ **Explainable in one sentence**
on a phone screen, in local languages

✓ **Payable how customers get paid**
daily, weekly or monthly, by wallet or debit

✓ **Claims paid in days**
with clear evidence rules set up front

✓ **Data the customer agrees to share**
and nothing more

✓ **Priced for persistency**
value that customers keep paying for

Complicated products need people to sell them. Simple products, well designed, can be sold by AI to millions, and bought with confidence.

The 100-day quick-wins plan for new business

Do not start with a large programme. Pick one product and one channel, deliver a working journey in 100 days, measure it against today's baseline, then reuse the

template. Each cycle is run by one small team that includes new business, underwriting, compliance and technology people.

EXHIBIT 8

One journey in 100 days: design, build, go live, measure

DAYS 1-30

Discover and design

- Map today's journey and baseline its numbers
- Agree the new journey and underwriting rules
- Plan data, security and compliance checks

DAYS 31-60

Build

- Digital journey on WhatsApp and web
- Identity checks and document reading connected
- Underwriting triage rules set up

DAYS 61-100

Go live and measure

- Test with real customers and agents
- Go live on one product and one channel
- Report results against the baseline

Source: KaribuTech AI quick-wins approach.

Three quick wins we recommend

QUICK WIN 1

Quote and apply on WhatsApp

- Instant quote for one simple product
- Save and resume across channels
- Agent handover for questions
- First premium by mobile money

MEASURE

Quote time, completion rate, cost per policy

QUICK WIN 2

Automated identity checks

- ID and selfie matching
- Automatic sanctions screening
- Documents read and filed
- Referrals with deadlines

MEASURE

Time to verify, drop-off, rework

QUICK WIN 3

Underwriting triage

- Three lanes with agreed rules
- Simple cases straight through
- Complex cases pre-summarised
- Weekly quality sampling

MEASURE

Time to issue, straight-through rate, referral rate

The business case for AI in new business

The case for AI in new business rests on four numbers: more applications completed, lower cost per policy, faster issue and fewer early lapses. The example below

is illustrative, to show how the numbers add up for a mid-sized insurer; your own baseline will differ.

EXHIBIT 9

Illustrative example: 100,000 applications a year, average annual premium of \$120

Measure	Before	After AI	Effect
Applications completed	55%	65%	+10,000 policies
Cost per policy	\$40	\$28	–\$12 per policy
Time to issue	10 days	1 day	Fewer drop-outs
Lapse in year one	25%	18%	+7 points retained
Premium in force after a year	\$5.0m	\$6.4m	≈ +\$1.4m a year

Source: Illustrative KaribuTech AI example, not a forecast. Policies in force after a year: 41,250 before (55,000 less 25% lapse) and 53,300 after (65,000 less 18% lapse), at \$120 each. Lower cost per policy saves a further ≈ \$0.8m a year.

Where the value comes from

- **More completed applications:** fewer drop-outs at quote, identity checks and underwriting.
- **Lower cost per policy:** less manual work, fewer resubmissions, cheaper channels.
- **Faster issue:** customers who get a yes today do not shop elsewhere tomorrow.
- **Fewer early lapses:** better fit at sale and first premium collected on payday.

Where the money goes

Journey design and change	30–40%
Integration with core systems	20–25%
Identity and document tools	10–15%
Cloud and AI services	10–15%
Training and governance	10%

Typical split of a first-year budget (illustrative)

THREE RULES FOR A CREDIBLE CASE

- 1 Baseline every funnel stage before you start
- 2 Count only benefits the sales and underwriting heads sign
- 3 Report value and cost to the executive team every quarter


What you need underneath: the new business engine

The head of new business does not need to be a technologist, but should know what the engine looks like, because it decides how fast new products and

channels can launch. The key idea is one shared engine behind every channel, rather than a separate system for each.


EXHIBIT 10

One engine behind every channel



Source: KaribuTech AI analysis.

Build, buy or partner



Ask one question of every technology decision: will the next product or channel be faster and cheaper to launch because of it?

The new business team in the AI era

AI does not remove the need for people in new business; it changes what they do. The work moves from keying, checking and chasing towards advising

customers, judging complex risks and improving the journey. The head of new business leads that shift.

EXHIBIT 11

How the main roles change

ROLE	TODAY	WITH AI
Sales agent	Fills forms and chases documents	→ Advises customers, with an AI assistant preparing quotes
Underwriting assistant	Keys data from applications and reports	→ Checks AI-extracted facts and handles exceptions
Underwriter	Reviews every application manually	→ Sets the rules, decides complex cases, samples quality
New business administrator	Issues documents and posts contracts	→ Manages the digital journey and customer queries
Product manager	Designs products on paper, launches yearly	→ Tests offers and journeys with data, improves monthly
Sales manager	Tracks volumes sold	→ Tracks the full funnel, persistency and conduct

Source: KaribuTech AI analysis.

Who you work with

- **Actuaries:** set and sign off pricing and underwriting rules.
- **Compliance:** agree identity, consent and fair-treatment checks.
- **Data and AI team:** builds and monitors the models.
- **Claims:** feeds back which new business turns into bad claims.
- **IT and partners:** connect channels and core systems.

Skills to build in the team

Reading and questioning AI recommendations

Designing customer journeys

Using data to test offers

Explaining decisions to customers

Spotting unfair or unusual outcomes

Risks and guardrails at the point of sale

Selling faster must not mean selling worse. AI can help prevent mis-selling and fraud, but it can also scale

mistakes quickly. These are the risks a head of new business owns, and the controls that manage them.

EXHIBIT 12

Five new business risks, and the control for each

RISK	WHAT CAN GO WRONG	THE CONTROL
Mis-selling	Customers sold cover they do not need or understand	Suitability checks; plain-language summaries; commission on persistency
Application fraud	Fake identities, forged documents and synthetic voices or photos	Liveness checks, document forensics, call-back verification
Unfair pricing or declines	AI treats groups of customers differently without good reason	Fairness testing; explainable decisions; human review of declines
Data and consent	Customer or partner data used without clear permission	Consent recorded; data kept in-country where required; approved tools only
Model errors	Triage or pricing models that get worse over time	Weekly quality sampling; monitoring; the ability to switch a model off

Source: KaribuTech AI analysis; IAIS, Application Paper on the supervision of AI (2025).

Eight questions to ask your team

01 Can we explain every automated decline to the customer?

02 How many applications did we lose last month, and at which step?

03 Which products are most often lapsed or cancelled in year one?

04 Have our triage rules been tested for fairness?

05 How would we spot a fake ID or a synthetic voice?

06 Do we have consent for every piece of data we use?


07 Who reviews the cases the AI refers, and how fast?

08 Are agents paid for policies that stay, or only for sales?

The new business scorecard and 12-month roadmap

EXHIBIT 13


Six numbers to run new business, plus four for the AI era



Source: KaribuTech AI analysis. Green arrow: higher is better; red arrow: lower is better.

EXHIBIT 14

Four quarters from first journey to AI-first new business



Source: KaribuTech AI analysis.

WHERE TO START ON MONDAY

Pull last quarter's numbers for every stage of the funnel, find the single biggest leak, and launch the first 100-day journey on the product and channel where fixing it will matter most.

About the authors



Kennedy Chengeta, PhD

Co-founder, KaribuTech AI

Enterprise Architect and AI Leader · Camunda Chapter Lead, Africa

Leads AI, automation and integration programmes for insurers and banks across Southern Africa.



Justice Muguda

Chief Technology Officer, KaribuTech AI


Cloud, integration and cybersecurity

Leads KaribuTech AI's technology and engineering, from cloud platforms to secure connections between systems.

With contributions from the KaribuTech AI Engineering Team. To discuss this report, or to run a 100-day new business diagnostic, contact KaribuTech AI at kaributech.co.za.

References

1. Accenture, underwriting survey, as reported by Insurance Business.
2. Capgemini, research on underwriters' time and workload.
3. European Union (2024), Regulation (EU) 2024/1689, the Artificial Intelligence Act.
4. GSMA (2026), State of the Industry Report on Mobile Money.
5. IAIS (2025), Application Paper on the supervision of artificial intelligence, July 2025.
6. McKinsey & Company, research on the value of AI in insurance.
7. Microinsurance Network (2024), Landscape of Microinsurance.
8. NAIC (2023), Model Bulletin: Use of Artificial Intelligence Systems by Insurers.
9. Ogborn, M. E. (1962), Equitable Assurances.
10. Histories of industrial assurance in Britain, including Prudential.
11. Safaricom (2007), launch of M-Pesa in Kenya.
12. KaribuTech AI, The AI Blueprint for the Insurance CEO (2026).
13. KaribuTech AI, onboarding and quick-wins approach (proprietary).



About KaribuTech AI

KaribuTech AI is a South African technology company specialising in artificial intelligence, machine learning, intelligent automation and integration for insurers and banks across Africa. Our teams design and build intelligent workflows, decision engines, document and computer-vision AI, and cloud-native core platforms, including the **BluFountain Insure** platform: an insurance AI

marketplace with core systems for life, short-term, medical, funeral and investment business. We combine deep process expertise (Camunda, IBM BPM, MuleSoft, webMethods and document automation platforms) with practical AI engineering, so that AI moves out of the lab and into the operations that customers feel.

Disclaimer: this report provides general information and does not constitute legal, actuarial or investment advice. Third-party figures are cited beneath each exhibit and have not been independently verified by KaribuTech AI. Illustrative figures are not forecasts.



KaribuTech AI

INSURANCE · LEADERSHIP SERIES · SEPTEMBER 2026

The AI Blueprint for the Head of Claims

Faster, fairer claims with less leakage: how AI helps insurers keep their promise


By Kennedy Chengeta, PhD, and Justice Muguda
with the KaribuTech AI Engineering Team

kaributech.co.za



Contents

01	Introduction	86
02	Why claims must change	87
03	Lessons from history	88
04	Where claims leak	89
05	Six moves for the Head of Claims	90
06	Triage with AI: three lanes	92
07	Fighting fraud fairly	93
08	Generative AI and AI agents in claims	94
09	AI in claims, line by line	95
10	Catastrophe and surge response	96
11	The 100-day quick-wins plan	97
12	The business case	98
13	What you need underneath	99
14	The claims team	100
15	Risks and guardrails in claims	101
16	The scorecard and 12-month roadmap	102
17	About the authors and KaribuTech AI	103



About this report

This report is written for heads of claims and their teams at insurers in developing markets, and for the executives they report to.

It shows how AI can make every claim faster, fairer and cheaper to handle, from the first report to payment and recovery: where claims leak, six moves to fix them, triage, assessment and fraud with AI, claims by line of business, catastrophe response, a 100-day plan, the business case, the team, the risks and a scorecard.

External data are sourced beneath each exhibit, and a full list of references appears at the end. Frameworks and illustrative figures without an external source are the authors' own.

WRITTEN FOR

Heads of claims

Claims managers

Fraud teams

Executives

Introduction

A claim is the moment an insurer proves what its promise is worth. It is also where most of the money goes, where fraud and leakage hide, and where customers decide whether to stay.

Too many claims still depend on phone calls, paper forms, repeated requests for documents and weeks of waiting. Claims handlers spend their days keying data and chasing evidence; honest customers wait while suspicious claims slip through; and complaints rise when the process is slow and unclear.

AI changes this. It can take a claim on WhatsApp in minutes, read documents and photos, sort claims by complexity, screen every claim for fraud, pay simple claims the same day and keep customers informed at every step, while people decide the hard cases and every decline.

This report sets out **six moves** for the head of claims, a practical view of **triage, assessment and fraud** with AI, claims by line of business, **catastrophe response**, a **100-day plan**, the business case and a **scorecard**, in plain language for business leaders.

THREE QUESTIONS THIS REPORT ANSWERS

01 Where do our claims leak time and money?

02 Which claims could be paid the same day, safely?

03 How do we fight fraud without delaying honest customers?

AT A GLANCE: SIX MOVES FOR CLAIMS

- 1** Make reporting a claim effortless
- 2** Triage every claim on arrival
- 3** Assess from photos and documents
- 4** Screen every claim for fraud, fairly
- 5** Pay in hours, not weeks
- 6** Close the loop: recover, learn, prevent

Why claims must change

Five pressures are converging on claims teams. Each one is also a chance for the

insurers that move first to win trust and margin.

Customers judge you at the claim

A slow or confusing claim undoes years of marketing. Customers now expect updates like a delivery app.

Leakage hides in manual files

Overpayments from errors, inflated invoices and missed recoveries add up quietly across thousands of claims.

Fraud is getting smarter

Forged documents, staged accidents and now AI-generated photos and voices are harder to spot by eye.

Handling costs are too high

Every phone call, paper form and site visit adds cost, which small policies cannot carry.

Catastrophes come in waves

Floods, storms and droughts create thousands of claims at once, beyond normal capacity.

EXHIBIT 1

What leading insurers have achieved with AI in claims

23 days

faster liability assessment on complex claims at Aviva

65%

fewer customer complaints at Aviva after rewiring claims

80%

faster settlement of simple claims with Allianz's AI agents

5–10%

claims cost savings typical from reducing leakage

Source: McKinsey, Rewired in action: Aviva; Allianz (2025); PwC, Stopping the leaks. PwC also reports leakage of up to 25% in some life portfolios.

In insurance, the claim is the product. Every day saved and every dollar of leakage recovered builds both trust and profit.

Lessons from history: how claims built the industry

Insurers have always been judged by how they handle the worst day in a customer's life. Three moments from the past show why claims are a strategic function, not a

back office, and why the next leap will come from preventing and paying losses faster.

1680s



Insurers who fought the fire

After the Great Fire of London in 1666, early fire insurers such as Nicholas Barbon's Fire Office ran their own fire brigades. Preventing a loss was cheaper than paying for it.

Today: weather alerts and telematics warn customers before a loss.

THE LESSON

The best claim is the one that never happens.

1906



"Pay all our policyholders in full"

After the 1906 San Francisco earthquake, Lloyd's underwriter Cuthbert Heath told his agents to pay policyholders in full, whatever the fine print. The reputation lasted a century.

Today: simple claims paid the same day by mobile money.

THE LESSON

Paying fairly and fast is the strongest brand an insurer can have.

2010



Paid by the index

In northern Kenya, index-based livestock insurance paid herders when satellite data showed pasture had failed, with no one visiting an animal.

Today: parametric cover pays on satellite and weather data.

THE LESSON

Data can replace the loss adjuster for whole classes of claims.

Today



AI at the claim

AI can take a claim by message, assess it from photos, screen it for fraud and pay by mobile money in hours, while people focus on the hard cases.

Today: photo assessment, fraud screening and live updates.

THE LESSON

The fire brigade's prevention and Heath's speed, at scale.

Source: Histories of the Fire Office and early London fire insurance; Lloyd's of London, history of Cuthbert Heath; International Livestock Research Institute (ILRI), Index-Based Livestock Insurance.

The insurers remembered for centuries were the ones that paid quickly and fairly when it mattered. AI makes that possible for every claim, not just the famous ones.

Where claims leak, and how AI plugs the gaps

Every claim passes through the same eight steps. At each one, time, money or trust can leak away. The

exhibit shows the typical leak and what AI does about it. Your own data will show which leaks are biggest.

EXHIBIT 2

The claims journey: eight steps from first report to lessons learned

	WHERE IT LEAKS	WHAT AI DOES
Report Customer tells you about the loss	Long calls; forms; missing details	Report by WhatsApp, voice or app; details captured once
Triage Claim sorted by type and complexity	Every claim queued the same way	Automatic routing; simple, standard or complex
Evidence Documents and photos collected	Documents requested again and again	Upload once on any channel; AI reads and files it
Assessment Damage or loss valued	Site visits for small claims	Photo and document assessment; remote assessors
Fraud checks Claim screened for fraud	Checks only when someone is suspicious	Every claim screened; suspicious ones referred
Decision Claim accepted, adjusted or declined	Unclear reasons; slow approvals	Rules within limits; a person on every decline
Payment Customer or supplier paid	Cheques and bank delays	Mobile money or instant payment, same day
Recovery and learning Recoveries and feedback	Recoveries missed; lessons lost	Recovery tracking; insights sent to underwriting

Source: KaribuTech AI analysis.

Most delay is waiting, not working. Measure how long each claim sits between steps, and fix the longest wait first.

Six moves for the Head of Claims

Claims leaders who use AI well do six things differently. Each is about the customer's

experience and the way the team works, not about choosing a technology.

MOVE 1

01

Make reporting a claim effortless

The first notice of loss sets the tone for the whole claim. Customers should be able to report in minutes, on the channel they already use, and give every detail only once.

What leaders do

- Offer claim reporting on WhatsApp, voice, USSD and app, not only by phone.
- Let AI ask the right questions for each type of claim and capture photos on the spot.
- Confirm receipt instantly, with a reference and the next step.

ASK YOUR TEAM

How long does it take a customer to report a claim today, and how many times do they repeat themselves?

MOVE 2

02

Triage every claim on arrival

Not every claim needs the same handling. AI can sort claims the moment they arrive by type, value and complexity, and send each to the right lane and the right person.

What leaders do

- Set three lanes: fast-track, standard and complex.
- Route specialist claims, such as injury or liability, to experts immediately.
- Review triage rules monthly with claims and actuarial teams.

ASK YOUR TEAM

What share of our claims could be paid on a fast track, and why are they not today?

MOVE 3

03

Assess from photos and documents

AI can read invoices, medical reports and police reports, and estimate damage from photos and videos. Assessors spend their time on complex claims rather than driving to small ones.

What leaders do

- Photo-based damage estimates for motor and property claims.
- Document reading for medical, funeral and life claims.
- Remote video assessment before any site visit.

ASK YOUR TEAM

How many site visits did we make last month for claims under our fast-track limit?

MOVE 4

04

Screen every claim for fraud, fairly

Manual fraud checks catch only the claims someone already suspects. AI can screen every claim for unusual patterns and links between people, garages, doctors and addresses, and send only the suspicious ones to investigators.

What leaders do

- Score every claim; refer only the highest-risk ones to the fraud team.
- Check photos and documents for signs of editing or AI generation.
- Measure false alarms, so honest customers are not delayed.

ASK YOUR TEAM

How many honest customers did our fraud checks delay last month?

MOVE 5

Pay in hours, not weeks

For simple, clean claims, the fastest safe decision is an automatic one within agreed limits, paid straight to the customer's mobile wallet or account. Speed is the strongest proof of the promise.

What leaders do

- Pay simple claims straight through within limits the claims team sets.
- Pay by mobile money or instant payment, not cheque.
- Keep a person on every decline and every payment above the limit.

ASK YOUR TEAM

What is the smallest claim we still take more than a week to pay?

05

MOVE 6

Close the loop: recover, learn, prevent

A claim is also data. AI can spot recovery opportunities, feed patterns back to underwriting and pricing, and warn customers before losses happen, such as flood alerts or maintenance reminders.

What leaders do

- Track every recovery from third parties and reinsurers until collected.
- Send claims insights to underwriting and product teams every month.
- Use claims data to warn customers before losses: weather, theft or health alerts.

ASK YOUR TEAM

Which three lessons from last quarter's claims have we fed back into pricing and products?

06

FROM OUR WORK

Claims adjudication and pre-authorisation, automated

The KaribuTech AI engineering team, working with client teams, has built claims adjudication and pre-authorisation workflows for health insurers, and a digital health platform with real-time insurer pre-authorisation. Rules decide the routine cases instantly, documents are read automatically, and clinical and claims experts see only the cases that need their judgement.

Real time

pre-authorisation at the point of care

Rules first

routine claims decided within agreed limits

Experts on exceptions

people decide the complex cases


Triage with AI: three lanes, one set of rules

The quickest way to cut cycle time without adding risk is to stop handling every claim the same way. AI sorts claims into three lanes the moment they arrive. The

claims team sets the rules for each lane, and every decline or exception still goes to a person.

EXHIBIT 3

Three claims lanes: let AI speed up the simple claims so experts can focus on the hard ones



MORE SPEED →

Source: KaribuTech AI analysis.

← MORE PERSONAL JUDGEMENT

What decides the lane

- **Value:** below or above the fast-track limit for that product.
- **Evidence:** complete, consistent documents and photos.
- **Fraud score:** low, medium or high risk.
- **Type:** injury, liability and disputes always go to experts.

Guardrails the claims team owns

- **Limits set by people:** claims and actuarial leaders agree the rules.
- **Weekly sampling:** fast-track claims checked for accuracy.
- **Fairness checks:** outcomes compared across customer groups.
- **Human decline:** no claim is declined by AI alone.


Fighting fraud with AI, without punishing honest customers

Fraud is a cost every honest customer pays through higher premiums. AI can screen every claim, see links that no person could, and spot forged or AI-generated

evidence. The goal is to catch more fraud while delaying fewer honest customers.

EXHIBIT 4

Screen everything, refer a few, investigate fewer, decide fairly



Source: KaribuTech AI analysis; widths are illustrative.

What AI can spot

Unusual patterns

Claims soon after a policy starts, repeated claims, amounts just under limits

Hidden links

The same phone, address, garage, doctor or bank account across many claims

Edited evidence

Altered invoices, reused photos, and AI-generated images, documents or voices

Inflated costs

Repair, medical or funeral costs far above the usual for that area and type

KEEP FRAUD CHECKS FAIR

- Measure false alarms as closely as fraud caught
- Never decline a claim on an AI score alone
- Tell customers what happens next if their claim is checked
- Test fraud models for bias across customer groups

Generative AI and AI agents in claims

Generative AI reads and writes; AI agents plan and carry out tasks. In claims, that means taking the first report in conversation, summarising long files, drafting

letters and keeping customers informed. Give AI more freedom only where the risk is low and the rules are clear.

EXHIBIT 5

Six jobs AI can take on in claims today

Take the first report Guides the customer through reporting on WhatsApp or voice, asking the right questions for each claim type ACT WITHIN LIMITS	Read the evidence Extracts details from invoices, medical, police and assessor reports into the claim file ASSIST	Summarise the file Turns a long claim file into a one-page summary with timeline and open questions RECOMMEND
Keep customers informed Sends clear updates at every step and answers "where is my claim?" any time ACT WITHIN LIMITS	Draft letters and decisions Prepares settlement, request and decline letters for a handler to approve RECOMMEND	Chase suppliers Follows up garages, hospitals and assessors for missing quotes and reports ACT WITHIN LIMITS

Source: KaribuTech AI analysis. Labels show the level of AI autonomy we recommend.

What stays with people

- Every decline and every partial payment.
- Injury, bereavement and other sensitive claims.
- Disputes and complaints.
- Any claim where the customer is vulnerable.

Guardrails that make it safe

- **Approved sources only:** AI answers from the policy and claim file.
- **No customer data in public AI tools.**
- **Every message logged** for quality and audit.
- **Tested and sampled** for accuracy and tone.

WHERE TO START

Start with claim summaries and customer updates. They save handlers hours every week, and customers notice the difference straight away.

AI in claims, line by line

The same AI building blocks work across every line of business, but the evidence, the risks and the quick wins

differ. The table shows where AI helps most in each line common in our markets.

EXHIBIT 6

Where AI helps most, by line of business

LINE	WHAT AI DOES	FRAUD TO WATCH	QUICK WIN
Motor	Photo damage estimates; garage quotes checked against usual costs; accident details from trip data	Staged accidents; inflated repairs	Photo estimates for small claims
Property and home	Photo and video assessment; satellite and weather data after storms and floods	Pre-existing damage; inflated contents	Remote assessment for small losses
Health and medical	Pre-authorisation in real time; invoices and medical codes checked automatically	Over-servicing; billing for care not given	Automated checking of routine claims
Life and funeral	Death verification from registries; documents read; beneficiaries paid by mobile money	Fake deaths; forged certificates	Funeral claims paid within 48 hours
Agriculture	Satellite and weather data; index-based payouts without field visits	Basis risk; disputed yields	Automatic payouts on weather triggers
Commercial and liability	Long reports and contracts summarised; timelines built; reserves suggested for review	Complex disputes; late reporting	AI summaries for handlers

Source: KaribuTech AI analysis.

Connecting the claims supply chain

Garages and panel beaters Quotes and photos sent digitally; AI checks costs against usual repair prices	Hospitals and doctors Pre-authorisation and invoices checked in real time against cover and usual tariffs	Funeral providers Claims lodged by the provider with documents; payment released within hours	Assessors and investigators Work allocated by portal; reports read by AI and added to the file
---	---	---	--

Start with the line where you have the most small, similar claims. That is where AI pays back first, and where customers feel the difference soonest.


When catastrophe strikes: AI for surge response

Floods, storms, droughts and fires can create more claims in a week than a claims team sees in a year. AI lets insurers estimate exposure before claims arrive,

assess damage at scale from satellite and drone images, and pay the most urgent claims first.


EXHIBIT 7

The first 30 days after a catastrophe, with AI



Source: KaribuTech AI analysis.

Three tools that change catastrophe claims



The insurer that pays quickly after a disaster earns loyalty that no marketing campaign can buy. Plan your surge response before the next storm, not during it.

The 100-day quick-wins plan for claims

Do not start with a large programme. Pick one claim type and one channel, deliver a working journey in 100 days, measure it against today's baseline, then reuse

the template. Each cycle is run by one small team that includes claims, fraud, compliance and technology people.

EXHIBIT 8

One claims journey in 100 days: design, build, go live, measure

DAYS 1-30

Discover and design

- Map today's journey and baseline its numbers
- Agree lanes, limits and evidence rules
- Plan data, security and compliance checks

DAYS 31-60

Build

- Claim reporting on WhatsApp and web
- Document and photo reading connected
- Triage and fraud screening set up

DAYS 61-100

Go live and measure

- Test with real claims and handlers
- Go live on one claim type and channel
- Report results against the baseline

Source: KaribuTech AI quick-wins approach.

Three quick wins we recommend

QUICK WIN 1

Digital claim reporting

- Report on WhatsApp, USSD or app
- Photos and documents on the spot
- Instant reference and next steps
- Status updates at every step

MEASURE

Time to report, calls avoided, customer satisfaction

QUICK WIN 2

Photo and document assessment

- Motor or property photos assessed
- Invoices and reports read
- Fast-track lane for small claims
- Weekly quality sampling

MEASURE

Days to pay, site visits avoided, accuracy

QUICK WIN 3

Fraud screening and triage

- Every claim scored on arrival
- Links across claims highlighted
- Referrals with deadlines
- False alarms tracked

MEASURE

Fraud found, false-alarm rate, leakage

The business case for AI in claims

AI in claims pays back through four levers: less leakage, lower handling cost, more fraud caught and fewer complaints. The example below is illustrative, to show

how the numbers add up for a mid-sized insurer; your own baseline will differ.

EXHIBIT 9

Illustrative example: 50,000 claims a year, \$75m of claims paid

Measure	Before	After AI	Annual value
Claims leakage	8% of claims paid	5% of claims paid	≈ \$2.25m
Handling cost per claim	\$120	\$85	≈ \$1.75m
Fraud detected	0.5% of claims paid	1.0% of claims paid	≈ \$0.38m
Days to pay (simple claims)	21 days	2 days	Fewer complaints
Total annual value at scale			≈ \$4.4m

Source: Illustrative KaribuTech AI example, not a forecast. Leakage: 3% of \$75m. Handling: \$35 saved on 50,000 claims. Fraud: an extra 0.5% of \$75m identified before payment. Leakage range informed by PwC, Stopping the leaks.

Where the value comes from

- **Less leakage:** every claim checked against policy terms and usual costs.
- **Lower handling cost:** fewer calls, forms, site visits and manual checks.
- **More fraud caught:** every claim screened, not just the suspicious ones.
- **Fewer complaints:** faster payment and clear updates.

Where the money goes

Journey design and change	30–40%
Integration with core systems	20–25%
Document, photo and fraud tools	15–20%
Cloud and AI services	10–15%
Training and governance	10%

Typical split of a first-year budget (illustrative)

THREE RULES FOR A CREDIBLE CASE

- 1 Baseline leakage, cost and cycle time before you start
- 2 Count only benefits the claims and finance heads sign
- 3 Report value, complaints and fairness every quarter


What you need underneath: the claims engine

The head of claims does not need to be a technologist, but should know what the engine looks like, because it decides how fast claims can be paid and how easily

new partners can plug in. The key idea is one shared claims engine behind every channel and every partner.

EXHIBIT 10

One claims engine behind every channel and partner



Source: KaribuTech AI analysis.

Build, buy or partner

Buy

Document reading, photo estimation, identity checks, chat tools and payment connections

Build

Your lanes, limits, claims rules and fraud models, trained on your own claims data

Partner

Garage, hospital and assessor networks; shared industry fraud data

Ask one question of every technology decision: will the next claim type, partner or catastrophe be faster to handle because of it?

The claims team in the AI era

AI does not remove the need for people in claims; it changes what they do. Work moves from keying, chasing and checking towards caring for customers,

judging complex claims and investigating real fraud. The head of claims leads that shift.

EXHIBIT 11

How the main roles change

ROLE	TODAY	WITH AI
Claims clerk	Keys data and files documents	→ Checks AI-prepared files and handles exceptions
Claims handler	Works every claim from start to finish	→ Focuses on standard and complex claims; approves AI drafts
Assessor	Drives to inspect most claims	→ Assesses remotely from photos and video; visits only complex losses
Fraud investigator	Reviews claims someone suspects	→ Investigates the few high-risk claims AI refers
Call-centre agent	Answers "where is my claim?"	→ Supports vulnerable customers and complex queries
Claims manager	Tracks backlogs and volumes	→ Tracks cycle time, leakage, fairness and customer outcomes

Source: KaribuTech AI analysis.

Who you work with

- **Underwriting and product:** receive claims insights every month.
- **Actuaries:** set reserves and agree fast-track limits.
- **Fraud and compliance:** agree referral rules and fair treatment.
- **Reinsurance:** track recoveries and catastrophe claims.
- **IT and partners:** connect garages, hospitals and payments.

Skills to build in the team

Reading and questioning AI recommendations

Empathy with customers in distress

Remote assessment

Explaining decisions clearly

Spotting unfair or unusual outcomes

Risks and guardrails in claims

Paying faster must not mean paying wrongly, or declining unfairly. AI can reduce both errors and fraud, but it can also repeat a mistake thousands of times.

These are the risks a head of claims owns, and the controls that manage them.

EXHIBIT 12

Five claims risks, and the control for each

RISK	WHAT CAN GO WRONG	THE CONTROL
Unfair declines	Valid claims declined, or some customer groups treated worse	A person on every decline; clear reasons; fairness testing
False fraud alarms	Honest customers delayed or accused	Track false alarms; never decline on a score alone; quick review
Fake evidence	Forged documents and AI-generated photos, videos or voices	Evidence forensics; liveness checks; call-back verification
Sensitive data	Health and bereavement data misused or leaked	Consent recorded; access controls; in-country data where required
Model errors	Estimates or triage that get worse over time	Weekly sampling; monitoring; the ability to switch a model off

Source: KaribuTech AI analysis; IAIS, Application Paper on the supervision of AI (2025).

Eight questions to ask your team

01

Can we explain every declined claim to the customer?

02

How long do simple claims wait between steps?

03

How many honest customers did fraud checks delay?

04

Have our triage and fraud rules been tested for fairness?

05

How would we spot an AI-generated photo or voice?

06

Who reviews the claims AI refers, and how fast?

07

Which recoveries did we miss last quarter?


08

Is our surge plan ready for the next flood or storm?

The claims scorecard and 12-month roadmap

EXHIBIT 13


Six numbers to run claims, plus four for the AI era



Source: KaribuTech AI analysis. Green arrow: higher is better; red arrow: lower is better.

EXHIBIT 14

Four quarters from first journey to AI-first claims



Source: KaribuTech AI analysis.

WHERE TO START ON MONDAY

Pull last quarter's claims data, find the claim type with the most small, similar claims and the longest waits, and launch the first 100-day journey there.

About the authors



Kennedy Chengeta, PhD

Co-founder, KaribuTech AI

Enterprise Architect and AI Leader · Camunda Chapter Lead, Africa

Leads AI, automation and integration programmes for insurers and banks across Southern Africa.



Justice Muguda

Chief Technology Officer, KaribuTech AI


Cloud, integration and cybersecurity

Leads KaribuTech AI's technology and engineering, from cloud platforms to secure connections between systems.

With contributions from the KaribuTech AI Engineering Team. To discuss this report, or to run a 100-day claims diagnostic, contact KaribuTech AI at kaributech.co.za.

References

1. Allianz (2025), announcements on AI agents in claims handling.
2. European Union (2024), Regulation (EU) 2024/1689, the Artificial Intelligence Act.
3. IAIS (2025), Application Paper on the supervision of artificial intelligence, July 2025.
4. IBM (2025), Cost of a Data Breach Report 2025.
5. International Livestock Research Institute (ILRI), Index-Based Livestock Insurance, Kenya (from 2010).
6. Lloyd's of London, history of Cuthbert Heath and the 1906 San Francisco earthquake.
7. McKinsey & Company, Rewired in action: Aviva.
8. NAIC (2023), Model Bulletin: Use of Artificial Intelligence Systems by Insurers.
9. PwC, Stopping the leaks: claims leakage in insurance.
10. Histories of the Fire Office and early London fire insurance (1680s).
11. KaribuTech AI, The AI Blueprint for the Insurance CEO (2026).
12. KaribuTech AI, claims and quick-wins approach (proprietary).



About KaribuTech AI


KaribuTech AI is a South African technology company specialising in artificial intelligence, machine learning, intelligent automation and integration for insurers and banks across Africa. Our teams design and build intelligent workflows, decision engines, document and computer-vision AI, and cloud-native core platforms, including the **BluFountain Insure** platform: an insurance AI

marketplace with core systems for life, short-term, medical, funeral and investment business. We combine deep process expertise (Camunda, IBM BPM, MuleSoft, webMethods and document automation platforms) with practical AI engineering, so that AI moves out of the lab and into the operations that customers feel.

Disclaimer: this report provides general information and does not constitute legal, actuarial or investment advice. Third-party figures are cited beneath each exhibit and have not been independently verified by KaribuTech AI. Illustrative figures are not forecasts.



KaribuTech AI



INSURANCE · LEADERSHIP SERIES · SEPTEMBER 2026



The AI Blueprint for the Head of Medical Claims

Faster pre-authorisation, fairer adjudication and less fraud, waste and abuse: AI for medical insurance claims

By Kennedy Chengeta, PhD, and Justice Muguda
with the KaribuTech AI Engineering Team

kaributech.co.za





About this report

This report is written for heads of medical claims, medical scheme administrators, clinical advisers and health insurers in developing markets.

It shows how AI can speed up pre-authorisation, adjudicate routine claims automatically, check coding and tariffs, fight fraud, waste and abuse, and work better with providers, while clinicians decide the hard cases: the pressures, six moves, adjudication lanes, provider partnerships, new models of care, a 100-day plan, the business case, the team, the risks and a scorecard.

External data are sourced beneath each exhibit, and a full list of references appears at the end. Frameworks and illustrative figures without an external source are the authors' own.

WRITTEN FOR

Heads of medical claims

Clinical advisers

Scheme administrators

Fraud teams

Contents

01	Introduction	106
02	Why medical claims must change	107
03	Lessons from history	108
04	Where medical claims leak	109
05	Six moves for the Head of Medical Claims	110
06	Three adjudication lanes	112
07	Fraud, waste and abuse	113
08	Generative AI and AI agents in medical claims	114
09	Working with providers	115
10	New models of care	116
11	The 100-day quick-wins plan	117
12	The business case	118
13	What you need underneath	119
14	The medical claims team	120
15	Risks and guardrails	121
16	The scorecard and 12-month roadmap	122
17	About the authors and KaribuTech AI	123

Introduction

A medical claim is not just a payment. Behind it is a patient waiting for treatment, a doctor waiting for approval and a hospital waiting to be paid.

Medical claims teams handle huge volumes of small, similar claims from doctors, pharmacies, labs and hospitals, alongside complex, expensive cases. Pre-authorisations wait in queues, invoices are checked by hand, coding errors slip through, and fraud, waste and abuse quietly raise costs for every member.

AI can check eligibility and benefits instantly, approve routine pre-authorisations in real time, read invoices and medical reports, check codes and tariffs, and flag unusual billing, while clinical advisers decide anything that affects a patient's care.

This report sets out **six moves** for the head of medical claims, **adjudication lanes, fraud, waste and abuse**, working with **providers, new models of care** such as telemedicine and tele-robotic surgery, a **100-day plan** and a **scorecard**.

THREE QUESTIONS THIS REPORT ANSWERS

01 Which pre-authorisations could be approved instantly?

02 How much of our claims volume could be decided automatically?

03 How do we cut fraud, waste and abuse fairly?

AT A GLANCE: SIX MOVES FOR MEDICAL CLAIMS

- 1** Pre-authorise in real time
- 2** Adjudicate routine claims automatically
- 3** Check coding and tariffs on every claim
- 4** Fight fraud, waste and abuse fairly
- 5** Manage provider networks with data
- 6** Prevent claims through care management

Why medical claims must change

Five pressures are squeezing medical schemes and health insurers. Each is also a

chance to lower costs and improve care at the same time.

Medical costs keep rising

New treatments, chronic disease and higher provider prices push claims up faster than premiums.

Pre-authorisation delays care

Patients and doctors wait for approvals that are often routine.

Fraud, waste and abuse

Over-servicing, upcoding and billing for care not given raise costs for every member.

Paper and manual checks

Invoices, codes and reports are still checked by hand, one by one.

Health data is sensitive

Medical records need the strongest privacy and security, and breaches are costly.

EXHIBIT 1

Evidence from claims automation

80%

faster settlement of simple claims with Allianz's AI agents

5–10%

claims cost savings typical from reducing leakage

\$4.44m

global average cost of a data breach (2025)

Real time

pre-authorisation at the point of care is now possible

Source: Allianz (2025); PwC, Stopping the leaks; IBM, Cost of a Data Breach Report 2025; KaribuTech AI project experience.

Every hour a routine pre-authorisation waits is an hour a patient waits. Automation is not only cheaper; it is better care.

Lessons from history: how health cover was built

Modern health insurance grew from a few breakthroughs in how care was funded, organised and

recorded. Each one shows what AI can now do for medical claims.

1883

Sickness insurance for workers

Germany's Health Insurance Act of 1883, under Otto von Bismarck, required sickness funds for many workers, creating one of the first national health insurance systems.

Today: automatic eligibility and instant benefit checks.

THE LESSON

Cover that is simple and automatic reaches the most people.

1929

The hospital prepayment plan

At Baylor University Hospital in Dallas, schoolteachers paid a small monthly amount for guaranteed hospital care. The idea grew into the Blue Cross plans.

Today: real-time links with provider networks.

THE LESSON

Partnering directly with providers makes care affordable and predictable.

1990

A common language for diagnoses

The World Health Assembly endorsed ICD-10, the international classification of diseases used to code diagnoses on claims worldwide.

Today: AI checks every code against diagnosis and rules.

THE LESSON

Common coding makes claims comparable and checkable.

Today

AI at the point of care

AI can check eligibility, approve routine pre-authorisations and adjudicate claims in seconds, using coded data and clinical rules.

Today: routine approvals in minutes; clinicians on the complex cases.

THE LESSON

Clinicians focus on the cases that need their judgement.

Source: Histories of German sickness insurance (1883) and the Baylor plan (1929); World Health Organization, ICD-10 (1990).

Each step made health cover easier to fund and to manage. AI is the next step: making every routine decision instant, so clinicians can focus on patients.

Where medical claims leak, and how AI plugs the gaps

Every medical claim passes through the same eight steps, from checking cover to auditing payment. At

each one, time, money or trust can leak away. The exhibit shows the typical leak and what AI does.

EXHIBIT 2

The medical claims journey: eight steps from eligibility to audit

	WHERE IT LEAKS	WHAT AI DOES
Eligibility Is the member covered?	Calls to check benefits	Instant checks at the point of care
Pre-authorisation Approval before treatment	Queues for routine requests	Routine requests approved in real time
Treatment Care is delivered	Care beyond what was approved	Alerts when care exceeds the plan
Submission Provider sends the claim	Paper invoices; missing details	Electronic claims read automatically
Coding check Diagnosis and procedure codes	Upcoding and unbundling missed	Codes checked against diagnosis and rules
Adjudication Claim approved or adjusted	Every claim checked by hand	Routine claims decided automatically
Payment Provider or member paid	Slow payments strain providers	Fast payment; clear remittances
Audit and recovery Checks after payment	Recoveries missed	Patterns found; overpayments recovered

Source: KaribuTech AI analysis.

Most medical claims are routine. Decide them instantly, and give clinicians time for the ones that matter.

Six moves for the Head of Medical Claims

Medical claims leaders who use AI well do six things differently. Each is about faster, fairer

decisions and better care, not about choosing a technology.

MOVE 1

01

Pre-authorise in real time

Many pre-authorisation requests are routine and meet clear clinical rules. AI can check the member's cover, the diagnosis and the requested treatment against agreed guidelines and approve routine requests instantly, sending the rest to a clinical adviser.

What leaders do

- Approve routine requests in real time against agreed clinical rules.
- Send complex or high-cost requests to a clinical adviser with the file summarised.
- Tell the provider and member the outcome immediately.

ASK YOUR TEAM

How long does a routine pre-authorisation wait today, and why?

MOVE 2

02

Adjudicate routine claims automatically

GP visits, pharmacy scripts and lab tests make up most claims by number. With clean electronic claims and clear rules, most can be checked and paid without a person touching them.

What leaders do

- Receive claims electronically from as many providers as possible.
- Check eligibility, benefits, codes and tariffs automatically.
- Send only exceptions to assessors, with the reason flagged.

ASK YOUR TEAM

What share of our claims could be paid today without a person touching them?

MOVE 3

03

Check coding and tariffs on every claim

Coding errors and tariff mistakes are a large, quiet source of leakage. AI can check that procedure codes match the diagnosis, that items are not billed separately when they should be combined, and that prices match agreed tariffs.

What leaders do

- Check every claim's codes against its diagnosis and clinical rules.
- Check prices against agreed tariffs and flag outliers.
- Give providers clear feedback so errors fall over time.

ASK YOUR TEAM

How much did we pay last year above agreed tariffs?

MOVE 4

04

Fight fraud, waste and abuse fairly

Fraud, waste and abuse range from honest mistakes to deliberate billing for care never given. AI can compare each provider with its peers, spot unusual patterns and send the riskiest cases to investigators, without delaying honest providers.

What leaders do

- Compare providers with peers on cost, volume and patterns of care.
- Screen every claim; refer only the highest-risk ones.
- Measure false alarms, so honest providers are paid on time.

ASK YOUR TEAM

Which providers bill very differently from their peers, and do we know why?

MOVE 5

05

Manage provider networks with data

Providers are partners, not adversaries. Data on quality, cost and outcomes helps schemes build networks, agree fair tariffs and reward providers who deliver good care efficiently.

What leaders do

- Score providers on cost, quality and member experience.
- Use data in tariff negotiations and network design.
- Give providers dashboards showing how they compare.

ASK YOUR TEAM

Which of our providers deliver the best outcomes for the cost?

MOVE 6

06

Prevent claims through care management

The best medical claim is one that is prevented. AI can identify members with rising risk, such as uncontrolled diabetes or high blood pressure, and connect them with care programmes before they need hospital treatment.

What leaders do

- Identify members at rising risk from claims and, with consent, health data.
- Offer care programmes, reminders and telemedicine.
- Measure hospital admissions avoided, not just claims paid.

ASK YOUR TEAM

How many of last year's hospital admissions could have been prevented?

FROM OUR WORK

Real-time pre-authorisation and adjudication

The KaribuTech AI engineering team, working with client teams, has built claims adjudication and pre-authorisation workflows for health insurers, and a digital health platform with real-time insurer pre-authorisation. Clinical rules decide routine requests instantly, documents are read automatically, and clinical advisers see only the cases that need their judgement. BluFountain Insure's medical core, Medishield, and its HL7 compliance support this kind of connected health data.

Real time

pre-authorisation at the point of care

Rules first

routine claims decided within agreed limits

HL7

standard health data exchange


Three adjudication lanes: let AI decide the routine, clinicians the complex

Medical claims range from a routine pharmacy script to a complex cancer treatment. AI sorts every claim and pre-authorisation into three lanes. Clinical advisers and

the claims team set the rules, and every decline goes to a person.

EXHIBIT 3

Three lanes for medical claims and pre-authorisations



MORE AUTOMATION →

← MORE CLINICAL JUDGEMENT

Source: KaribuTech AI analysis.

What decides the lane

- **Cost:** below or above the automatic limit.
- **Clinical complexity:** routine or specialist care.
- **Coding and tariff checks:** clean or flagged.
- **Fraud risk score:** low, medium or high.

Guardrails the team owns

- **Clinical rules agreed** with clinical advisers.
- **No decline by AI alone.**
- **Urgent care never delayed** by automated checks.
- **Fairness checks** across member groups.

Fraud, waste and abuse: what AI can spot

Fraud, waste and abuse cost every member through higher contributions. Some is deliberate; much is error or poor practice. AI can spot patterns across millions of

claims that no assessor could see, so investigators can focus on the cases that matter.

EXHIBIT 4

Six common patterns, and how AI detects them

PATTERN	WHAT IT LOOKS LIKE	HOW AI DETECTS IT
Upcoding	Billing a more expensive procedure than was performed	Codes compared with diagnosis and peer patterns
Unbundling	Billing separately for items that should be one charge	Clinical coding rules applied to every claim
Services not rendered	Billing for care that was never given	Claims matched to visits, pre-authorisations and members
Duplicate claims	The same service billed twice or by two providers	Automatic matching across all claims
Over-servicing	More tests or visits than the condition needs	Provider behaviour compared with peers
Member fraud	Cards shared with non-members; false claims	Identity checks; unusual usage patterns

Source: KaribuTech AI analysis.

Keep it fair

01

Educate first

Many errors stop when providers get clear feedback.

02

Refer the few

Only the highest-risk cases go to investigators.

03

Protect patients

Never delay urgent care because of a fraud flag.


Generative AI and AI agents in medical claims

Generative AI reads and writes; AI agents plan and carry out tasks. In medical claims, that means reading medical reports and invoices, summarising cases for

clinical advisers, answering members' benefit questions and drafting letters, while clinicians make every clinical decision.

EXHIBIT 5

Six jobs AI can take on in medical claims today



Source: KaribuTech AI analysis. Labels show the level of AI autonomy we recommend.

What stays with clinicians and people

- Every clinical decision about a patient's care.
- Every decline and partial payment.
- Urgent and emergency cases.
- Disputes with providers and members.

Guardrails that make it safe

- **Medical data protected:** secure, approved tools only.
- **No patient data in public AI tools.**
- **Every AI output logged** for audit.
- **Clinical rules** reviewed by clinical advisers.

WHERE TO START

Start with case summaries for clinical advisers and benefit answers for members. Both save time without touching clinical decisions.

Working with providers: data that helps both sides

Providers want to be paid quickly and fairly; schemes want good care at a fair price. Shared data and clear

rules reduce disputes and speed up payment. The table shows what AI can do for each type of provider.

EXHIBIT 6

What AI does for each type of provider

PROVIDER	WHAT AI DOES	WHAT BOTH SIDES GAIN
Hospitals	Real-time pre-authorisation; case summaries; tariff checks	Faster approvals; fewer disputes
GPs and clinics	Instant eligibility; automatic adjudication	Same-day payment for routine visits
Pharmacies	Script checks against benefits and chronic lists	Fewer rejected scripts; safer medicines
Laboratories and radiology	Tests matched to diagnosis and guidelines	Fewer unnecessary tests
Specialists	Referral and treatment plans checked against rules	Clearer approvals for complex care

Source: KaribuTech AI analysis.

Three rules for provider partnerships

01

Pay fast

Clean claims paid within days, not weeks.

02

Share the data

Give providers dashboards on their own claims.

03

Reward quality

Better outcomes earn better terms.

New models of care: what they mean for claims

Care is moving out of the consulting room: to phones, homes and even across hundreds of kilometres by robot. In September 2026, a surgeon in Ogun State remotely removed a cancerous kidney from a patient in

Abuja, about 500 km away, in what was described as West Africa's first tele-robotic surgery. Claims rules must keep up.

EXHIBIT 7

Four new models of care, and what changes for medical claims

MODEL	WHAT CHANGES	WHAT CLAIMS TEAMS NEED
Telemedicine	Consultations by video, phone or chat	Benefits and tariffs for remote visits; identity checks
Remote monitoring	Devices track chronic conditions at home	Cover for devices; data consent; care-programme rules
Tele-robotic surgery	A surgeon operates remotely through a robot and a live link	Rules for two sites, two teams and liability; connectivity risk
Home and pharmacy delivery	Medicines and care delivered to the home	Delivery claims checked; chronic medicines managed

Source: Tribune Online; Pulse Nigeria; Leadership (September 2026), reporting on Redeemer's Health Village, RoboMed Global and Nisa Premier Hospital; KaribuTech AI analysis.

What to do now

01

Update benefit rules

Cover new care models clearly, with fair tariffs.

02

Agree liability

Know who is responsible when care spans sites.

03

Use the data

Remote care creates data that can prevent claims.

The 100-day quick-wins plan for medical claims

Do not start with a large programme. Pick one benefit or claim type, deliver a working improvement in 100 days, measure it against today, then reuse the approach.

Each cycle is run by a small team from claims, clinical advice, provider relations and technology.

EXHIBIT 8

One medical claims journey in 100 days: design, build, go live, measure

DAYS 1-30

Discover and design

- Baseline turnaround and costs
- Agree clinical and tariff rules
- Plan data privacy and security

DAYS 31-60

Build

- Electronic claims and eligibility
- Automatic rules and code checks
- Fraud screening set up

DAYS 61-100

Go live and measure

- Go live on one benefit or provider group
- Compare with today's process
- Report results to the executive team

Source: KaribuTech AI quick-wins approach.

Three quick wins we recommend

QUICK WIN 1

Real-time pre-authorisation

- One benefit, such as scans or day surgery
- Routine requests approved instantly
- Complex ones summarised for clinicians
- Providers told immediately

MEASURE

Pre-authorisation time, share approved instantly

QUICK WIN 2

Auto-adjudication

- GP, pharmacy and lab claims
- Eligibility, codes and tariffs checked
- Exceptions to assessors
- Weekly quality sampling

MEASURE

Auto-adjudication rate, days to pay

QUICK WIN 3

Fraud, waste and abuse screening

- Every claim scored
- Providers compared with peers
- Highest-risk cases referred
- Feedback to providers

MEASURE

Savings found, false-alarm rate

The business case for AI in medical claims

AI in medical claims pays back through less fraud, waste and abuse, lower administration costs and faster turnaround that providers and members notice. The

example below is illustrative, for a mid-sized health insurer or medical scheme.

EXHIBIT 9

Illustrative example: 400,000 claims a year, \$50m of claims paid

Measure	Before	After AI	Annual value
Fraud, waste and abuse paid	6% of claims	4% of claims	≈ \$1.0m
Administration cost per claim	\$6.00	\$3.50	≈ \$1.0m
Routine pre-authorisation time	48 hours	Minutes	Faster care
Days to pay providers	14 days	3 days	Fewer disputes
Total annual value			≈ \$2.0m + better care

Source: Illustrative KaribuTech AI example, not a forecast. Fraud, waste and abuse: 2% of \$50m. Administration: \$2.50 saved on 400,000 claims.

Where the value comes from

- **Less fraud, waste and abuse:** every claim checked, not a sample.
- **Lower costs:** fewer manual checks and calls.
- **Faster care:** routine approvals in minutes.
- **Happier providers:** paid faster, with fewer disputes.

Where the money goes

Electronic claims and provider links	25–35%
Rules, coding and tariff engines	20–25%
Fraud analytics	15–20%
Privacy and security	10–15%
Training and clinical governance	10%

Typical split of a first-year budget (illustrative)

THREE RULES FOR A CREDIBLE CASE

- 1 Baseline turnaround, costs and leakage first
- 2 Agree clinical rules with clinical advisers
- 3 Report savings, turnaround and complaints quarterly


What you need underneath: the medical claims engine

The head of medical claims does not need to be a technologist, but should know what the engine looks like. The key idea is one engine connected to every

provider, using standard health data formats such as HL7, so eligibility, pre-authorisation and claims all use the same rules and data.


EXHIBIT 10

One medical claims engine, connected to every provider



Source: KaribuTech AI analysis.

Build, buy or partner



Ask one question of every technology decision: will the next routine claim be decided faster, and the next complex one reviewed better, because of it?

The medical claims team in the AI era

AI takes over routine checking and paying. The team's work moves to clinical judgement on complex cases, managing providers, investigating real fraud and

helping members. The head of medical claims leads that shift.

EXHIBIT 11

How the main roles change

ROLE	TODAY	WITH AI
Claims assessor	Checks every claim by hand	→ Reviews exceptions flagged by AI
Clinical adviser	Reviews routine and complex requests	→ Focuses on complex care and clinical rules
Medical coder	Codes and checks claims manually	→ Tunes coding rules and reviews flagged claims
Fraud investigator	Reviews tip-offs and samples	→ Investigates the highest-risk referrals
Provider relations	Handles disputes and queries	→ Uses data to manage quality and tariffs
Head of medical claims	Tracks backlogs and costs	→ Tracks turnaround, savings, fairness and care outcomes

Source: KaribuTech AI analysis.

Who you work with

- **Clinical advisers:** set and review clinical rules.
- **Actuaries:** benefits, tariffs and medical loss ratio.
- **Provider networks:** contracts and quality.
- **Compliance and privacy:** health data and fair treatment.
- **IT and partners:** provider connections and HL7 data.

Skills to build in the team

Reading AI recommendations critically

Clinical coding

Data-driven provider management

Explaining decisions clearly

Health data privacy

Risks and guardrails in medical claims

In medical claims, a wrong automated decision can delay care. AI must make care faster, never slower, and must protect some of the most sensitive data people

have. These are the risks a head of medical claims owns.

EXHIBIT 12

Five medical claims risks, and the control for each

RISK	WHAT CAN GO WRONG	THE CONTROL
Delayed or denied care	Automated checks slow urgent treatment or wrongly decline it	Urgent cases bypass checks; a clinician on every decline
Health data exposure	Medical records leaked, misused or sent to unapproved tools	Strong security; consent; in-country data where required
Unfair outcomes	Some member groups or providers treated worse	Fairness testing; explanations; appeal routes
Provider disputes	Automated adjustments that providers do not understand	Clear reasons; provider feedback; fast review
Clinical rule errors	Out-of-date rules approve or decline the wrong care	Clinical advisers review rules regularly; monitoring

Source: KaribuTech AI analysis; IAIS, Application Paper on the supervision of AI (2025).


Eight questions to ask your team

- 01 How long does an urgent pre-authorisation take?
- 02 Can every decline be explained to the member and doctor?
- 03 Where is our members' health data stored and processed?
- 04 Have our rules been tested for fairness?
- 05 Which providers dispute our decisions most, and why?
- 06 When were our clinical rules last reviewed?
- 07 How many honest providers did fraud checks delay?
- 08 How many admissions could care management have prevented?

The medical claims scorecard and 12-month roadmap

EXHIBIT 13


Six numbers to run medical claims, plus four for the AI era



Source: KaribuTech AI analysis. Gold: keep within a target band; green: higher is better; red: lower is better.

EXHIBIT 14

Four quarters to AI-driven medical claims



Source: KaribuTech AI analysis.

WHERE TO START ON MONDAY

Pull last month's pre-authorisations and count how many met clear clinical rules. That number is how many patients could have been approved in minutes.

About the authors



Kennedy Chengeta, PhD

Co-founder, KaribuTech AI

Enterprise Architect and AI Leader · Camunda Chapter Lead, Africa

Leads AI, automation and integration programmes for insurers and banks across Southern Africa.



Justice Muguda

Chief Technology Officer, KaribuTech AI


Cloud, integration and cybersecurity

Leads KaribuTech AI's technology and engineering, from cloud platforms to secure connections between systems.

With contributions from the KaribuTech AI Engineering Team. To discuss this report, or to run a 100-day medical claims diagnostic, contact KaribuTech AI at kaributech.co.za.

References

1. Allianz (2025), announcements on AI agents in claims handling.
2. PwC, Stopping the leaks: claims leakage in insurance.
3. IBM (2025), Cost of a Data Breach Report 2025.
4. World Health Organization, International Classification of Diseases, ICD-10 (endorsed 1990).
5. Histories of German sickness insurance (1883) and the Baylor hospital plan (1929).
6. Tribune Online, Pulse Nigeria and Leadership (September 2026), reports on West Africa's first tele-robotic surgery.
7. IAIS (2025), Application Paper on the supervision of artificial intelligence, July 2025.
8. HL7 International, health data exchange standards.
9. KaribuTech AI, BluFountain Insure platform (Medishield medical core).
10. KaribuTech AI, The AI Blueprint for the Head of Claims (2026).



About KaribuTech AI

KaribuTech AI is a South African technology company specialising in artificial intelligence, machine learning, intelligent automation and integration for insurers and banks across Africa. Our teams design and build intelligent workflows, decision engines, document and computer-vision AI, and cloud-native core platforms, including the **BluFountain Insure** platform: an insurance AI

marketplace with core systems for life, short-term, medical, funeral and investment business. We combine deep process expertise (Camunda, IBM BPM, MuleSoft, webMethods and document automation platforms) with practical AI engineering, so that AI moves out of the lab and into the operations that customers feel.

Disclaimer: this report provides general information and does not constitute legal, actuarial or investment advice. Third-party figures are cited beneath each exhibit and have not been independently verified by KaribuTech AI. Illustrative figures are not forecasts.



KaribuTech AI

INSURANCE · LEADERSHIP SERIES · SEPTEMBER 2026

The AI Blueprint for the Head of Premium Collections

Collect more, lose fewer policies and treat customers fairly: AI for premium collection and persistency


By Kennedy Chengeta, PhD, and Justice Muguda
with the KaribuTech AI Engineering Team

kaributech.co.za



Contents

01	Introduction	126
02	Why premium collection must change	127
03	Lessons from history	128
04	Where premiums leak	129
05	Six moves for the Head of Premium Collections	130
06	Three customer lanes	132
07	Choosing the right payment rails	133
08	Generative AI and AI agents in collections	134
09	Collecting fairly	135
10	Premiums through brokers and employers	136
11	The 100-day quick-wins plan	137
12	The business case	138
13	What you need underneath	139
14	The collections team	140
15	Risks and guardrails	141
16	The scorecard and 12-month roadmap	142
17	About the authors and KaribuTech AI	143



About this report

This report is written for heads of premium collections, billing, credit control and customer retention at insurers in developing markets.

It shows how AI can raise collection rates, cut lapses and failed payments, match every payment to the right policy and treat customers in difficulty fairly: where premiums leak, six moves, customer lanes, payment rails, fair collections, broker and employer premiums, a 100-day plan, the business case, the team, the risks and a scorecard.

External data are sourced beneath each exhibit, and a full list of references appears at the end. Frameworks and illustrative figures without an external source are the authors' own.

WRITTEN FOR

Heads of collections

Billing

Retention

Finance

Introduction

An insurer only keeps its promise if it keeps its customers. Every premium that goes uncollected is a policy at risk and a family that may be left without cover.

Many insurers still collect premiums as they did decades ago: a fixed debit date, a cash desk, a letter when a payment fails. Customers with irregular incomes miss payments, policies lapse quietly, and payments sit in suspense accounts because no one can match them to a policy.

AI changes this. It can collect on the day customers are paid, retry failed payments at the right moment, predict who is likely to lapse, match every payment automatically and send the right reminder in the right language.

This report sets out **six moves** for the head of premium collections, how to choose **payment rails**, how to collect **fairly**, how to handle **broker and employer** premiums, a **100-day plan**, the business case and a **scorecard**.

THREE QUESTIONS THIS REPORT ANSWERS

01 Why do customers miss premiums, and when?

02 How much money sits unmatched, and why?

03 How do we cut lapses while treating customers fairly?

AT A GLANCE: SIX MOVES FOR PREMIUM COLLECTIONS

- 1** Offer the ways customers already pay
- 2** Collect on payday, not on a fixed date
- 3** Predict and prevent lapses
- 4** Match every payment automatically
- 5** Make arrears gentle and effective
- 6** Win customers back

Why premium collection must change

Five pressures are squeezing traditional premium collection. Each one is also a chance

to keep more customers and more premium.

Incomes are irregular

Many customers are paid weekly, daily or by the season. A fixed monthly debit date misses them.

Failed payments are costly

Every bounced debit costs bank fees, staff time and often the policy itself.

Money sits unmatched

Payments arrive without a clear policy reference and wait in suspense accounts for weeks.

Lapses erode the book

Policies lost in the first year waste the cost of selling them and leave customers uncovered.

Customers have moved to mobile money

Customers who manage money on their phones expect to pay premiums the same way.

EXHIBIT 1

Payments are moving to the phone

2.3bn

registered mobile money accounts worldwide

\$2tn+

moved through mobile money in 2025

15%

of African micro products used mobile money in 2023, up from 12% in 2022

17%

relied on cash in 2023, down from 22% in 2022

Source: GSMA, State of the Industry Report on Mobile Money (2026); Microinsurance Network, Landscape of Microinsurance (2024).

The cheapest policy to keep is the one already on the books. Every point of persistency is premium you do not have to sell again.

Lessons from history: collecting small premiums has always needed new ideas

Insuring ordinary families has always depended on collecting small amounts, regularly, at low cost. Each

breakthrough in reaching the mass market came from a better way to collect.

1700s

Friendly societies

In 18th-century Britain, working people formed friendly societies, paying small regular contributions into a shared fund for sickness and burial. Collection happened where members already met, often at the local inn.

Today: collect through the apps, shops and employers customers already use.

THE LESSON

Collect where customers already gather.

1850s

The collector at the door

Industrial assurance companies sent agents to working families every week to collect a few pennies. The weekly visit kept millions of small policies in force.

Today: payday-timed collection and friendly reminders on WhatsApp.

THE LESSON

Collect on the customer's rhythm, not the insurer's.

1967

Direct debit

The UK introduced direct debit in 1967, letting companies collect regular payments automatically from bank accounts. Collection costs fell and lapses from forgotten payments dropped.

Today: smart retries and automatic matching make collection nearly effortless.

THE LESSON

Automation keeps policies alive.

2007

Premiums by phone

Safaricom launched M-Pesa in Kenya in 2007. Mobile money made it possible to pay tiny premiums with no bank account and no collector.

Today: premiums from daily to monthly, with no bank account.

THE LESSON

The phone is today's collector at the door.

Source: Histories of friendly societies and industrial assurance in Britain; Bacs, history of Direct Debit (1967); Safaricom (2007).

AI is the next step: the weekly collector's knowledge of each customer, and the reliability of the direct debit, at the cost of a text message.

Where premiums leak, and how AI plugs the gaps

Every premium passes through the same eight steps, from the bill to a policy that stays in force. At each one,

money or customers can leak away. The exhibit shows the typical leak and what AI does about it.

EXHIBIT 2

The collection cycle: eight steps from billing to reinstatement

	WHERE IT LEAKS	WHAT AI DOES
Bill Premium due is set up	Wrong amounts; unclear schedules	Schedules checked; clear bills by message
Payment method How the customer pays	One method offered for everyone	Method matched to how the customer is paid
Collect The payment is attempted	Fixed date misses payday	Collection timed to each customer's payday
Failed payment Debit bounces or is missed	Blind retries and bank fees	Smart retries at the best moment
Allocate Money matched to a policy	Payments left in suspense	Automatic matching of every payment
Remind Customer told about arrears	Late, generic letters	Early, friendly messages in local languages
Grace and arrears Time to catch up	No help for customers in difficulty	Payment holidays, smaller cover, hardship support
Lapse and reinstate Policy ends or restarts	Lapsed customers forgotten	Win-back offers and easy reinstatement

Source: KaribuTech AI analysis.

Most lapses are not decisions. They are missed payments that no one followed up in time.

Six moves for the Head of Premium Collections

Collection leaders who use AI well do six things differently. Each is about fitting how

customers are paid and keeping cover in force, not about choosing a technology.

MOVE 1

01

Offer the ways customers already pay

Customers pay for everything else by mobile money, debit order, card, payroll or at a local shop. Premiums should be the same. Every extra payment option removes a reason to miss a payment.

What leaders do

- Accept mobile money, debit orders, cards, payroll deduction and agent or shop payments.
- Let customers change how they pay on WhatsApp or USSD in minutes.
- Use AI to suggest the method each customer is most likely to keep using.

ASK YOUR TEAM

How many of our customers could only pay in a way they no longer use every day?

MOVE 2

02

Collect on payday, not on a fixed date

A debit on the 1st fails for someone paid on the 25th. AI can learn when each customer's money arrives, from payment history and, with consent, mobile-money patterns, and collect then.

What leaders do

- Let customers choose their collection date, or suggest one from their history.
- Retry failed payments when money is most likely to be there, not blindly.
- Split premiums into weekly or daily amounts where incomes are irregular.

ASK YOUR TEAM

How many failed debits last month were simply attempted on the wrong day?

MOVE 3

03

Predict and prevent lapses

Lapses rarely come without warning: a missed payment, a changed phone number, a drop in income, a complaint. AI can score every policy for lapse risk and trigger help before the policy is lost.

What leaders do

- Score every policy for lapse risk every month.
- Contact high-risk customers early, with options rather than demands.
- Feed lapse insights back to sales, so products fit customers from the start.

ASK YOUR TEAM

Which of our customers are most likely to lapse next month, and what are we doing about it?

MOVE 4

04

Match every payment automatically

Payments that cannot be matched to a policy leave customers wrongly in arrears and money idle in suspense. AI can read references, amounts, phone numbers and names to match almost every payment, and send the rest to a person.

What leaders do

- Match payments automatically using every clue available, not just the reference.
- Clear the suspense account daily, not monthly.
- Tell customers immediately that their payment was received.

ASK YOUR TEAM

How much money is sitting in our suspense account today, and for how long?

MOVE 5

05

Make arrears gentle and effective

Harsh collection damages trust and rarely works with customers who cannot pay. AI can tailor the message, timing and offer to each customer: a reminder for the forgetful, a payment plan for the struggling.

What leaders do

- Send early, friendly reminders in the customer's language and channel.
- Offer options: a later date, smaller payments or reduced cover instead of a lapse.
- Flag customers who may be in hardship for a person to help.

ASK YOUR TEAM

Do we offer customers in difficulty any option other than paying in full or losing their cover?

MOVE 6

06

Win customers back

Many lapsed customers would come back if asked at the right time with a simple offer. AI can identify who is worth contacting, when and with what.

What leaders do

- Contact recently lapsed customers with a simple way to restart.
- Make reinstatement quick, with minimal new forms.
- Learn from why customers lapsed and fix the causes.

ASK YOUR TEAM

How many customers who lapsed last year did we ever invite back?

FROM OUR WORK

Premiums collected where customers already pay

KaribuTech AI's BluFountain Insure platform is built mobile-first, with WhatsApp and self-service portals for brokers, merchants and employers. Customers can pay and track premiums on the channels they already use, employers and brokers can reconcile their premiums online, and every payment is matched to the right policy in the core system.

Mobile first

premiums paid and tracked on the phone

Portals

for brokers, employers and merchants

Matched

every payment linked to its policy


Three customer lanes: the right help at the right time

Not every customer who misses a payment needs the same response. AI sorts customers into three lanes, based on payment history and lapse risk, so the team

spends its time where it makes a difference. People set the rules, and every customer in difficulty can reach a person.

EXHIBIT 3

Three lanes for premium collection, from light touch to personal help



MORE AUTOMATION →

← MORE PERSONAL HELP

Source: KaribuTech AI analysis.

What decides the lane

- **Payment history:** on time, late or missed.
- **Lapse score:** the AI's estimate of lapse risk.
- **Signals of difficulty:** complaints, changed details, requests.
- **Policy value:** size and type of cover at stake.

Guardrails the team owns

- **Rules set by people:** collections and compliance agree the lanes.
- **Fair treatment:** outcomes checked across customer groups.
- **A person for hardship:** no customer in difficulty handled by AI alone.
- **No surprise debits:** customers always know when and how much.

Choosing the right payment rails

Each way of paying has its own cost, reliability and fit for different customers. Offering several rails, and using

AI to steer each customer to the one they are most likely to keep, is the simplest way to lift collection rates.

EXHIBIT 4

Six ways customers pay premiums, and what AI adds

RAIL	BEST FOR	WHAT AI ADDS	WATCH
Mobile money	Customers without bank accounts; small, frequent premiums	Collection timed to wallet top-ups; instant confirmations	Wallet balances; SIM swaps
Debit orders	Salaried customers with bank accounts	Payday prediction; smart retries	Bank fees on failed debits
Payroll deduction	Group schemes and employees of partner employers	Automatic matching of employer schedules	Late employer payments
Cards and instant payments	Digital customers; one-off payments	Payment links in reminders; card updates	Expired cards; chargebacks
Airtime	Very small cover sold with phone plans	Deduction timed to top-ups	Customer understanding of cover
Agents and shops	Customers who prefer cash	Receipts matched automatically	Cash handling; delays

Source: KaribuTech AI analysis.

Three rules for payment rails

01

Offer at least three

No customer should have only one way to pay.

02

Make switching easy

Customers change rails on WhatsApp or USSD in minutes.

03

Track every rail

Measure cost and success rate for each rail every month.

The best payment method is the one the customer is already using to get paid.

Generative AI and AI agents in collections

Generative AI reads and writes; AI agents plan and carry out tasks. In premium collection, that means friendly reminders in local languages, a payment

assistant on WhatsApp, explanations of unmatched payments and help for staff in difficult conversations.

EXHIBIT 5

Six jobs AI can take on in collections today

Send reminders Writes clear, friendly reminders in the customer's language, timed to their payday ACT WITHIN LIMITS	Answer payment questions Tells customers what they owe, when and how to pay, on WhatsApp any time ACT WITHIN LIMITS	Explain unmatched payments Suggests which policy a payment belongs to, with reasons, for staff to confirm RECOMMEND
Prepare hardship calls Summarises the customer's history and options before a retention call ASSIST	Reconcile broker statements Reads broker and employer schedules and flags differences RECOMMEND	Draft win-back offers Prepares tailored offers for lapsed customers for a person to approve RECOMMEND

Source: KaribuTech AI analysis. Labels show the level of AI autonomy we recommend.

What stays with people

- Every conversation with a customer in hardship.
- Decisions to lapse or cancel a policy.
- Disputes about payments.
- Setting the rules for reminders and offers.

Guardrails that make it safe

- **Plain, respectful language:** no threats or pressure.
- **No customer data in public AI tools.**
- **Every message logged** for quality and complaints.
- **Limits on contact:** no more reminders than the rules allow.

WHERE TO START

Start with payday-timed reminders and a WhatsApp payment assistant. They are low risk and customers notice straight away.

Collecting fairly: good outcomes for customers in difficulty

Customers who miss premiums are often dealing with a lost job, illness or a bad harvest. Treating them fairly is right, and it also works: customers helped through a

hard month are more likely to stay. AI can help spot difficulty early, but people must handle it.

EXHIBIT 6

Five fair-collection practices, and how AI supports them

PRACTICE	WHAT IT MEANS	HOW AI HELPS
Clear information	Customers always know what they owe, when and what happens if they miss a payment	Plain-language messages; instant answers on WhatsApp
Reasonable contact	Reminders at sensible times and frequencies, never harassing	Contact limits built into the rules; timing matched to payday
Options before lapse	A later date, smaller payments or reduced cover before cancelling	Suggests the option most likely to keep cover in force
Help for hardship	Customers in difficulty reach a trained person	Flags signs of hardship for a person to follow up
Fair across groups	No group of customers treated worse than others	Outcomes compared by age, gender, region and income

Source: KaribuTech AI analysis.

Three measures of fair collection

01

Complaints from collections

Tracked monthly by channel and customer group.

02

Lapses after hardship

How many customers who asked for help still lost their cover.

03

Options taken up

How many customers used a payment plan or reduced cover instead of lapsing.

A customer helped through a hard month is often a customer for life. A customer pushed into a lapse rarely comes back.

Premiums through brokers and employers

For many insurers, a large share of premium arrives through brokers, employers and other intermediaries, often in bulk, late and with schedules that do not match.

AI can read the schedules, match every line and chase what is missing.

EXHIBIT 7

Where intermediary premiums get stuck, and what AI does

SOURCE	WHERE IT GETS STUCK	WHAT AI DOES
Brokers	Premiums held for weeks; statements that do not match policies	Reads statements; matches lines; flags old balances
Employers and group schemes	Payroll schedules with joiners, leavers and errors	Reconciles schedules with the member list automatically
Banks and partners	Bulk payments with little detail	Splits bulk payments to policies using every clue
Merchants and platforms	Many tiny premiums from sales	Matches premiums to sales records daily
Government and unions	Deductions paid late or in arrears	Tracks what is due and chases automatically

Source: KaribuTech AI analysis.

Three rules for intermediary premiums

01

Self-service portals

Brokers and employers upload schedules and see balances online.

02

Agreed payment terms

Clear deadlines for passing on premiums, tracked automatically.

03

Balances by age

Report what each intermediary holds and for how long.

The 100-day quick-wins plan for collections

Do not start with a large programme. Pick one product and one payment rail, deliver a working improvement in 100 days, measure it against today, then reuse the

approach. Each cycle is run by a small team from collections, finance, customer service and technology.

EXHIBIT 8

One collection journey in 100 days: design, build, go live, measure

DAYS 1-30

Discover and design

- Baseline collection rate and lapses
- Map failed payments and suspense
- Agree lanes, messages and limits

DAYS 31-60

Build

- Payday prediction and smart retries
- Automatic payment matching
- Reminders on WhatsApp and SMS

DAYS 61-100

Go live and measure

- Go live on one product and rail
- Compare with a control group
- Report results to the executive team

Source: KaribuTech AI quick-wins approach.

Three quick wins we recommend

QUICK WIN 1

Payday collection and smart retries

- Collection dates matched to paydays
- Retries at the best moment
- Customers choose their date
- Fees on failed debits tracked

MEASURE

Collection rate, failed debit rate, bank fees

QUICK WIN 2

Automatic payment matching

- Every clue used to match payments
- Suspense account cleared daily
- Instant payment confirmations
- Exceptions sent to a person

MEASURE

Unallocated cash, days in suspense

QUICK WIN 3

Lapse prediction and nudges

- Every policy scored monthly
- Friendly reminders and options
- Hardship flagged for a person
- Win-back offers for recent lapses

MEASURE

Persistency, lapses saved, reinstatements

The business case for AI in premium collection

The value of AI in collections comes from premium collected that would otherwise be lost, lower collection costs and money matched faster. The example below is

illustrative, for a mid-sized insurer; your own baseline will differ.

EXHIBIT 9

Illustrative example: 200,000 policies, average premium of \$10 a month

Measure	Before	After AI	Annual value
Collection rate	88%	93%	≈ \$1.2m more premium
Cost to collect each premium	\$1.00	\$0.70	≈ \$0.72m saved
Money in suspense	\$2.0m	\$0.4m	\$1.6m matched faster
Lapse in year one	25%	20%	More policies kept
Total annual value			≈ \$1.9m + policies kept

Source: Illustrative KaribuTech AI example, not a forecast. Premium due: 200,000 × \$10 × 12 = \$24m a year; five points more collected = \$1.2m. Costs: \$0.30 saved on 2.4 million premium payments.

Where the value comes from

- **More premium collected:** payments timed to paydays and retried smartly.
- **Lower costs:** fewer failed debits, calls and manual matching.
- **Faster matching:** money reaches the right policy the same day.
- **More policies kept:** lapses prevented by early help.

Where the money goes

Payment connections	25–35%
Matching and billing tools	20–25%
Lapse prediction and messaging	15–20%
Integration with core systems	15–20%
Training and governance	10%

Typical split of a first-year budget (illustrative)

THREE RULES FOR A CREDIBLE CASE

- 1 Baseline collection, lapses and suspense first
- 2 Test against a control group before rolling out
- 3 Report value and complaints every quarter


What you need underneath: the collections engine

The head of premium collections does not need to be a technologist, but should know what the engine looks like. The key idea is one collections engine connected

to every payment rail and every intermediary, so every payment is collected, matched and confirmed the same way.

EXHIBIT 10

One collections engine, connected to every way customers pay



Source: KaribuTech AI analysis.

Build, buy or partner

Buy

Payment gateways, mobile-money connections, messaging tools

Build

Your lanes, retry rules, lapse models and messages, on your own data

Partner

Mobile-money providers, banks, employers and retail networks

Ask one question of every technology decision: will the next payment rail or partner be faster to add because of it?

The collections team in the AI era

AI takes over the routine work of collecting, retrying and matching. The team's work moves to helping customers in difficulty, managing payment partners and

fixing the causes of lapses. The head of premium collections leads that shift.

EXHIBIT 11

How the main roles change

ROLE	TODAY	WITH AI
Collections clerk	Chases failed debits by phone	→ Handles exceptions AI cannot resolve
Cash-allocation clerk	Matches payments by hand	→ Reviews unmatched payments with AI suggestions
Call-centre agent	Reads scripts about arrears	→ Retention advisor helping customers in difficulty
Debit-order administrator	Loads and fixes debit files	→ Manages payment rails and partner performance
Broker accounts clerk	Reconciles statements in spreadsheets	→ Manages intermediary balances and terms
Head of collections	Tracks collection rate	→ Tracks collection, persistency, fairness and cost

Source: KaribuTech AI analysis.

Who you work with

- **Sales:** products and payment methods that fit customers.
- **Finance:** reconciliation and the general ledger.
- **Customer service:** complaints and hardship.
- **Compliance:** fair treatment and contact rules.
- **IT and partners:** banks, mobile money and employers.

Skills to build in the team

Empathy in hardship conversations

Reading AI suggestions critically

Managing payment partners

Explaining options simply

Spotting unfair outcomes

Risks and guardrails in collections

Collecting more must not mean pressuring customers or taking the wrong money. AI can scale good practice, but it can also scale mistakes. These are the risks a

head of premium collections owns, and the controls that manage them.

EXHIBIT 12

Five collection risks, and the control for each

RISK	WHAT CAN GO WRONG	THE CONTROL
Customer harm	Too many reminders, pressure or lapses for customers in difficulty	Contact limits; hardship options; a person for vulnerable customers
Wrong debits	Double debits, wrong amounts or wrong accounts	Checks before every collection run; instant refunds
Misallocated money	Payments matched to the wrong policy	Confidence thresholds; people review uncertain matches
Payment fraud	Payments redirected or accounts changed by fraudsters	Verification of account changes; alerts on unusual activity
Unfair lapse scores	Some customer groups flagged or treated worse	Fairness tests; explanations; regular model review

Source: KaribuTech AI analysis; IAIS, Application Paper on the supervision of AI (2025).


Eight questions to ask your team

- 01 How many reminders does a customer in arrears receive?
- 02 What options do we offer before lapsing a policy?
- 03 How much money is in suspense, and for how long?
- 04 Have our lapse scores been tested for fairness?
- 05 How quickly do we refund a wrong debit?
- 06 Which intermediaries hold premiums the longest?
- 07 Which payment rail fails most often?
- 08 How many lapsed customers did we invite back last year?

The collections scorecard and 12-month roadmap

EXHIBIT 13


Six numbers to run premium collections, plus four for the AI era



Source: KaribuTech AI analysis. Gold: keep within a target band; green: higher is better; red: lower is better.

EXHIBIT 14

Four quarters to AI-driven premium collection



Source: KaribuTech AI analysis.

WHERE TO START ON MONDAY

Pull last month's failed debits and check how many were attempted before the customer's payday. That number is where the first 100 days should start.

About the authors



Kennedy Chengeta, PhD

Co-founder, KaribuTech AI

Enterprise Architect and AI Leader · Camunda Chapter Lead, Africa

Leads AI, automation and integration programmes for insurers and banks across Southern Africa.



Justice Muguda

Chief Technology Officer, KaribuTech AI


Cloud, integration and cybersecurity

Leads KaribuTech AI's technology and engineering, from cloud platforms to secure connections between systems.

With contributions from the KaribuTech AI Engineering Team. To discuss this report, or to run a 100-day premium collections diagnostic, contact KaribuTech AI at kaributech.co.za.

References

1. GSMA (2026), State of the Industry Report on Mobile Money.
2. Microinsurance Network (2024), Landscape of Microinsurance.
3. Histories of friendly societies in 18th-century Britain.
4. Histories of industrial assurance in Britain, including Prudential.
5. Bacs, history of Direct Debit in the UK (from 1967).
6. Safaricom (2007), launch of M-Pesa in Kenya.
7. IAIS (2025), Application Paper on the supervision of artificial intelligence, July 2025.
8. KaribuTech AI, BluFountain Insure platform (mobile, WhatsApp and self-service portals).
9. KaribuTech AI, The AI Blueprint for the Insurance CEO (2026).
10. KaribuTech AI, quick-wins approach (proprietary).



About KaribuTech AI

KaribuTech AI is a South African technology company specialising in artificial intelligence, machine learning, intelligent automation and integration for insurers and banks across Africa. Our teams design and build intelligent workflows, decision engines, document and computer-vision AI, and cloud-native core platforms, including the **BluFountain Insure** platform: an insurance AI

marketplace with core systems for life, short-term, medical, funeral and investment business. We combine deep process expertise (Camunda, IBM BPM, MuleSoft, webMethods and document automation platforms) with practical AI engineering, so that AI moves out of the lab and into the operations that customers feel.

Disclaimer: this report provides general information and does not constitute legal, actuarial or investment advice. Third-party figures are cited beneath each exhibit and have not been independently verified by KaribuTech AI. Illustrative figures are not forecasts.



KaribuTech AI

INSURANCE · LEADERSHIP SERIES · SEPTEMBER 2026

The AI Blueprint for the Head of Actuarial

Sharper prices, faster reserves and models you can explain: AI for the actuarial function


By Kennedy Chengeta, PhD, and Justice Muguda
with the KaribuTech AI Engineering Team

kaributech.co.za



Contents

01	Introduction	146
02	Why the actuarial function must change	147
03	Lessons from history	148
04	Where actuarial time goes	149
05	Six moves for the Head of Actuarial	150
06	Current models and AI models	152
07	Explainable, governed AI	153
08	Generative AI and AI agents for actuaries	154
09	IFRS 17 and reporting	155
10	Climate, catastrophe and new risks	156
11	The 100-day quick-wins plan	157
12	The business case	158
13	What you need underneath	159
14	The actuarial team	160
15	Risks and guardrails	161
16	The scorecard and 12-month roadmap	162
17	About the authors and KaribuTech AI	163



About this report

This report is written for heads of actuarial, chief actuaries and their teams at insurers, reinsurers and pension funds in developing markets.

It shows how AI can sharpen pricing, speed up reserving and valuation, strengthen IFRS 17 reporting and help actuaries govern every model in the business: the pressures on the function, six moves, current and AI models side by side, explainability, a 100-day plan, the business case, the team, the risks and a scorecard.

External data are sourced beneath each exhibit, and a full list of references appears at the end. Frameworks and illustrative figures without an external source are the authors' own.

WRITTEN FOR

Chief actuaries

Pricing teams

Reserving teams

Model risk

Introduction

Actuaries have always been the insurer's data scientists. Now the tools have changed: more data, faster computers and AI models that can find patterns no table or triangle could.

Yet many actuarial teams still spend most of their time preparing data, reconciling spreadsheets and running valuations by hand. IFRS 17 has multiplied the reporting load. Business leaders want sharper prices and faster answers, and regulators want every model explained.

AI can take on much of this work: cleaning and checking data, finding new rating factors, projecting claims and lapses, running and explaining valuations, and drafting reports. The judgement, the assumptions and the sign-off stay with actuaries.

This report sets out **six moves** for the head of actuarial, a side-by-side view of **current and AI models**, how to keep AI **explainable and governed**, what AI means for **IFRS 17** and **climate risk**, a **100-day plan**, the business case and a **scorecard**, in plain language.

THREE QUESTIONS THIS REPORT ANSWERS

01 Where does actuarial time go today, and what could AI take on?

02 When should AI models lead, and when should they only support?

03 How do we explain, validate and sign off AI models?

AT A GLANCE: SIX MOVES FOR THE ACTUARIAL FUNCTION

- 1** Fix the data first
- 2** Price with richer data, explainably
- 3** Reserve faster and smarter
- 4** Automate the valuation run
- 5** Own model risk for all AI
- 6** Look ahead: climate and new risks

Why the actuarial function must change

Five pressures are reshaping the actuarial function at once. Each one is also an

opportunity for actuaries to move from producing numbers to shaping decisions.

IFRS 17 raised the bar

Policy-level data, annual cohorts and a full analysis of change every period have multiplied the reporting workload.

The business wants sharper prices

Competitors price with more factors and update more often. Broad rating tables leave good risks overpriced and bad risks underpriced.

Data is thin and messy

In our markets, histories are short, informal incomes are common and data sits in many systems.

New risks are harder to model

Climate, cyber, pandemics and AI itself do not follow the patterns of the past.

Every model must be explained

Regulators and boards expect actuaries to explain and stand behind every model the insurer uses, including AI built outside the actuarial team.

EXHIBIT 1

Why now: the tools have changed faster than the workload

280x

cheaper to run a GPT-3.5-level AI model in 2024 than in 2022

1 Jan 2023

IFRS 17 in force, with policy-level and cohort-level data

\$400bn

AI value at stake in pricing, underwriting and promotion

2025

IAIS guidance on supervising AI in insurance published

Source: Stanford HAI, AI Index 2025; IASB, IFRS 17 Insurance Contracts; McKinsey & Company; IAIS, Application Paper on the supervision of AI (2025).

The actuary's job is not to run the model. It is to know whether the model is right, and to explain why.

Lessons from history: the profession was born from data

The actuarial profession exists because, at each turning point, someone found a better way to turn data into

prices and promises. AI is the latest turning point, and the lessons from the first ones still apply.

1693



Halley's life table

The astronomer Edmond Halley used birth and death records from the city of Breslau to build one of the first life tables and showed how to use it to price annuities. Guesswork gave way to evidence.

Today: mobile, satellite and health data are the new parish records.

THE LESSON

New data sources change what can be priced.

1762



Premiums by age

The Equitable Society in London priced life cover by age using mortality data, and in 1775 appointed William Morgan as its actuary. Pricing became a profession.

Today: machine learning finds factors that tables miss.

THE LESSON

Better models create new, fairer products.

1848



A profession with standards

The Institute of Actuaries was founded in London, setting exams, standards and professional judgement at the heart of the work.

Today: validation and sign-off for every AI model.

THE LESSON

Trust comes from standards and sign-off, not only from calculations.

Today



Machine learning

AI models can learn from millions of records and hundreds of factors. The actuary's role shifts to choosing, testing, explaining and signing off those models.

Today: actuaries move from building every model to governing all of them.

THE LESSON

The tools change; the duty to explain and stand behind the numbers does not.

Source: E. Halley (1693), Philosophical Transactions of the Royal Society; M. E. Ogborn, Equitable Assurances (1962); Institute and Faculty of Actuaries, history of the profession.

Actuaries led the first data revolution in insurance. They are best placed to lead the AI one, if they move from building every model to governing all of them.

Where actuarial time goes, and what AI changes

Most actuarial work follows the same cycle, from data to decisions. In many teams, the early steps consume most of the time, leaving too little for analysis and

advice. The exhibit shows where time typically goes at each step and what AI does.

EXHIBIT 2

The actuarial cycle: eight steps from data to decisions

	WHERE TIME GOES TODAY	WHAT AI DOES
Data preparation Extracting, cleaning, reconciling	Weeks of manual cleaning and checks	Automated checks; issues flagged at source
Experience analysis Actual versus expected	Small samples; slow refresh	Continuous monitoring; early trend alerts
Pricing Rates and rating factors	Broad tables; yearly updates	More factors, tested; faster updates
Reserving Outstanding and future claims	Triangles rebuilt by hand	Machine-learning estimates alongside triangles
Capital Solvency and stress tests	Few scenarios; slow runs	Many scenarios run quickly
Valuation and IFRS 17 Liabilities, CSM, reporting	Long runs; manual analysis of change	Automated runs; drafted explanations
Asset and liability Matching assets to promises	Periodic studies	Faster projections; climate scenarios
Review and advice Board papers and sign-off	Squeezed into the last days	More time for judgement and advice

Source: KaribuTech AI analysis.

Shift the balance: less time preparing numbers, more time explaining them and advising the business.

Six moves for the Head of Actuarial

Actuarial leaders who use AI well do six things differently. Each is about the quality of the

numbers and the way the team works, not about choosing a technology.

MOVE 1

01

Fix the data first

No model is better than its data. Actuarial teams lose weeks every quarter extracting, cleaning and reconciling. AI can check data automatically, spot errors at the source and keep one trusted set of policy, claims and premium data.

What leaders do

- Automate data checks and reconciliations for every valuation and pricing run.
- Agree one trusted data source with finance and claims, built for IFRS 17 cohorts.
- Report data quality to the board alongside the results.

ASK YOUR TEAM

How many days of each valuation cycle do we spend preparing data?

MOVE 2

02

Price with richer data, explainably

Machine-learning models can find interactions and new rating factors that traditional tables miss. Run alongside the current pricing model, they sharpen prices and show which factors really matter.

What leaders do

- Run an AI challenger model alongside the current pricing model on one product.
- Use AI to find new factors, then test and adopt them in an explainable model.
- Check every price change for fairness across customer groups.

ASK YOUR TEAM

Which good risks are we overpricing, and which bad risks are we underpricing?

MOVE 3

03

Reserve faster and smarter

Reserving still relies heavily on triangles rebuilt by hand. Machine-learning models can estimate reserves claim by claim and spot changes in trends earlier, as a second opinion beside traditional methods.

What leaders do

- Run machine-learning reserve estimates alongside chain-ladder methods.
- Use claim-level data to spot trends before they reach the triangle.
- Explain every difference between methods in the reserving report.

ASK YOUR TEAM

How quickly would we notice if claims inflation changed direction this quarter?

MOVE 4

04

Automate the valuation run

IFRS 17 valuations involve many steps: data, projections, CSM roll-forward, analysis of change and disclosures. Automating the run and drafting the explanations frees days for review and judgement.

What leaders do

- Automate the valuation run end to end, with checks at every step.
- Let AI draft the analysis-of-change narrative for actuaries to review.
- Set a target: close the valuation in days, not weeks.

ASK YOUR TEAM

What would our team do with an extra week each quarter?

MOVE 5

Own model risk for all AI

AI models are appearing across the insurer: in claims, fraud, underwriting and marketing. Actuaries have the skills to test and challenge them. The best functions act as the independent reviewer for every model that affects price, reserves or customers.

What leaders do

- Keep a register of every model, with an owner and a review date.
- Validate AI models independently: accuracy, fairness and stability.
- Sign off only what you can explain to the board and the regulator.

ASK YOUR TEAM

Which AI models in our business have never been reviewed by an actuary?

05

MOVE 6

Look ahead: climate and new risks

The past is a weaker guide to climate, cyber and emerging risks. AI can combine satellite, weather and economic data and generate many more scenarios, helping actuaries price and reserve for risks without long histories.

What leaders do

- Use satellite and weather data in property and agriculture pricing.
- Run more climate and economic scenarios in capital and ALM work.
- Build pricing approaches for new covers such as cyber and AI liability.

ASK YOUR TEAM

Which of our biggest risks has the least historical data behind it?

06

FROM OUR WORK

Actuarial models and valuation, built into the platform

KaribuTech AI's BluFountain Insure platform includes shared actuarial-model and valuation engines that sit beside the life, short-term, medical and funeral cores, alongside IFRS 17 and IFRS 9 compliance. Policy, premium and claims data flow into the models without spreadsheets in between, so actuaries spend their time on assumptions, review and sign-off.

One data flow

from policy and claims systems to models

IFRS 17 ready

valuation built for cohorts and CSM

Actuaries sign off

models run automatically; people decide

Current models and AI models: use each where it wins

AI does not replace the actuary's models; it adds new ones. Traditional models are transparent, stable and trusted by regulators. AI models find patterns in large,

messy data. The best teams combine them, and decide area by area which should lead.

EXHIBIT 3

Six actuarial areas: today's model, the AI model and who leads

AREA	CURRENT MODEL	AI MODEL	WHO LEADS
Pricing	Rating tables and generalised linear models	Machine-learning models that find new factors and interactions	HYBRID
Reserving	Chain-ladder and other triangle methods	Claim-by-claim machine-learning estimates	HYBRID
Mortality and lapse	Standard tables adjusted for experience	Models using behaviour, payment and health data	AI SUPPORTS
Capital and stress tests	Standard formula and a few scenarios	Thousands of scenarios generated and run quickly	AI SUPPORTS
Catastrophe and climate	Vendor catastrophe models	Satellite and weather data; forward-looking scenarios	HYBRID
Customer behaviour	Fixed assumptions for take-up and lapse	Predictions of lapse, take-up and claims by customer	AI LEADS

Source: KaribuTech AI analysis.

Three ways to combine them

01

AI finds, the model prices

AI searches the data for new factors; the best ones are tested and added to the explainable pricing model.

02

Champion and challenger

The current model stays in charge; an AI model runs beside it. Switch only when the challenger wins consistently and can be explained.

03

AI around the model

AI prepares data, runs checks and drafts explanations, while the core actuarial model stays unchanged.

Opening the black box: explainable, governed AI

An actuary must be able to explain every number to the board, the regulator and, increasingly, the customer. AI models can be explained, if the right checks are built in

from the start. These are the tools and controls that make AI acceptable in actuarial work, in plain language.

EXHIBIT 4

Five ways to explain and check an AI model

1

Which factors drove this result?

Shows, for any price or estimate, how much each factor pushed it up or down. Lets an actuary explain a single customer's price.

2

How does the result change with each factor?

Charts how the prediction moves as one factor changes, so the team can check the shape makes sense.

3

Does it beat the current model?

Runs the AI model beside the current one on the same data and compares results over time.

4

Does it hold up on past data?

Tests the model on periods it was not built on, to check it predicts well out of sample.

5

What happens in a stress?

Runs extreme but plausible scenarios to see whether the model behaves sensibly.

Source: KaribuTech AI analysis.

The sign-off ladder for every model

Built

Model built and documented by its owner

Tested

Accuracy, stability and fairness checked

Validated

Independent review by a second actuary or model-risk team

Approved

Signed off by the head of actuarial or model committee

Monitored

Tracked monthly; reviewed at least once a year

If you cannot explain it to the board, do not sign it off. Explainability is not a brake on AI; it is what lets actuaries use it with confidence.

Generative AI and AI agents for actuaries

Generative AI reads and writes; AI agents plan and carry out tasks. For actuaries, that means help with the words and the routine work around the models:

documentation, data checks, report drafting and research. The calculations and the judgement stay with the team.

EXHIBIT 5

Six jobs AI can take on in the actuarial team today

Check the data Flags missing, inconsistent and unusual records before a run starts, with a short report ACT WITHIN LIMITS	Document the models Writes and updates model documentation from code and specifications ASSIST	Explain the movements Drafts the analysis-of-change and variance commentary for actuaries to review RECOMMEND
Draft board papers Turns results into clear board and committee papers in plain language RECOMMEND	Check regulatory returns Compares returns with previous periods and rules, and flags inconsistencies RECOMMEND	Research assumptions Summarises industry studies, regulations and market data for assumption setting ASSIST

Source: KaribuTech AI analysis. Labels show the level of AI autonomy we recommend.

What stays with actuaries

- Every assumption and every calculation.
- Sign-off of results and models.
- Advice to the board and the regulator.
- Professional judgement on uncertainty.

Guardrails that make it safe

- **Approved sources only:** AI works from the insurer's own data and documents.
- **No confidential data in public AI tools.**
- **Every draft reviewed** before it leaves the team.
- **Numbers never taken from AI text:** always from the models.

WHERE TO START

Start with data checks and model documentation. They save many hours, carry little risk and make every later AI step safer.


IFRS 17 and reporting: from long closes to clear explanations

IFRS 17, in force since 1 January 2023, turned insurers into data companies. It requires detailed data by portfolio, year of sale and profitability, and a full

explanation of how profit emerges. AI and automation make that workload manageable, and free actuaries to explain the results.


EXHIBIT 6

The IFRS 17 liability in three building blocks, and where AI helps



Source: IASB, IFRS 17 Insurance Contracts (2017, amended 2020); KaribuTech AI analysis.

The three measurement approaches



A good IFRS 17 process is measured by how clearly the board understands the result, not by how many spreadsheets it took to produce it.

Climate, catastrophe and new risks

Actuarial methods rely on the past to predict the future. For climate, cyber and other emerging risks, the past is a weak guide. AI helps by bringing in new data, such as

satellite and weather records, and by generating far more forward-looking scenarios than a team could build by hand.

EXHIBIT 7

Where AI helps actuaries with risks that have little history

RISK	WHY IT IS HARD	WHAT AI ADDS	WHERE IT HELPS
Floods and storms	Short, patchy loss history; changing climate	Satellite, elevation and weather data to map exposure street by street	Property pricing; accumulation limits
Drought and crops	Few claims records for smallholders	Rainfall and vegetation data as parametric triggers	Index-based agriculture cover
Heat and health	Changing disease and heat patterns	Climate scenarios linked to health claims data	Health pricing and reserves
Cyber	Fast-changing threats; little data	Scenarios built from incident data and company profiles	Cyber pricing for business customers
AI and model failure	New liability with almost no history	Scenario analysis of how AI errors turn into losses	Pricing AI liability cover

Source: KaribuTech AI analysis.

Scenarios, not just averages

More scenarios

Generate and run thousands of climate and economic scenarios for capital and ALM work.

Faster runs

Use fast approximate models to test many options, then confirm the chosen ones in full models.

Clear narratives

Turn scenario results into plain-language stories the board can act on.

For new risks, the actuary's value is judgement under uncertainty. AI widens the range of futures you can test; it does not tell you which one will happen.

The 100-day quick-wins plan for actuaries

Do not start with a large programme. Pick one process, such as the quarterly valuation or one product's pricing, deliver a working improvement in 100 days, measure it

against today, then reuse the approach. Each cycle is run by a small team of actuaries, data engineers and finance colleagues.

EXHIBIT 8

One actuarial process in 100 days: design, build, run, measure

DAYS 1-30

Discover and design

- Map the process and time each step
- Agree data sources and checks
- Agree model and sign-off rules

DAYS 31-60

Build

- Automated data checks and loads
- Challenger model or automated run
- Documentation drafted with AI

DAYS 61-100

Run and measure

- Run in parallel with today's process
- Compare results and explain differences
- Report time saved and accuracy

Source: KaribuTech AI quick-wins approach.

Three quick wins we recommend

QUICK WIN 1

Automated data checks

- Checks run before every valuation
- Errors flagged at the source
- Reconciliation to finance automated
- Data-quality report for the board

MEASURE

Days to prepare data, errors found, rework

QUICK WIN 2

AI challenger pricing model

- One product, one line of business
- Runs beside the current model
- New factors tested and explained
- Fairness checked before use

MEASURE

Prediction accuracy, loss ratio by segment

QUICK WIN 3

Automated valuation and commentary

- Valuation run automated end to end
- Analysis of change drafted by AI
- Actuaries review and sign off
- Audit trail kept

MEASURE

Days to close, share of movement explained

The business case for AI in the actuarial function

The value of AI in actuarial work comes from two places: better numbers, which improve pricing and reserving, and freed time, which lets actuaries advise

the business. The example below is illustrative, for a mid-sized insurer; your own baseline will differ.

EXHIBIT 9

Illustrative example: \$150m of premium, a 20-person actuarial team

Measure	Before	After AI	Annual value
Loss ratio on repriced products	65%	64%	≈ \$1.5m
Valuation close	20 working days	8 working days	Faster decisions
Time spent preparing data	40% of team time	15% of team time	5 people freed for analysis
Reserve surprises	Found at year end	Found each quarter	Fewer shocks
Total measurable value			≈ \$1.5m + capacity

Source: Illustrative KaribuTech AI example, not a forecast. Loss ratio: one point on \$150m of premium. Freed time: 25% of a 20-person team, redeployed to pricing, experience analysis and advice.

Where the value comes from

- **Sharper prices:** fewer good risks lost and bad risks underpriced.
- **Earlier warnings:** reserve and experience trends spotted each quarter.
- **Faster closes:** results and explanations to the board sooner.
- **Freed capacity:** actuaries move from data work to advice.

Where the money goes

Data preparation and quality	30–40%
Model and valuation tools	20–25%
Cloud computing	10–15%
Training and new skills	15–20%
Model governance	10%

Typical split of a first-year budget (illustrative)

THREE RULES FOR A CREDIBLE CASE

- 1 Time every step of today's process before you start
- 2 Run new models in parallel before switching
- 3 Report accuracy, time saved and fairness every quarter


What you need underneath: the actuarial engine

The head of actuarial does not need to be a technologist, but should shape the platform the team works on. The key idea is one trusted flow of data into a

shared model library, with every run recorded, rather than a web of spreadsheets passed between people.

EXHIBIT 10

One data flow, one model library, every run recorded



Source: KaribuTech AI analysis.

Build, buy or partner

Buy

Valuation and IFRS 17 tools, cloud computing, catastrophe and scenario models

Build

Your pricing and reserving models, assumptions and data checks, on your own data

Partner

Industry mortality and claims studies; universities; reinsurers' data and models

Ask one question of every technology decision: will the next valuation or pricing review be faster, and easier to explain, because of it?

The actuarial team in the AI era

AI does not replace actuaries; it changes the mix of their work and the skills they need. Less time goes on preparing data and running models, more on testing,

explaining and advising. The head of actuarial leads that shift, and decides which skills to build and which to borrow.

EXHIBIT 11

How the main roles change

ROLE	TODAY	WITH AI
Actuarial analyst	Prepares data and runs spreadsheets	→ Builds and tests models with data and AI tools
Pricing actuary	Updates rating tables once a year	→ Tests new factors and challenger models continuously
Reserving actuary	Rebuilds triangles each quarter	→ Reviews AI estimates and explains differences
Valuation actuary	Runs the valuation by hand	→ Reviews automated runs and AI-drafted commentary
Model validator	Reviews a few actuarial models	→ Reviews every model that affects price, reserves or customers
Head of actuarial	Signs off the numbers	→ Signs off the numbers and governs the insurer's models

Source: KaribuTech AI analysis.

Who you work with

- **Data and AI engineers:** build data flows and model tools.
- **Finance:** reconcile results and share the IFRS 17 close.
- **Underwriting and claims:** use prices and feed back experience.
- **Risk and compliance:** oversee model risk and fair treatment.
- **Reinsurance:** share data and catastrophe models.

Skills to build in the team

Python or R alongside spreadsheets

Machine-learning basics

Explaining models simply

Testing for fairness

Scenario thinking

Risks and guardrails for actuarial AI

AI can make actuarial numbers sharper, but it can also make them wrong with great confidence. These are the

risks a head of actuarial owns, and the controls that manage them.

EXHIBIT 12

Five actuarial AI risks, and the control for each

RISK	WHAT CAN GO WRONG	THE CONTROL
Overfitting	Models that fit past data perfectly and predict the future badly	Test on data the model has not seen; prefer simpler models when results are close
Unfair pricing	Prices that disadvantage groups of customers, directly or through proxies	Fairness tests by gender, age, region and income; remove unfair proxies
Black-box results	Numbers no one can explain to the board or regulator	Explanation tools; champion and challenger; sign-off only when explainable
Poor data	Errors and gaps flowing into prices, reserves and capital	Automated checks; data-quality reporting; reconciliation to finance
Model drift	Models that get worse as the world changes	Monthly monitoring; yearly review; the ability to switch back

Source: KaribuTech AI analysis; IAIS, Application Paper on the supervision of AI (2025).

Eight questions to ask your team

01 Can we explain every price and reserve to the board?

02 How many days of each close go on data preparation?

03 Which models have not been independently validated?

04 Have our pricing models been tested for fairness?

05 How would we know if a model started to drift?

06 Which AI models outside our team affect prices or reserves?


07 Do we run new models in parallel before switching?

08 Which risks have too little history to trust the past?

The actuarial scorecard and 12-month roadmap

EXHIBIT 13


Six numbers to run the actuarial function, plus four for the AI era



Source: KaribuTech AI analysis. Gold: keep within a target band; green: higher is better; red: lower is better.

EXHIBIT 14

Four quarters from first win to an AI-enabled actuarial function



Source: KaribuTech AI analysis.

WHERE TO START ON MONDAY

Time every step of your last valuation, find where the days went, and launch the first 100-day cycle on the step that took the longest.

About the authors



Kennedy Chengeta, PhD

Co-founder, KaribuTech AI

Enterprise Architect and AI Leader · Camunda Chapter Lead, Africa

Leads AI, automation and integration programmes for insurers and banks across Southern Africa.



Justice Muguda

Chief Technology Officer, KaribuTech AI


Cloud, integration and cybersecurity

Leads KaribuTech AI's technology and engineering, from cloud platforms to secure connections between systems.

With contributions from the KaribuTech AI Engineering Team. To discuss this report, or to run a 100-day actuarial diagnostic, contact KaribuTech AI at kaributech.co.za.

References

1. Halley, E. (1693), estimate of the degrees of mortality, Philosophical Transactions of the Royal Society.
2. Ogborn, M. E. (1962), Equitable Assurances.
3. Institute and Faculty of Actuaries, history of the profession (Institute of Actuaries founded 1848).
4. IASB (2017, amended 2020), IFRS 17 Insurance Contracts.
5. IAIS (2025), Application Paper on the supervision of artificial intelligence, July 2025.
6. European Union (2024), Regulation (EU) 2024/1689, the Artificial Intelligence Act.
7. McKinsey & Company, research on the value of AI in insurance.
8. Stanford HAI (2025), Artificial Intelligence Index Report 2025.
9. NAIC (2023), Model Bulletin: Use of Artificial Intelligence Systems by Insurers.
10. KaribuTech AI, BluFountain Insure platform (actuarial and valuation engines).
11. KaribuTech AI, The AI Blueprint for the Insurance CEO (2026).
12. KaribuTech AI, quick-wins approach (proprietary).



About KaribuTech AI

KaribuTech AI is a South African technology company specialising in artificial intelligence, machine learning, intelligent automation and integration for insurers and banks across Africa. Our teams design and build intelligent workflows, decision engines, document and computer-vision AI, and cloud-native core platforms, including the **BluFountain Insure** platform: an insurance AI

marketplace with core systems for life, short-term, medical, funeral and investment business. We combine deep process expertise (Camunda, IBM BPM, MuleSoft, webMethods and document automation platforms) with practical AI engineering, so that AI moves out of the lab and into the operations that customers feel.

Disclaimer: this report provides general information and does not constitute legal, actuarial or investment advice. Third-party figures are cited beneath each exhibit and have not been independently verified by KaribuTech AI. Illustrative figures are not forecasts.



KaribuTech AI

INSURANCE · LEADERSHIP SERIES · SEPTEMBER 2026

The AI Blueprint for the Head of Risk

Seeing risk earlier, governing AI safely and keeping the insurer inside its risk appetite


By Kennedy Chengeta, PhD, and Justice Muguda
with the KaribuTech AI Engineering Team

kaributech.co.za



Contents

01	Introduction	166
02	Why risk management must change	167
03	Lessons from history	168
04	Where the risk cycle breaks	169
05	Six moves for the Head of Risk	170
06	Tiering AI by risk	172
07	Early warning: risk indicators with AI	173
08	Generative AI and AI agents for risk teams	174
09	The AI rulebook for insurers	175
10	Cyber, fraud and operational resilience	176
11	The 100-day quick-wins plan	177
12	The business case	178
13	What you need underneath	179
14	The risk team	180
15	Risks of using AI in risk management	181
16	The scorecard and 12-month roadmap	182
17	About the authors and KaribuTech AI	183



About this report

This report is written for chief risk officers, heads of risk and compliance, and their teams at insurers, reinsurers and pension funds in developing markets.

It shows how AI can help risk teams see problems earlier and test more scenarios, and how to govern AI itself as a risk: the pressures on the function, six moves, tiering AI by risk, early-warning indicators, the regulatory rulebook, cyber resilience, a 100-day plan, the business case, the team, the pitfalls and a scorecard.

External data are sourced beneath each exhibit, and a full list of references appears at the end. Frameworks and illustrative figures without an external source are the authors' own.

WRITTEN FOR

Chief risk officers

Compliance

Model risk

Boards

Introduction

The head of risk now has two jobs: using AI to see risk earlier, and governing AI as a risk in its own right.

Risk teams still spend much of their time collecting data, chasing control owners and building reports that describe last quarter. Meanwhile, AI is spreading through claims, underwriting and marketing, often faster than the risk framework can follow.

AI can help: watching risk indicators continuously, testing thousands of scenarios, reading regulations and incident reports, and flagging when a control stops working. The same technology brings new risks that the board expects the risk function to own.

This report sets out **six moves** for the head of risk, a way to **tier AI by risk**, **early-warning indicators**, the **AI rulebook**, **cyber resilience**, a **100-day plan**, the business case and a **scorecard**, in plain language.

THREE QUESTIONS THIS REPORT ANSWERS

01

Which AI models are in use, and how risky are they?

02

How do we see problems before they hit the results?

03

What will supervisors expect of us on AI?

AT A GLANCE: SIX MOVES FOR THE RISK FUNCTION

- 1 Build a live view of risk
- 2 Put every AI model on the register
- 3 Stress-test with more scenarios
- 4 Strengthen cyber and operational resilience
- 5 Make compliance continuous
- 6 Govern third parties and data

Why risk management must change

Five pressures are reshaping the risk function. Each one is also a chance for risk leaders to

move from reporting the past to shaping decisions.

AI is spreading faster than controls

Models appear in claims, pricing and marketing, sometimes bought as part of software, before the risk team knows.

Regulators are setting AI rules

Supervisors expect governance, explainability and fair treatment for AI decisions, with deadlines approaching.

Cyber threats are growing

Breaches are costly, and AI makes phishing, deepfakes and attacks cheaper and more convincing.

Climate risk is changing the picture

Floods, droughts and heat affect claims, investments and operations at the same time.

Reports arrive too late

Quarterly risk reports describe what has already happened. Boards want to know what is about to happen.

EXHIBIT 1

The numbers behind the pressure

\$4.44m

global average cost of a data breach (2025)

1 in 5

organisations had a breach involving unapproved "shadow" AI

97%

of organisations with AI breaches lacked proper AI access controls

Dec 2027

EU high-risk AI rules apply, after the Digital Omnibus delay

Source: IBM, Cost of a Data Breach Report 2025; European Union, AI Act and Digital Omnibus (2024, 2026).

The best risk functions do not slow AI down. They make it safe enough for the business to go faster.

Lessons from history: when risk was missed

Some of the costliest failures in financial history came from risks that were visible in the data but not in the

reports, or from technology that moved faster than its controls. Each holds a lesson for AI.

1720

The South Sea Bubble

A crash in South Sea Company shares ruined many investors. The Bubble Act of 1720 followed, chartering the Royal Exchange and London Assurance to write marine insurance.

Today: AI rules are arriving; build controls before they are required.

THE LESSON

Rules follow crises. Better to build controls before the regulator has to.

2008

AIG and the credit crisis

AIG, then one of the world's largest insurers, needed a government rescue after losses on credit-default swaps written outside its core insurance business. The risk was concentrated where oversight was weakest.

Today: shadow AI and vendor AI sit outside the risk team's view.

THE LESSON

Risk hides in the parts of the business the risk function sees least.

2012

Knight Capital's 45 minutes

A faulty software release at the US trading firm Knight Capital sent millions of unintended orders in about 45 minutes, causing a loss of around \$440 million and the firm's sale.

Today: AI decisions repeat thousands of times, so monitoring must be live.

THE LESSON

Automated systems fail fast. Controls must work at machine speed.

Today

AI across the insurer

AI now makes or shapes decisions on price, cover and claims. Its errors can repeat thousands of times before anyone notices.

Today: register, tier, test and monitor every AI model.

THE LESSON

Govern AI like any material risk: inventory, test, monitor, report.

Source: Histories of the South Sea Bubble and the Bubble Act (1720); US government and AIG disclosures (2008); US Securities and Exchange Commission, Knight Capital (2013).

The common thread is not bad luck. It is risk that grew faster than the controls around it. AI makes that gap easier to open, and easier to close.

Where the risk cycle breaks, and how AI helps

Risk management follows a cycle, from identifying risks to learning from what went wrong. At each step, information arrives late or incomplete. The exhibit

shows where the cycle typically breaks and what AI does.

EXHIBIT 2

The risk cycle: eight steps from identifying risk to learning from it

	WHERE IT BREAKS TODAY	WHAT AI DOES
Identify Spot new and emerging risks	Annual workshops; risks found late	Scans news, claims and regulation for new risks
Assess Size likelihood and impact	Judgement with little data	Estimates from internal and external data
Set appetite Agree limits with the board	Limits rarely linked to live data	Limits tracked automatically against data
Monitor Track risk indicators	Quarterly spreadsheets	Live dashboards and early alerts
Stress test Test severe scenarios	A handful of scenarios a year	Thousands of scenarios, run quickly
Respond Act on breaches and incidents	Slow escalation by email	Automatic escalation with deadlines
Report Inform executives and board	Weeks to compile	Reports drafted from live data
Learn Fix root causes	Lessons lost between incidents	Patterns found across incidents

Source: KaribuTech AI analysis.

The goal is not more reports. It is earlier warnings, faster action and fewer surprises.

Six moves for the Head of Risk

Risk leaders who use AI well do six things differently. Each is about seeing risk sooner

and governing it better, not about choosing a technology.

MOVE 1

01

Build a live view of risk

Most risk reports describe the past. AI can pull data from every system, track key risk indicators continuously and warn when a limit is about to be breached, so the risk team talks about what is coming, not what happened.

What leaders do

- Link key risk indicators to live data from finance, claims, IT and HR.
- Set early-warning thresholds below every appetite limit.
- Give executives one live risk dashboard, not ten reports.

ASK YOUR TEAM

Which of our risk limits could be breached this month without anyone noticing?

MOVE 2

02

Put every AI model on the register

AI is arriving through many doors: projects, software updates and staff using public tools. The risk function should know every model in use, what it decides and how risky it is.

What leaders do

- Create one register of every AI model and tool, including those inside vendor software.
- Tier each model by its impact on customers and the balance sheet.
- Set controls by tier: the higher the risk, the stronger the testing and oversight.

ASK YOUR TEAM

How many AI models are in use in our business today, and who owns each one?

MOVE 3

03

Stress-test with more scenarios

A few annual stress tests cannot capture climate, cyber, economic and pandemic risks together. AI can generate and run many more scenarios, and show which combinations hurt most.

What leaders do

- Run many combined scenarios: climate plus economic, cyber plus operational.
- Use fast approximate models to screen, then full models to confirm.
- Turn results into plain stories the board can act on.

ASK YOUR TEAM

Which combination of events would hurt our solvency most, and how likely is it?

MOVE 4

04

Strengthen cyber and operational resilience

Cyber attacks and system failures are among the biggest risks for insurers, and AI makes attacks cheaper. AI also helps defend: spotting unusual activity, fake identities and failing controls faster than people can.

What leaders do

- Use AI to detect unusual activity across systems and users.
- Test staff against AI-generated phishing and voice fraud.
- Rehearse recovery from a major outage or breach every year.

ASK YOUR TEAM

How many hours would it take us to notice an intruder in our systems?

MOVE 5

05

Make compliance continuous

Compliance checks are often sample-based and periodic. AI can check every transaction, screen every customer and track every regulatory change, so issues are found in days, not at the next audit.

What leaders do

- Screen every customer and payment for sanctions and money-laundering risk.
- Track regulatory changes and map them to policies and controls.
- Test key controls automatically and continuously.

ASK YOUR TEAM

What share of our key controls are tested automatically today?

MOVE 6

06

Govern third parties and data

Much of an insurer's AI and data now sits with cloud providers, software vendors and partners. The risk function must know where data goes and what happens if a key provider fails.

What leaders do

- Review every vendor that uses AI or holds customer data, before and during the contract.
- Keep sensitive data in-country where regulation requires it.
- Have an exit plan for every critical provider.

ASK YOUR TEAM

Which single provider could stop our business if it failed tomorrow?

FROM OUR WORK

Compliance built into the platform

KaribuTech AI's BluFountain Insure platform builds risk and compliance into everyday processes: customer-checking (KYC) and anti-money-laundering (AML) engines screen every customer, AI agents help detect fraud, and IFRS 17 and IFRS 9 compliance is built in, with ISO certification in progress. Every step is recorded, so the risk team can see controls working rather than sampling them after the event.

Every customer

screened by KYC and AML engines

Every step

recorded for audit and review

Built in

IFRS 17 and IFRS 9 compliant


Tiering AI by risk: controls that match the stakes

Not every AI use case needs the same level of control. A drafting assistant is not a pricing model. Sorting AI into three tiers lets the business move quickly on

low-risk uses while the strongest controls go where customers and capital are affected.

EXHIBIT 3

Three tiers of AI risk, with controls that rise with the stakes



LIGHTER CONTROLS →

← STRONGER CONTROLS

Source: KaribuTech AI analysis.

How to decide the tier

- **Who is affected:** staff only, or customers?
- **What is decided:** advice, or a final decision?
- **How much is at stake:** money, cover or rights?
- **How reversible:** can a mistake be undone quickly?

Controls every tier needs

- **A named owner** for every model.
- **An entry in the register** with its tier.
- **Clear data rules:** what data it may use.
- **A way to switch it off** if it goes wrong.

Early warning: risk indicators with AI

Key risk indicators only help if they arrive in time to act. AI can watch many more signals, spot unusual movements and warn before a limit is breached. The

table shows what AI can monitor for each main risk type.

EXHIBIT 4

What AI can watch for each risk type

RISK TYPE	WHAT AI MONITORS	EARLY WARNING	ACTION
Insurance risk	Claims frequency and size by product and region; lapse rates	Claims rising faster than priced	Review pricing and reinsurance
Market and liquidity	Asset prices, interest rates, cash flows	Liquidity buffer falling	Adjust investments or funding
Credit and counterparty	Reinsurer ratings, broker balances, investment issuers	Counterparty downgraded or late	Reduce exposure; call collateral
Operational	System outages, process delays, errors and complaints	Backlogs and error rates rising	Escalate to process owners
Cyber	Unusual logins, data movement, phishing attempts	Suspicious activity detected	Isolate and investigate
Conduct	Declines, complaints and outcomes by customer group	Unfair patterns emerging	Review rules and models
Climate	Weather forecasts, exposure maps, catastrophe events	Severe weather approaching insured areas	Warn customers; prepare claims surge

Source: KaribuTech AI analysis.

Three rules for useful indicators

01

Fewer, better indicators

Track the ten that predict trouble, not the hundred that describe it.

02

Thresholds before limits

Warn at amber, well before the red limit is breached.

03

An owner for every alert

Each warning goes to a named person with a deadline to respond.

A risk indicator that arrives after the loss is a history lesson. The value of AI is in the warning that arrives in time to act.


Generative AI and AI agents for risk teams

Generative AI reads and writes; AI agents plan and carry out tasks. For risk and compliance teams, that means help with the reading and writing that fills much

of the week: regulations, policies, incident reports, control tests and board papers.

EXHIBIT 5

Six jobs AI can take on in the risk team today



Source: KaribuTech AI analysis. Labels show the level of AI autonomy we recommend.

What stays with the risk team

- Setting risk appetite with the board.
- Judging whether a risk is acceptable.
- Challenging the business and escalating.
- Signing off reports and attestations.

Guardrails that make it safe

- **Approved sources only:** AI works from the insurer's own documents and data.
- **No confidential data in public AI tools.**
- **Every draft reviewed** before it reaches the board.
- **Every automated test logged** for audit.

WHERE TO START

Start with regulatory-change scanning and incident summaries. They save many hours each month and carry little risk.

The AI rulebook for insurers is taking shape

Rules for AI in insurance are converging around the same principles: govern, explain, be fair and stay

secure. A risk function that builds to these principles now will be ready for whichever national rules follow.

EXHIBIT 6

Key milestones for AI regulation in insurance

WHEN	WHAT	WHY IT MATTERS FOR INSURERS
Dec 2023	NAIC model bulletin on insurers' use of AI (US)	Sets expectations for AI governance, testing and third-party oversight; adopted by many US states
Jul 2024	African Union Continental AI Strategy	Signals the direction for national AI rules across Africa
Aug 2024	EU AI Act enters into force	Classes life and health risk pricing as high-risk AI
Mar 2025	Kenya National AI Strategy 2025–2030	An example of national AI plans in our markets
Jul 2025	IAIS Application Paper on the supervision of AI	Guidance for insurance supervisors worldwide, built on the Insurance Core Principles
Dec 2027	EU high-risk AI rules apply	Deferred from August 2026 by the Digital Omnibus

Source: NAIC (2023); African Union (2024); EU AI Act and Digital Omnibus (2024, 2026); Government of Kenya (2025); IAIS (2025).

What supervisors expect, in plain language

01

Governance

A named owner, board oversight and a register of AI models.

02

Fairness and explanation

Decisions that can be explained and do not treat groups unfairly.

03

Security and resilience

Protected data, robust models and a plan when things go wrong.

Build once to the common principles, then map each new national rule to the controls you already have.

Cyber, fraud and operational resilience

AI makes attacks cheaper and more convincing, from AI-written phishing to cloned voices. It also gives

defenders new tools. The table shows the main threats and the controls that work against them.

EXHIBIT 7

Five threats, and the controls that work

THREAT	WHAT HAPPENS	THE CONTROL
AI-written phishing	Convincing emails and messages trick staff into giving access	Regular phishing tests; strong sign-in checks; AI filtering of messages
Cloned voices and deepfakes	Fake calls from executives or customers authorise payments or changes	Call-back rules for payments; code words; liveness checks
Shadow AI	Staff paste customer data into unapproved AI tools	Approved tools; data-loss monitoring; clear rules and training
Attacks on AI models	Models manipulated or data stolen through AI systems	Access controls on AI; testing; monitoring of AI use
Major outage	Cloud, core system or provider failure stops operations	Tested recovery plans; backups; exit plans for key providers

Source: IBM, Cost of a Data Breach Report 2025; KaribuTech AI analysis.

What good looks like

01

Detect in hours

Know about an intrusion in hours, not months.

02

Recover in days

Restore critical services within agreed times.

03

Learn every time

Every incident leads to a fixed root cause.

Resilience is not avoiding every incident. It is noticing quickly, limiting the damage and recovering fast.

The 100-day quick-wins plan for risk

Do not start with a large programme. Pick one area, such as the AI register or one set of risk indicators, deliver a working improvement in 100 days, measure it

against today, then build on it. Each cycle is run by a small team from risk, compliance, IT and the business.

EXHIBIT 8

One risk improvement in 100 days: design, build, run, measure

DAYS 1-30

Discover and design

- Map today's process and data
- Agree indicators, tiers or controls
- Agree owners and escalation

DAYS 31-60

Build

- Connect data sources
- Build dashboards, register or tests
- Set thresholds and alerts

DAYS 61-100

Run and measure

- Run beside today's process
- Review alerts with owners
- Report results to the risk committee

Source: KaribuTech AI quick-wins approach.

Three quick wins we recommend

QUICK WIN 1

AI register and tiering

- Every AI model and tool listed
- Owner and tier for each
- Controls set by tier
- Shadow AI identified

MEASURE

Share of models registered, share tested

QUICK WIN 2

Live risk dashboard

- Ten key indicators on live data
- Amber and red thresholds
- Alerts with owners and deadlines
- Board view in one page

MEASURE

Warnings before breaches, report time

QUICK WIN 3

Continuous control testing

- Key controls tested automatically
- Exceptions sent to owners
- Regulatory changes tracked
- Audit trail kept

MEASURE

Share of controls tested, findings closed on time

The business case for AI in risk management

The value of AI in risk comes from losses avoided, capital used more wisely and time freed for judgement. Losses avoided are hard to prove, so a credible case

combines measurable efficiency with a clear view of risk reduced. The example below is illustrative.

EXHIBIT 9

Illustrative example: a mid-sized insurer with a 15-person risk and compliance team

Measure	Before	After AI	Annual value
Operational losses	\$3.0m a year	\$2.4m a year	≈ \$0.6m
Time to produce risk reports	15 working days	5 working days	Faster decisions
Controls tested automatically	10%	60%	Fewer audit findings
Chance of a major breach in a year	10%	7%	≈ \$0.13m expected
Total measurable value			≈ \$0.7m + time freed

Source: Illustrative KaribuTech AI example, not a forecast. Breach value: three-point lower chance of a breach costing about \$4.44m, the 2025 global average (IBM).

Where the value comes from

- **Fewer losses:** problems caught before they become incidents.
- **Faster reporting:** the board sees risk sooner.
- **Stronger controls:** tested continuously, not by sample.
- **Freed time:** risk staff move from reporting to advising.

Where the money goes

Risk data and dashboards	30–40%
Control-testing tools	20–25%
AI register and model testing	15–20%
Cyber tools	10–15%
Training	10%

Typical split of a first-year budget (illustrative)

THREE RULES FOR A CREDIBLE CASE

- 1 Baseline losses, findings and report times first
- 2 Count only benefits the CRO and CFO sign
- 3 Report value and incidents every quarter

What you need underneath: the risk engine

The head of risk does not need to be a technologist, but should know what the engine looks like. The key idea is

one risk data hub fed by every system, so indicators, registers and reports all use the same live data.


EXHIBIT 10

One risk data hub, fed by every system



Source: KaribuTech AI analysis.

Build, buy or partner



Ask one question of every technology decision: will we see the next problem sooner, and act on it faster, because of it?

The risk team in the AI era

AI changes the risk function from a producer of reports into a live early-warning and governance team. Less time goes on collecting data and sampling controls;

more on challenge, advice and oversight of AI across the insurer.

EXHIBIT 11

How the main roles change

ROLE	TODAY	WITH AI
Risk analyst	Collects data and builds reports	→ Interprets live indicators and investigates alerts
Operational risk manager	Runs annual risk workshops	→ Tracks incidents and indicators continuously
Model risk validator	Reviews a few actuarial models	→ Reviews every tier 2 and tier 3 AI model
Compliance officer	Tests samples each quarter	→ Oversees continuous testing and follows up exceptions
Cyber risk lead	Reviews security reports	→ Works with IT on live threat detection and response
Chief risk officer	Reports last quarter's risks	→ Advises on what is coming and governs AI across the firm

Source: KaribuTech AI analysis.

Who you work with

- **Board and risk committee:** agree appetite and AI oversight.
- **Actuarial:** capital, reserving and model validation.
- **IT and security:** cyber defence and resilience.
- **Business heads:** own risks, controls and AI use cases.
- **Internal audit:** independent assurance.

Skills to build in the team

Reading AI and data outputs

Model validation basics

Scenario thinking

Clear board communication

Regulatory interpretation

Risks of using AI in risk management

AI can make the risk function sharper, but it brings its own traps. A dashboard can create false comfort; too

many alerts can hide the one that matters. These are the pitfalls to manage.

EXHIBIT 12

Five pitfalls of AI in risk management, and the control for each

RISK	WHAT CAN GO WRONG	THE CONTROL
False comfort	Green dashboards hide risks the data does not capture	Keep judgement and challenge; review what indicators miss
Alert fatigue	Too many alerts, so real warnings are ignored	Fewer indicators; tuned thresholds; owners for every alert
Risk models that fail	Early-warning or scenario models that are wrong	Validate risk models like any other; test against real events
Data exposure	Sensitive data gathered in one place becomes a target	Access controls; encryption; in-country data where required
Over-automation	Decisions made by rules no one reviews	Human review of escalations; regular rule reviews

Source: KaribuTech AI analysis; IAIS, Application Paper on the supervision of AI (2025).


Eight questions to ask your team

- Which risk limits could we breach without an early warning?
- How many AI models are not on our register?
- Which controls are still tested only by sample?
- How quickly would we detect an intruder?
- Which vendors hold our customer data, and where?
- Are our tier 3 models independently validated?
- How many alerts did we ignore last month, and why?
- Could we explain any AI decision to the regulator tomorrow?

The risk scorecard and 12-month roadmap

EXHIBIT 13


Six numbers to run the risk function, plus four for the AI era



Source: KaribuTech AI analysis. Gold: keep within a target band; green: higher is better; red: lower is better.

EXHIBIT 14

Four quarters to an AI-ready risk function



Source: KaribuTech AI analysis.

WHERE TO START ON MONDAY

Ask for a list of every AI model and tool in use across the business. Whatever the number is, it will be larger than expected, and it is where the first 100 days should start.

About the authors



Kennedy Chengeta, PhD

Co-founder, KaribuTech AI

Enterprise Architect and AI Leader · Camunda Chapter Lead, Africa

Leads AI, automation and integration programmes for insurers and banks across Southern Africa.



Justice Muguda

Chief Technology Officer, KaribuTech AI


Cloud, integration and cybersecurity

Leads KaribuTech AI's technology and engineering, from cloud platforms to secure connections between systems.

With contributions from the KaribuTech AI Engineering Team. To discuss this report, or to run a 100-day risk and AI governance diagnostic, contact KaribuTech AI at kaributech.co.za.

References

1. IBM (2025), Cost of a Data Breach Report 2025.
2. IAIS (2025), Application Paper on the supervision of artificial intelligence, July 2025.
3. European Union (2024), Regulation (EU) 2024/1689, the Artificial Intelligence Act; Digital Omnibus (2026).
4. NAIC (2023), Model Bulletin: Use of Artificial Intelligence Systems by Insurers.
5. African Union (2024), Continental Artificial Intelligence Strategy.
6. Government of Kenya (2025), Kenya Artificial Intelligence Strategy 2025–2030.
7. Histories of the South Sea Bubble and the Bubble Act (1720).
8. US government and AIG disclosures on the 2008 rescue.
9. US Securities and Exchange Commission (2013), order concerning Knight Capital Americas.
10. KaribuTech AI, BluFountain Insure platform (KYC and AML engines).
11. KaribuTech AI, The AI Blueprint for the Insurance CEO (2026).
12. KaribuTech AI, quick-wins approach (proprietary).



About KaribuTech AI

KaribuTech AI is a South African technology company specialising in artificial intelligence, machine learning, intelligent automation and integration for insurers and banks across Africa. Our teams design and build intelligent workflows, decision engines, document and computer-vision AI, and cloud-native core platforms, including the **BluFountain Insure** platform: an insurance AI

marketplace with core systems for life, short-term, medical, funeral and investment business. We combine deep process expertise (Camunda, IBM BPM, MuleSoft, webMethods and document automation platforms) with practical AI engineering, so that AI moves out of the lab and into the operations that customers feel.

Disclaimer: this report provides general information and does not constitute legal, actuarial or investment advice. Third-party figures are cited beneath each exhibit and have not been independently verified by KaribuTech AI. Illustrative figures are not forecasts.



KaribuTech AI

INSURANCE · LEADERSHIP SERIES · SEPTEMBER 2026

The AI Blueprint for the Insurance Broker

Better advice, faster placement and stronger client relationships: AI for brokers and agencies


By Kennedy Chengeta, PhD, and Justice Muguda
with the KaribuTech AI Engineering Team

kaributech.co.za



Contents

01	Introduction	186
02	Why broking must change	187
03	Lessons from history	188
04	Where the broking cycle breaks	189
05	Six moves for the insurance broker	190
06	Three client lanes	192
07	AI across broking tasks	193
08	Generative AI and AI agents for brokers	194
09	Advice, conduct and the client's interest	195
10	Working with insurers and partners	196
11	The 100-day quick-wins plan	197
12	The business case	198
13	What you need underneath	199
14	The broking team	200
15	Risks and guardrails	201
16	The scorecard and 12-month roadmap	202
17	About the authors and KaribuTech AI	203



About this report

This report is written for broker principals, managing directors and account executives at brokerages and agencies in developing markets.

Brokers are caught between insurers going direct, platforms embedding cover and clients who expect instant answers. This report shows how AI can make brokers more valuable, not less: the pressures, six moves, client lanes, AI across broking tasks, conduct and the client's interest, working with insurers, a 100-day plan, the business case, the team, the risks and a scorecard.

External data are sourced beneath each exhibit, and a full list of references appears at the end. Frameworks and illustrative figures without an external source are the authors' own.

WRITTEN FOR

Broker principals

Account executives

Placement teams

Agencies

Introduction

Brokers win when they know the client and the market better than anyone else. AI can give every broker that edge, or give it to the platforms that want to replace them.

Much of a broker's week still goes on retyping client details into insurer forms, chasing quotes, comparing long policy wordings and preparing renewal packs by hand. That leaves too little time for what clients actually pay for: advice, negotiation and support when they claim.

AI can read submissions, match risks to insurers' appetite, compare quotes and wordings, warn when a client may not renew and prepare the file for a claim, so brokers spend their time advising.

This report sets out **six moves** for the insurance broker, **three client lanes**, **AI across broking tasks**, **advice and conduct**, **working with insurers**, a **100-day plan** and a **scorecard**, in plain language.

THREE QUESTIONS THIS REPORT ANSWERS

01

How do we spend more time advising and less on admin?

02

How do we compare quotes and wordings without errors?

03

How do we keep clients at renewal and win new ones?

AT A GLANCE: SIX MOVES FOR BROKERS

1

Advise, don't just transact

2

Automate submissions and placement

3

Compare quotes and wordings with AI

4

Win renewals early

5

Be the client's claims advocate

6

Reach new clients digitally and through partners

Why broking must change

Five pressures are squeezing the traditional broking model. Each is also a chance for

brokers who use AI well to become more valuable to clients and insurers.

Insurers and platforms go direct

Simple cover is sold online, in bank apps and at checkouts, often without a broker.

Clients expect instant answers

Business owners used to digital banking expect quotes and advice in hours, not weeks.

Admin eats the margin

Submissions, rekeying, chasing and renewal packs take time that clients will not pay for.

Conduct rules are tightening

Advice, suitability, disclosure and commission rules demand better records and fairer outcomes.

Errors are expensive

A missed exclusion or wrong sum insured can mean a professional indemnity claim against the broker.

EXHIBIT 1

Why speed and advice now decide who wins

10%

of submissions some underwriters manage to answer

40%

of underwriters' time on non-core and admin work

\$2tn+

moved through mobile money in 2025, a channel for new clients

280x

cheaper to run capable AI in 2024 than in 2022

Source: Accenture underwriting survey, as reported by Insurance Business; Capgemini; GSMA (2026); Stanford HAI, AI Index 2025.

Clients will always pay for good advice. They will not pay for retyping. AI lets brokers do more of the first and less of the second.

Lessons from history: broking has always been about information

Brokers have always earned their place by knowing the market and carrying the client's risk to the right

underwriter. Each change in how information moves has reshaped the broker's role.

1688

Lloyd's coffee house

Brokers carried their clients' risks from underwriter to underwriter in Edward Lloyd's coffee house, where the best shipping news was found. The broker's value was knowing who would write what.

Today: AI matches each risk to the insurers most likely to write it.

THE LESSON

Brokers win by knowing the market better than anyone.

2002

Comparison websites

Online comparison sites such as Confused.com, launched in the UK in 2002, let customers compare simple cover themselves. Brokers who relied on simple personal lines lost ground.

Today: brokers compete on advice for complex risks, not simple cover.

THE LESSON

Simple products go direct. Advice on complex risks does not.

2007

Mobile money

M-Pesa in Kenya made it possible to sell and collect premiums on a phone, opening mass-market distribution through telcos and agents.

Today: WhatsApp and mobile channels bring brokers new clients.

THE LESSON

New channels create new clients for those who use them.

Today

AI as the broker's assistant

AI can prepare submissions, compare markets and wordings and track renewals, freeing the broker to advise.

Today: AI does the admin; brokers spend the time with clients.

THE LESSON

Use AI to do the admin, and put the time into the client.

Source: Lloyd's of London, history; Confused.com (2002); Safaricom (2007).

Every time information became easier to get, brokers who only moved paper lost out, and brokers who gave real advice grew. AI is the biggest version of that shift yet.

Where the broking cycle breaks, and how AI helps

Every client relationship follows the same cycle, from first contact to renewal and claims. At each step, time,

accuracy or the client's trust can be lost. The exhibit shows where it typically breaks and what AI does.

EXHIBIT 2

The broking cycle: eight steps from prospect to renewal and claims

	WHERE IT BREAKS TODAY	WHAT AI DOES
Prospect Finding and qualifying clients	Cold lists; slow follow-up	Leads scored; follow-ups prompted
Needs analysis Understanding the client's risks	Long questionnaires	Pre-filled from public and client data
Submission Preparing the risk for insurers	Rekeyed into every insurer's form	Prepared once, sent in each format
Placement Getting quotes from markets	Sent to insurers who decline	Matched to insurers' appetite
Compare and advise Recommending the best option	Wordings compared by hand	Quotes and wordings compared side by side
Bind and documents Cover confirmed and issued	Errors in schedules and documents	Documents checked against the quote
Service Changes during the year	Slow endorsements	Requests handled and tracked
Renewal and claims Keeping and supporting the client	Renewals left late; claims unsupported	Early renewal alerts; claims files prepared

Source: KaribuTech AI analysis.

The broker's value lives in steps two, five and eight. AI should take over the rest.

Six moves for the insurance broker

Brokers who use AI well do six things differently. Each is about more advice and less

admin, not about choosing a technology.

MOVE 1

01

Advise, don't just transact

Clients can buy simple cover without a broker. What they cannot easily get elsewhere is a trusted adviser who understands their business and risks. AI frees time for that.

What leaders do

- Use AI to take over rekeying, chasing and document preparation.
- Spend the time saved on risk reviews and client meetings.
- Build advice into every renewal, not just the price.

ASK YOUR TEAM

How many hours a week do our account executives spend advising, rather than administering?

MOVE 2

02

Automate submissions and placement

Preparing the same risk in different insurer formats, and sending it to insurers who will decline, wastes days. AI can prepare one submission, reformat it for each market and target insurers whose appetite fits.

What leaders do

- Capture client data once and reuse it across insurers.
- Match risks to insurers' published appetite before sending.
- Track every quote request and chase automatically.

ASK YOUR TEAM

How many of last month's submissions went to insurers who were never going to quote?

MOVE 3

03

Compare quotes and wordings with AI

The biggest broker errors hide in policy wordings: exclusions, limits and conditions that differ between insurers. AI can compare wordings line by line and highlight what matters to the client.

What leaders do

- Compare every quote's wording, not just its price.
- Highlight exclusions and limits that matter for this client.
- Keep a record of the comparison for the file.

ASK YOUR TEAM

Could we show a client exactly how the wordings we compared differ?

MOVE 4

04

Win renewals early

Clients leave when renewals arrive late, with a price rise and no explanation. AI can flag clients at risk months ahead, prepare renewal packs and prompt a review conversation.

What leaders do

- Score every client for renewal risk, 90 days out.
- Prepare renewal packs automatically, with changes explained.
- Contact at-risk clients early with a review, not just a quote.

ASK YOUR TEAM

Which clients are most likely to leave at renewal, and have we spoken to them?

MOVE 5

05

Be the client's claims advocate

The moment clients value their broker most is when they claim. AI can prepare the claim file, check the policy for cover and track the insurer's response, so the broker can push for a fair outcome.

What leaders do

- Prepare claim files with the right documents and policy wording.
- Track every claim's progress with each insurer.
- Report to clients proactively, not only when they ask.

ASK YOUR TEAM

How many of our clients' claims last year were paid in full, and how fast?

MOVE 6

06

Reach new clients digitally and through partners

Brokers can use the same digital channels that threaten them: WhatsApp, online quote tools and partnerships with banks, associations and platforms, with AI handling first contact.

What leaders do

- Offer quotes and advice booking on WhatsApp and the web.
- Partner with associations, banks and platforms serving your target clients.
- Use AI to qualify leads and book adviser time.

ASK YOUR TEAM

Where do our best future clients already gather, online and offline?

HOW WE HELP

Self-service portals for brokers

KaribuTech AI's BluFountain Insure platform gives brokers their own self-service portal alongside portals for merchants, employers, doctors, underwriters and assessors. Brokers can quote, submit, track policies and claims and reconcile commission online, with WhatsApp and document engines behind the scenes, so less time goes on rekeying and chasing.

Broker portal

quote, submit and track online

Shared engines

documents, WhatsApp and KYC

Less rekeying

client data captured once


Three client lanes: the right service for each client

Not every client needs the same service. Sorting clients into three lanes lets brokers serve simple needs digitally

and profitably, and spend expert time where advice matters most.

EXHIBIT 3

Three client lanes, from digital self-service to expert advice



MORE DIGITAL →

← MORE EXPERT ADVICE

Source: KaribuTech AI analysis.

What decides the lane

- **Complexity:** standard cover or bespoke terms?
- **Premium size:** can the client support adviser time?
- **Client preference:** digital or face to face?
- **Risk:** how much could go wrong for the client?

Rules for every lane

- **Suitable advice:** cover must fit the client's needs.
- **A person on request:** no client trapped in a chatbot.
- **Clear records:** what was advised, and why.
- **Fair commission:** disclosed and consistent.

AI across broking tasks

AI helps most with the tasks that are repetitive, document-heavy and error-prone. The table shows

seven common broking tasks, what AI does and the benefit to the brokerage and the client.

EXHIBIT 4

Seven broking tasks, and what AI changes

TASK	WHAT AI DOES	BENEFIT
Submissions	Reads client data and documents; prepares each insurer's format	Hours saved per placement
Market matching	Matches risks to insurers' appetite and past quotes	Fewer declined submissions
Quote comparison	Puts prices, limits and deductibles side by side	Faster, clearer advice
Wording checks	Compares policy wordings and highlights differences	Fewer errors and E&O claims
Renewals	Flags at-risk clients; prepares renewal packs	Higher retention
Claims support	Prepares claim files; tracks insurer progress	Stronger client loyalty
Compliance	Checks KYC, advice records and disclosures	Cleaner audits

Source: KaribuTech AI analysis.

Three rules for AI in broking

01

The broker advises

AI prepares; the broker recommends and signs off.

02

Check the wording

Never bind on an AI comparison without review.

03

Keep the record

Every AI-assisted recommendation is documented.

Generative AI and AI agents for brokers

Generative AI reads and writes; AI agents plan and carry out tasks. For brokers, that means drafting proposals, comparing wordings, summarising client

risks, answering routine questions and preparing renewal and claim files. Advice and recommendations stay with the broker.

EXHIBIT 5

Six jobs AI can take on in a brokerage today

Draft proposals Prepares client proposals and cover summaries for the broker to review RECOMMEND	Compare wordings Compares policy wordings and flags differences in cover RECOMMEND	Summarise client risks Turns surveys, financials and past claims into a risk summary ASSIST
Answer routine questions Answers clients' questions about their cover and documents, any time ACT WITHIN LIMITS	Prepare renewals Builds renewal packs with changes in price and cover explained RECOMMEND	Chase insurers Follows up quotes, endorsements and claims with insurers ACT WITHIN LIMITS

Source: KaribuTech AI analysis. Labels show the level of AI autonomy we recommend.

What stays with the broker

- Every recommendation to a client.
- Negotiation with insurers.
- Binding cover and signing off documents.
- Support for clients in a claim or dispute.

Guardrails that make it safe

- **Approved tools only:** no client data in public AI tools.
- **Every comparison checked** before advice is given.
- **Records kept** of AI use in advice.
- **Clients told** when AI helps prepare their advice.

WHERE TO START

Start with quote and wording comparison. It saves hours on every placement and directly reduces the risk of errors.

Advice, conduct and the client's interest

Brokers owe their clients a duty of care. AI must strengthen that duty, not weaken it. Regulators increasingly expect firms to show that advice is

suitable, conflicts are managed and customers are treated fairly, whether or not AI is involved.

EXHIBIT 6

Five conduct duties, and how AI supports them

DUTY	WHAT IT MEANS	HOW AI HELPS
Suitable advice	Cover that fits the client's needs and budget	Needs captured fully; gaps in cover flagged
Clear disclosure	Clients understand cover, exclusions and costs	Plain-language summaries of each option
Managing conflicts	Recommendations not driven by commission	Recommendations recorded with reasons
Fair treatment	No group of clients treated worse than others	Outcomes compared across client groups
Good records	Evidence of what was advised and why	Automatic, complete advice files

Source: KaribuTech AI analysis.

Three habits of good AI-assisted advice

01

Explain the why

Every recommendation says why it suits this client.

02

Show the options

Clients see what was compared, not only the winner.

03

Review the file

Senior brokers sample AI-assisted files monthly.

AI should make every client file better documented and every recommendation easier to explain. That is also the broker's best defence.

Working with insurers and partners

Brokers sit between clients and insurers. Better data exchange with insurers, underwriting agencies and

partners speeds up placement and servicing for everyone. The table shows where connections pay off.

EXHIBIT 7

Where better connections pay off

CONNECTION	WHAT IT ENABLES	WHAT TO AGREE
Insurer quoting links	Instant quotes for standard risks	Data formats; appetite rules
Underwriting agencies	Faster placement for specialist risks	Authority limits; reporting
Policy and claims updates	Live status of policies and claims	Data sharing and privacy
Bank and association partners	Access to new client groups	Referral terms; conduct duties
Commission and accounts	Automatic reconciliation of premiums and commission	Statement formats and timing

Source: KaribuTech AI analysis.

Three rules for partnerships

01

Standard data

Agree formats once, use them with every insurer.

02

Clear roles

Who advises, who underwrites, who services.

03

Client first

Partnerships must serve the client's interest.

The 100-day quick-wins plan for brokers

Do not start with a large programme. Pick one line of business or client group, deliver a working improvement in 100 days, measure it, then reuse the approach. Each

cycle is run by a small team of brokers, operations staff and technology partners.

EXHIBIT 8

One broking improvement in 100 days: design, build, go live, measure

DAYS 1-30

Discover and design

- Time today's placement and renewal steps
- Choose one line and client group
- Agree rules for AI use and records

DAYS 31-60

Build

- Submission capture and formatting
- Quote and wording comparison
- Renewal risk scores

DAYS 61-100

Go live and measure

- Use on live placements
- Compare time and errors with before
- Report to partners and staff

Source: KaribuTech AI quick-wins approach.

Three quick wins we recommend

QUICK WIN 1

Submission automation

- Client data captured once
- Formatted for each insurer
- Matched to insurer appetite
- Quotes chased automatically

MEASURE

Hours per placement, quotes received

QUICK WIN 2

Quote and wording comparison

- Prices and limits side by side
- Wording differences highlighted
- Plain-language client summary
- Comparison kept on file

MEASURE

Time to advise, errors found

QUICK WIN 3

Renewal risk alerts

- Every client scored 90 days out
- Renewal packs prepared
- At-risk clients contacted early
- Reasons for loss recorded

MEASURE

Client retention, renewals saved

The business case for AI in broking

The value of AI for a brokerage comes from keeping more clients, freeing time to win new ones and making

fewer costly errors. The example below is illustrative, for a mid-sized brokerage; your own baseline will differ.

EXHIBIT 9

Illustrative example: 5,000 clients, \$50m of premium placed, \$6m of commission and fee income

Measure	Before	After AI	Annual value
Client retention	85%	89%	≈ \$0.24m more income
Hours per placement	5 hours	2 hours	≈ 18,000 hours freed
New clients from freed time	–	360 more	≈ \$0.43m more income
Wording and document errors	Baseline	Halved	Fewer E&O claims
Total annual value			≈ \$0.7m + fewer errors

Source: Illustrative KaribuTech AI example, not a forecast. Retention: 4 points of \$6m income. Hours: 3 saved on 6,000 placements. New clients: 9 advisers × 40 clients × \$1,200 average income.

Where the value comes from

- **More clients kept:** renewals handled early, with advice.
- **More time to sell:** admin hours turned into client meetings.
- **Fewer errors:** wordings and documents checked every time.
- **Better records:** stronger conduct and audit position.

Where the money goes

Broking platform and portal	25–35%
Insurer connections	20–25%
Document and comparison AI	15–20%
CRM and renewal tools	10–15%
Training	10%

Typical split of a first-year budget (illustrative)

THREE RULES FOR A CREDIBLE CASE

- 1 Time today's placement and renewal work first
- 2 Count only income the principals sign off
- 3 Report retention, errors and advice time quarterly


What you need underneath: the broking platform

The broker principal does not need to be a technologist, but should know what the platform looks like. The key idea is one client record connected to every insurer and

every channel, so client data is captured once and reused everywhere.

EXHIBIT 10

One broking platform, connected to clients and markets



Source: KaribuTech AI analysis.

Build, buy or partner

Buy

Broking platforms, CRM, document and comparison tools

Build

Your advice process, client lanes and renewal rules

Partner

Insurers, underwriting agencies, associations and banks

Ask one question of every technology decision: will it give our brokers more time with clients?

The broking team in the AI era

AI takes over much of the admin in broking. The team's work moves to advice, negotiation, relationships and

claims support. The broker principal leads that shift.

EXHIBIT 11

How the main roles change

ROLE	TODAY	WITH AI
Account executive	Rekeys data and chases quotes	→ Advises clients with AI-prepared files
Placement broker	Sends submissions widely	→ Targets the right markets and negotiates
Broker assistant	Prepares documents by hand	→ Checks AI outputs and handles exceptions
Claims advocate	Handles claims when asked	→ Tracks every claim and pushes for fair outcomes
Compliance officer	Samples files after the event	→ Oversees complete, automatic advice records
Broker principal	Manages volume and commission	→ Leads an advice-first, AI-enabled brokerage

Source: KaribuTech AI analysis.

Who you work with

- **Insurers and agencies:** appetite, quoting and data exchange.
- **Technology partners:** platforms and integration.
- **Regulators and associations:** conduct and standards.
- **Banks and partners:** new client channels.
- **Clients:** feedback on advice and service.

Skills to build in the team

Advising with AI

Reading policy wordings

Digital client service

Negotiation

Conduct and compliance

Risks and guardrails for AI in broking

A broker's reputation rests on good advice. AI can make advice faster and better documented, but it can also

spread errors quickly. These are the risks a broker principal owns, and the controls that manage them.

EXHIBIT 12

Five broking risks, and the control for each

RISK	WHAT CAN GO WRONG	THE CONTROL
Unsuitable advice	AI suggests cover that does not fit the client	Broker reviews every recommendation; needs captured fully
Wording errors	An AI comparison misses an exclusion or limit	Human check before binding; comparison kept on file
Client data exposure	Client or claims data leaked or put in public AI tools	Approved tools; access controls; consent
Conflicts of interest	Recommendations lean towards higher commission	Recorded reasons; commission disclosed; reviews
Over-reliance	Staff stop questioning AI outputs	Training; sample checks; clear accountability

Source: KaribuTech AI analysis; IAIS, Application Paper on the supervision of AI (2025).


Eight questions to ask your team

- 01 How many hours a week go on advice versus admin?
- 02 Could we explain every recommendation we made last month?
- 03 Which clients are most likely to leave at renewal?
- 04 How often do our submissions go to the wrong insurers?
- 05 Where is our client data stored and processed?
- 06 Do we check every AI comparison before advising?
- 07 How fast are our clients' claims paid?
- 08 Which new client groups could we reach digitally?

The broking scorecard and 12-month roadmap

EXHIBIT 13


Six numbers to run a brokerage, plus four for the AI era



Source: KaribuTech AI analysis. Gold: keep within a target band; green: higher is better; red: lower is better.

EXHIBIT 14

Four quarters to an AI-enabled brokerage



Source: KaribuTech AI analysis.

WHERE TO START ON MONDAY

Ask each account executive to log one week of work. The hours spent rekeying and chasing are the hours AI can give back to clients.

About the authors



Kennedy Chengeta, PhD

Co-founder, KaribuTech AI

Enterprise Architect and AI Leader · Camunda Chapter Lead, Africa

Leads AI, automation and integration programmes for insurers and banks across Southern Africa.



Justice Muguda

Chief Technology Officer, KaribuTech AI


Cloud, integration and cybersecurity

Leads KaribuTech AI's technology and engineering, from cloud platforms to secure connections between systems.

With contributions from the KaribuTech AI Engineering Team. To discuss this report, or to run a 100-day broking diagnostic, contact KaribuTech AI at kaributech.co.za.

References

1. Accenture, underwriting survey, as reported by Insurance Business.
2. Capgemini, research on underwriters' time and workload.
3. GSMA (2026), State of the Industry Report on Mobile Money.
4. Stanford HAI (2025), Artificial Intelligence Index Report 2025.
5. Lloyd's of London, history of Lloyd's.
6. Confused.com, launch of UK price comparison (2002).
7. Safaricom (2007), launch of M-Pesa in Kenya.
8. IAIS (2025), Application Paper on the supervision of artificial intelligence, July 2025.
9. KaribuTech AI, BluFountain Insure platform (broker portal).
10. KaribuTech AI, The AI Blueprint for the Insurance CEO (2026).



About KaribuTech AI


KaribuTech AI is a South African technology company specialising in artificial intelligence, machine learning, intelligent automation and integration for insurers and banks across Africa. Our teams design and build intelligent workflows, decision engines, document and computer-vision AI, and cloud-native core platforms, including the **BluFountain Insure** platform: an insurance AI

marketplace with core systems for life, short-term, medical, funeral and investment business. We combine deep process expertise (Camunda, IBM BPM, MuleSoft, webMethods and document automation platforms) with practical AI engineering, so that AI moves out of the lab and into the operations that customers feel.

Disclaimer: this report provides general information and does not constitute legal, actuarial or investment advice. Third-party figures are cited beneath each exhibit and have not been independently verified by KaribuTech AI. Illustrative figures are not forecasts.



KaribuTech AI



INSURANCE · LEADERSHIP SERIES · SEPTEMBER 2026



The AI Blueprint for the Reinsurance Executive

Better data, sharper catastrophe models and faster recoveries: AI for reinsurers and reinsurance buyers

By Kennedy Chengeta, PhD, and Justice Muguda
with the KaribuTech AI Engineering Team

kaributech.co.za





About this report

This report is written for reinsurance executives, chief underwriting officers and reinsurance buyers at insurers in developing markets.

Reinsurance depends on data that arrives late and incomplete, catastrophe models under strain from a changing climate, and manual accounting that delays recoveries. This report shows how AI helps: the pressures, six moves, underwriting lanes, cedant data, catastrophe and climate, closing the protection gap, a 100-day plan, the business case, the team, the risks and a scorecard.

External data are sourced beneath each exhibit, and a full list of references appears at the end. Frameworks and illustrative figures without an external source are the authors' own.

WRITTEN FOR

Reinsurance executives

Underwriters

Cat modellers

Reinsurance buyers

Contents

01	Introduction	206
02	Why reinsurance must change	207
03	Lessons from history	208
04	Where the reinsurance cycle breaks	209
05	Six moves for the reinsurance executive	210
06	Three underwriting lanes	212
07	Data from cedants	213
08	Generative AI and AI agents in reinsurance	214
09	Catastrophe and climate	215
10	Closing the protection gap	216
11	The 100-day quick-wins plan	217
12	The business case	218
13	What you need underneath	219
14	The reinsurance team	220
15	Risks and guardrails	221
16	The scorecard and 12-month roadmap	222
17	About the authors and KaribuTech AI	223

Introduction

Reinsurance is the insurer's insurer. Its value rests on understanding risk better than anyone, and paying quickly when catastrophe strikes.

Yet much of reinsurance still runs on spreadsheets. Bordereaux arrive in dozens of formats, exposure data is incomplete, treaty accounts are reconciled by hand and recoveries take months. Meanwhile, a changing climate is testing catastrophe models built on the past.

AI can read bordereaux and treaty wordings, clean and map exposure data, run many more catastrophe scenarios, reconcile accounts and speed up recoveries, while underwriters and modellers keep the judgement.

This report sets out **six moves** for the reinsurance executive, **underwriting lanes**, **cedant data**, **catastrophe and climate**, **closing the protection gap**, a **100-day plan** and a **scorecard**, in plain language.

THREE QUESTIONS THIS REPORT ANSWERS

01

How do we turn messy cedant data into clean, usable data?

02

How do we price catastrophe risk in a changing climate?

03

How do we settle accounts and recoveries faster?

AT A GLANCE: SIX MOVES FOR REINSURANCE

- 1 Fix cedant data at the source
- 2 Price with better models and more scenarios
- 3 Automate treaty accounts and recoveries
- 4 See accumulations live
- 5 Design new covers for new risks
- 6 Govern models and counterparties

Why reinsurance must change

Five pressures are reshaping reinsurance. Each is also a chance for reinsurers and

buyers with better data to win.

Catastrophes are changing

Floods, storms, droughts and heat are shifting in frequency and severity, straining models built on history.

Cedant data is poor

Bordereaux and exposure data arrive late, in many formats and with gaps.

Accounts and recoveries are slow

Treaty accounts reconciled by hand delay settlements and tie up capital.

New risks lack history

Cyber, pandemics and AI liability have little loss data to price from.

Capital and reporting demands

IFRS 17 and solvency rules require detailed data on reinsurance held and ceded.

EXHIBIT 1

Why data and models now decide the price

1 Jan 2023

IFRS 17 in force, with reinsurance held measured separately

\$4.44m

global average cost of a data breach, a driver of cyber losses

280x

cheaper to run capable AI in 2024 than in 2022

2025

IAIS guidance on supervising AI in insurance

Source: IASB, IFRS 17; IBM, Cost of a Data Breach Report 2025; Stanford HAI, AI Index 2025; IAIS (2025).

In reinsurance, the cedant with the best data gets the best terms, and the reinsurer with the best models picks the best risks.

Lessons from history: reinsurance grows after the big losses

Reinsurance has grown through major losses and the lessons they taught. Each turning point improved how

risk is shared, modelled and paid for.

1370

The first reinsurance contract

The oldest known reinsurance contract, written in Genoa in 1370, passed part of a marine risk from one underwriter to another.

Today: AI reads treaty wordings and bordereaux in any format.

THE LESSON

Sharing risk lets insurers take on more than they could alone.

1846

The first professional reinsurer

Cologne Re was founded in Germany in 1846 after a great fire in Hamburg showed that local insurers could not carry such losses alone.

Today: cleaner data lets capital go where it is needed.

THE LESSON

Big losses create the need for dedicated reinsurance capital.

1992

Hurricane Andrew

Hurricane Andrew caused record insured losses in the United States and several insurers failed. The industry turned to computer catastrophe models, and new capital followed.

Today: climate-adjusted scenarios and satellite data refine the models.

THE LESSON

Models, not only experience, are needed for catastrophes.

Today

AI and new data

Satellite, weather and exposure data, combined with AI, allow far more scenarios and faster responses after events.

Today: loss estimates within days of an event.

THE LESSON

Better data and more scenarios make better prices and faster payouts.

Source: Histories of early marine reinsurance (Genoa, 1370), Cologne Re (1846) and Hurricane Andrew (1992).

Each big loss taught reinsurers to measure risk better. Climate change is the next test, and AI with better data is how the industry passes it.

Where the reinsurance cycle breaks, and how AI helps

Reinsurance follows a cycle, from cedant data to capital and retrocession. At each step, time, accuracy or

money can be lost. The exhibit shows where it typically breaks and what AI does.

EXHIBIT 2

The reinsurance cycle: eight steps from cedant data to retrocession

	WHERE IT BREAKS TODAY	WHAT AI DOES
Cedant data Submissions and exposure data	Many formats; missing fields	Data read, cleaned and mapped automatically
Underwriting Assessing the programme	Limited time per submission	Key risks and changes summarised
Pricing Setting the price	Few scenarios; old data	Many scenarios with current data
Treaty terms Wordings and conditions	Clauses compared by hand	Wordings compared line by line
Bordereaux and accounts Premium and claims statements	Manual reconciliation	Statements matched automatically
Claims and recoveries Paying and collecting claims	Slow notices and settlements	Claims tracked; recoveries chased
Accumulation Total exposure by area and peril	Seen once a year	Live maps of exposure
Capital and retro Protecting the reinsurer	Decisions on stale data	Capital and retro tested on current data

Source: KaribuTech AI analysis.

Most of the friction in reinsurance is data friction. Fix the data, and pricing, accounting and recoveries all speed up.

Six moves for the reinsurance executive

Reinsurance leaders who use AI well do six things differently. Each is about better data,

better models and faster settlement, not about choosing a technology.

MOVE 1

01

Fix cedant data at the source

Poor cedant data leads to wrong prices, slow accounts and disputes. AI can read bordereaux and exposure schedules in any format, map them to a standard and flag gaps, and cedants that provide clean data can be rewarded.

What leaders do

- Agree data standards with cedants and brokers.
- Use AI to read and map bordereaux and exposure schedules.
- Score cedant data quality and reflect it in terms.

ASK YOUR TEAM

How many different bordereaux formats do we process, and how much time does that cost?

MOVE 2

02

Price with better models and more scenarios

Catastrophe and pricing models built on the past miss how risks are changing. AI helps combine vendor models with satellite, weather and claims data, and run many more scenarios.

What leaders do

- Blend vendor catastrophe models with your own claims and exposure data.
- Run more scenarios, including climate-adjusted ones.
- Check prices against actual losses after every major event.

ASK YOUR TEAM

How far off were our models on the last major catastrophe in our markets?

MOVE 3

03

Automate treaty accounts and recoveries

Manual reconciliation of accounts and claims delays settlements for months. AI can match statements, spot differences and chase outstanding balances, speeding up cash for both sides.

What leaders do

- Match premium and claims statements automatically.
- Flag differences and route them to the right person.
- Track recoveries and outstanding balances by age.

ASK YOUR TEAM

How long does it take us to settle a treaty account, and why?

MOVE 4

04

See accumulations live

Concentrations of risk in one flood plain, city or industry can threaten the balance sheet. AI with geospatial data can map exposure continuously, not once a year.

What leaders do

- Map exposure by location and peril from cedant data.
- Update accumulation views as new data arrives.
- Alert underwriters when limits are approached.

ASK YOUR TEAM

Where are our biggest concentrations of risk today, not last year?

MOVE 5

05

Design new covers for new risks

Parametric cover, cyber reinsurance and protection-gap partnerships need new data and new pricing. AI helps build triggers, test scenarios and price risks with little history.

What leaders do

- Develop parametric covers with satellite and weather triggers.
- Build cyber and emerging-risk scenarios.
- Partner with governments and development agencies on protection-gap solutions.

ASK YOUR TEAM

Which new risks could we write if we had better data and models?

MOVE 6

06

Govern models and counterparties

Reinsurers rely on vendor models, AI tools and many counterparties. Governance must cover model risk, data quality and the financial strength of cedants and retrocessionaires.

What leaders do

- Validate catastrophe and AI models independently.
- Track counterparty strength and concentration.
- Keep an audit trail for every pricing decision.

ASK YOUR TEAM

Which model or counterparty could surprise us most next year?

HOW WE HELP

Documents, data and workflows for reinsurance

KaribuTech AI builds the document reading, data mapping and workflow tools that reinsurance depends on: reading bordereaux, schedules and treaty wordings in any format, mapping them to a standard, reconciling accounts and routing exceptions to the right person with deadlines. The same engines power BluFountain Insure's document centre and IFRS 17-compliant valuation.

Any format

bordereaux and schedules read automatically

Reconciled

accounts matched, exceptions routed

IFRS 17

compliant data and valuation

Three underwriting lanes: automate the routine, focus experts on the complex

Not every reinsurance programme needs the same effort. Sorting business into three lanes lets reinsurers

process routine programmes efficiently and spend expert time on complex and catastrophe-exposed risks.

EXHIBIT 3

Three lanes for reinsurance underwriting



MORE AUTOMATION →

← MORE EXPERT JUDGEMENT

Source: KaribuTech AI analysis.

What decides the lane

- **Structure:** proportional or non-proportional?
- **Data quality:** clean and complete, or patchy?
- **Catastrophe exposure:** low or high?
- **Novelty:** familiar risk or new peril?

Rules for every lane

- **Underwriters own the terms.**
- **Models are validated** before use.
- **Accumulation checked** before binding.
- **Decisions recorded** with reasons.

Data from cedants: the foundation of everything

Reinsurance can only be as good as the data cedants send. AI can turn messy bordereaux, schedules and

notices into clean, standard data, and show both sides where the gaps are.

EXHIBIT 4

Five data flows, and what AI does with each

DATA FLOW	WHAT AI DOES	BENEFIT
Exposure schedules	Reads locations, values and occupancy; geocodes addresses	Better catastrophe pricing
Premium bordereaux	Maps any format to a standard; checks totals	Faster, cleaner accounts
Claims bordereaux and notices	Extracts claims and matches to treaties	Faster recoveries
Treaty wordings	Compares clauses with past years and market wordings	Fewer disputes
Event reports	Links catastrophe events to exposure	Faster loss estimates

Source: KaribuTech AI analysis.

Reward good data

01

Publish standards

Tell cedants exactly what data you need.

02

Score quality

Measure each cedant's data every renewal.

03

Price it in

Better data earns better terms.


Generative AI and AI agents in reinsurance

Generative AI reads and writes; AI agents plan and carry out tasks. In reinsurance, that means reading treaty wordings and slips, extracting bordereaux,

summarising submissions and claims, and drafting renewal materials. Pricing and terms stay with underwriters.

EXHIBIT 5

Six jobs AI can take on in reinsurance today



Source: KaribuTech AI analysis. Labels show the level of AI autonomy we recommend.

What stays with people

- Pricing and binding decisions.
- Treaty terms and negotiations.
- Catastrophe model choices and adjustments.
- Large claims and disputes.

Guardrails that make it safe

- **Confidential data protected:** approved tools only.
- **Every extraction checked** against totals.
- **Model outputs validated** before pricing.
- **Records kept** for audit and IFRS 17.

WHERE TO START

Start with bordereaux extraction and wording comparison. Both save many hours and reduce disputes straight away.

Catastrophe and climate: modelling a changing world

Catastrophe models are built on history, but the climate is changing. AI helps by bringing in new data and generating more forward-looking scenarios. The table

shows where AI adds most for the perils that matter in our markets.

EXHIBIT 6

Where AI adds most, by peril

PERIL	CHALLENGE	WHAT AI ADDS	USE
Floods	Patchy loss history; changing rainfall	Satellite and elevation data mapped to exposure	Pricing; accumulation
Tropical storms	Shifting tracks and intensity	Climate-adjusted scenarios	Capital; retro
Drought	Slow onset; crop and livelihood losses	Rainfall and vegetation indices	Parametric triggers
Earthquake	Rare events; building data gaps	Building characteristics from imagery	Exposure quality
Heat and wildfire	New and growing risks	Temperature and fire-risk models	Emerging-risk pricing

Source: KaribuTech AI analysis.

After every event

01

Estimate fast

Use imagery and exposure to estimate losses within days.

02

Compare

Check model estimates against actual losses.

03

Update

Feed lessons into models and prices.

Closing the protection gap

Much of the economic loss from disasters in our markets is uninsured. Reinsurers can help close that gap by backing new products and partnerships, and AI

makes many of them workable. The African Risk Capacity, set up in 2012, is one example of pooled disaster cover for governments.

EXHIBIT 7

Five ways reinsurance can close the protection gap

APPROACH	HOW IT WORKS	WHAT AI ADDS
Parametric cover	Pays automatically when an index crosses a trigger	Better triggers from satellite and weather data
Sovereign and regional pools	Governments pool disaster risk and buy cover together	Faster loss estimates and payouts
Microinsurance backing	Reinsurance for low-premium, high-volume products	Data to price and monitor small risks
Catastrophe bonds	Investors take on catastrophe risk	Scenario modelling for investors
Public-private partnerships	Shared risk between governments and insurers	Shared data and exposure mapping

Source: African Risk Capacity (established 2012); KaribuTech AI analysis.

What makes it work

01

Trusted triggers

Independent, transparent data that customers understand.

02

Fast payouts

Money within days, by mobile money where possible.

03

Shared data

Exposure and loss data pooled across partners.

The protection gap is the largest growth opportunity in reinsurance. Data and AI are what make small, fast, fair payouts possible.

The 100-day quick-wins plan for reinsurance

Do not start with a large programme. Pick one data flow or one line of business, deliver a working improvement in 100 days, measure it, then reuse the approach. Each

cycle is run by a small team of underwriters, modellers, technical accountants and technology staff.

EXHIBIT 8

One reinsurance improvement in 100 days: design, build, run, measure

DAYS 1-30

Discover and design

- Map data flows and formats
- Agree data standards and checks
- Baseline times and errors

DAYS 31-60

Build

- Bordereaux reading and mapping
- Account matching rules
- Exposure mapping for one peril

DAYS 61-100

Run and measure

- Run on live treaties
- Compare with manual processing
- Report to the executive team

Source: KaribuTech AI quick-wins approach.

Three quick wins we recommend

QUICK WIN 1

Bordereaux automation

- Any format read and mapped
- Totals checked
- Gaps reported to cedants
- Standard data stored

MEASURE

Days to process, errors found

QUICK WIN 2

Accumulation dashboard

- Exposure mapped by location
- Updated as data arrives
- Limits and alerts
- Board view

MEASURE

Exposure visibility, limit breaches

QUICK WIN 3

Wording comparison

- Treaty wordings compared
- Changes highlighted
- Past disputes linked
- Notes for underwriters

MEASURE

Time to review, disputes

The business case for AI in reinsurance

The value of AI in reinsurance comes from better pricing, faster accounts and recoveries, and lower processing costs. The example below is illustrative, for

a reinsurer or reinsurance department writing \$300m of premium.

EXHIBIT 9

Illustrative example: \$300m of reinsurance premium written

Measure	Before	After AI	Annual value
Loss ratio from better pricing	62%	61%	≈ \$3.0m
Days to process bordereaux	30 days	5 days	Faster accounts
Time to settle treaty accounts	90 days	45 days	Cash in sooner
Disputes over wordings and data	Baseline	Halved	Lower cost and friction
Total measurable value			≈ \$3.0m + faster cash

Source: Illustrative KaribuTech AI example, not a forecast. Loss ratio: one point on \$300m of premium.

Where the value comes from

- **Better pricing:** cleaner data and more scenarios.
- **Faster cash:** accounts and recoveries settled sooner.
- **Lower costs:** less manual processing.
- **Fewer disputes:** wordings and data checked.

Where the money goes

Data ingestion and standards	30–40%
Pricing and catastrophe models	20–25%
Accounts automation	15–20%
Geospatial data	10–15%
Training and governance	10%

Typical split of a first-year budget (illustrative)

THREE RULES FOR A CREDIBLE CASE

- 1 Baseline processing times and pricing accuracy first
- 2 Validate models before relying on them
- 3 Report value and model accuracy every quarter


What you need underneath: the reinsurance engine

The reinsurance executive does not need to be a technologist, but should know what the engine looks like. The key idea is one standard data layer fed by

every cedant and broker, so pricing, accounts, accumulation and reporting all use the same clean data.

EXHIBIT 10

One reinsurance engine, fed by the whole market



Source: KaribuTech AI analysis.

Build, buy or partner

Buy

Catastrophe models, geospatial data, document tools

Build

Your data standards, pricing adjustments and rules

Partner

Brokers, cedants, modelling firms and data providers

Ask one question of every technology decision: will the next bordereau, event or renewal be faster and more accurate because of it?

The reinsurance team in the AI era

AI takes over data preparation, reconciliation and first drafts. The team's work moves to judgement: pricing, terms, catastrophe modelling and relationships with

cedants and brokers. The reinsurance executive leads that shift.

EXHIBIT 11

How the main roles change

ROLE	TODAY	WITH AI
Treaty underwriter	Rekeys and reviews submissions	→ Prices and negotiates with AI-prepared files
Catastrophe modeller	Runs a few model scenarios	→ Runs and validates many climate-adjusted scenarios
Technical accountant	Reconciles accounts by hand	→ Reviews exceptions from automatic matching
Claims manager	Chases notices and balances	→ Focuses on large and catastrophe claims
Data analyst	Cleans cedant data manually	→ Sets standards and scores data quality
Reinsurance executive	Oversees capacity and results	→ Leads a data-driven, AI-enabled reinsurance business

Source: KaribuTech AI analysis.

Who you work with

- **Cedants and brokers:** data standards and renewals.
- **Actuarial:** pricing, reserving and IFRS 17.
- **Risk and capital:** accumulation and solvency.
- **Retrocessionaires:** protection for the reinsurer.
- **Technology partners:** data and models.

Skills to build in the team

Catastrophe modelling

Data standards

Reading AI outputs critically

Climate risk

Negotiation

Risks and guardrails in reinsurance

Reinsurance concentrates risk, so model and data errors can be very costly. AI can improve pricing and

processing, but it must be governed carefully. These are the risks a reinsurance executive owns.

EXHIBIT 12

Five reinsurance risks, and the control for each

RISK	WHAT CAN GO WRONG	THE CONTROL
Model error	Catastrophe or pricing models miss how risk is changing	Independent validation; compare with actual losses
Poor data	Wrong exposure data leads to wrong prices and limits	Data standards; quality scores; checks on totals
Hidden accumulation	Concentrations missed until an event	Live exposure mapping; limit alerts
Counterparty failure	A cedant or retrocessionaire cannot pay	Credit monitoring; limits; collateral
Cyber and confidentiality	Sensitive treaty and cedant data exposed	Access controls; approved tools; encryption

Source: KaribuTech AI analysis; IAIS, Application Paper on the supervision of AI (2025).

Eight questions to ask your team

01

How far off were our models on the last major event?

02

Which cedants send the poorest data?

03

Where are our largest accumulations today?

04

How long does it take to settle a treaty account?

05

Which counterparty worries us most?

06

Have our AI and catastrophe models been validated?

07

Which new risks could we write with better data?


08

How do we support closing the protection gap?

The reinsurance scorecard and 12-month roadmap

EXHIBIT 13


Six numbers to run reinsurance, plus four for the AI era



Source: KaribuTech AI analysis. Gold: keep within a target band; green: higher is better; red: lower is better.

EXHIBIT 14

Four quarters to AI-enabled reinsurance



Source: KaribuTech AI analysis.

WHERE TO START ON MONDAY

Count how many bordereaux formats you process and how many days they take. That is where the first 100 days will pay back fastest.

About the authors



Kennedy Chengeta, PhD

Co-founder, KaribuTech AI

Enterprise Architect and AI Leader · Camunda Chapter Lead, Africa

Leads AI, automation and integration programmes for insurers and banks across Southern Africa.



Justice Muguda

Chief Technology Officer, KaribuTech AI


Cloud, integration and cybersecurity

Leads KaribuTech AI's technology and engineering, from cloud platforms to secure connections between systems.

With contributions from the KaribuTech AI Engineering Team. To discuss this report, or to run a 100-day reinsurance diagnostic, contact KaribuTech AI at kaributech.co.za.

References

1. IASB (2017, amended 2020), IFRS 17 Insurance Contracts.
2. IBM (2025), Cost of a Data Breach Report 2025.
3. Stanford HAI (2025), Artificial Intelligence Index Report 2025.
4. IAIS (2025), Application Paper on the supervision of artificial intelligence, July 2025.
5. Histories of early marine reinsurance (Genoa, 1370).
6. Histories of Cologne Re (founded 1846).
7. Histories of Hurricane Andrew (1992) and catastrophe modelling.
8. African Risk Capacity (established 2012).
9. KaribuTech AI, BluFountain Insure platform (document centre and valuation).
10. KaribuTech AI, The AI Blueprint for the Insurance CEO (2026).



About KaribuTech AI

KaribuTech AI is a South African technology company specialising in artificial intelligence, machine learning, intelligent automation and integration for insurers and banks across Africa. Our teams design and build intelligent workflows, decision engines, document and computer-vision AI, and cloud-native core platforms, including the **BluFountain Insure** platform: an insurance AI

marketplace with core systems for life, short-term, medical, funeral and investment business. We combine deep process expertise (Camunda, IBM BPM, MuleSoft, webMethods and document automation platforms) with practical AI engineering, so that AI moves out of the lab and into the operations that customers feel.

Disclaimer: this report provides general information and does not constitute legal, actuarial or investment advice. Third-party figures are cited beneath each exhibit and have not been independently verified by KaribuTech AI. Illustrative figures are not forecasts.



KaribuTech AI

INSURANCE · LEADERSHIP SERIES · SEPTEMBER 2026


The AI Blueprint for the Insurance Regulator

Supervising AI in insurance, and using AI to supervise: a practical guide for regulators in developing markets

By Kennedy Chengeta, PhD, and Justice Muguda
with the KaribuTech AI Engineering Team

kaributech.co.za





Contents

01	Introduction	226
02	Why supervision must change	227
03	Lessons from history	228
04	Where supervision breaks	229
05	Six moves for the insurance regulator	230
06	Proportionate supervision of AI	232
07	SupTech: AI for supervisors	233
08	Generative AI and AI agents for supervisors	234
09	The AI rulebook for insurers	235
10	Regulating for inclusion	236
11	The 100-day quick-wins plan	237
12	The public value case	238
13	What you need underneath	239
14	The supervisory team	240
15	Risks of AI in supervision	241
16	The scorecard and 12-month roadmap	242
17	About the authors and KaribuTech AI	243



About this report

This report is written for insurance commissioners, supervisors and policy teams in developing markets, and for the insurers and partners who work with them.

It covers both sides of AI for regulators: how to supervise insurers' use of AI fairly and proportionately, and how to use AI to supervise better. It sets out the pressures on supervisors, six moves, proportionate rules, SupTech, the emerging rulebook, regulating for inclusion, a 100-day plan, the public value case, the team, the risks and a scorecard.

External data are sourced beneath each exhibit, and a full list of references appears at the end. Frameworks and illustrative figures without an external source are the authors' own.

WRITTEN FOR

Commissioners

Supervisors

Policy teams

Industry bodies

Introduction

AI is changing insurance faster than rules are being written. Regulators now have two jobs: making sure insurers use AI fairly and safely, and using AI themselves to supervise a faster, more digital market.

Many supervisors still receive returns on spreadsheets or paper, review them by hand weeks after the period ends, and learn about problems from complaints. Meanwhile, insurers are using AI to price, underwrite and settle claims, and new players are selling cover through phones and platforms.

Done well, AI helps both sides. Insurers can reach the uninsured at low cost, and supervisors can move from paper returns to live data, spot problems early and focus scarce staff where risk is highest.

This report sets out **six moves** for the insurance regulator, **proportionate supervision of AI**, **SupTech** use cases, the emerging **AI rulebook**, regulating for **inclusion**, a **100-day plan** and a **scorecard**, in plain language.

THREE QUESTIONS THIS REPORT ANSWERS

01 How should we supervise insurers' AI, proportionately?

02 How can AI help us supervise better and sooner?

03 How can rules enable inclusion while protecting customers?

AT A GLANCE: SIX MOVES FOR THE REGULATOR

1 Learn by doing: labs and sandboxes

2 Move from paper returns to live data

3 Use AI to supervise (SupTech)

4 Set proportionate AI rules

5 Build new skills and structures

6 Partner across borders and sectors

Why insurance supervision must change

Five pressures are reshaping the supervisor's job. Each is also a chance to protect

customers better and grow a healthier market.

Insurers are using AI already

Pricing, underwriting, claims and fraud decisions are increasingly made or shaped by models that supervisors rarely see.

Markets are going digital

Cover is sold through mobile money, telcos and platforms, often by partners the regulator does not license directly.

Data arrives too late

Paper and spreadsheet returns, reviewed by hand, show problems months after they start.

Skills are scarce

Few supervisors have data science or AI skills, and competing with industry salaries is hard.

Global standards are moving

The IAIS, the EU and national authorities are setting AI expectations that local markets will be measured against.

EXHIBIT 1

The rulebook and the tools are arriving together

Jul 2025

IAIS Application Paper on the supervision of AI

Aug 2024

EU AI Act in force; life and health pricing high-risk

Dec 2023

NAIC model bulletin on insurers' use of AI

280x

cheaper to run capable AI in 2024 than in 2022

Source: IAIS (2025); EU AI Act (2024); NAIC (2023); Stanford HAI, AI Index 2025.

A regulator that understands AI can say yes to innovation with confidence. One that does not will either say no too often, or yes too late.

Lessons from history: regulation follows crises, and good supervisors get ahead of them

Insurance regulation has usually been written after something went wrong. The best supervisors learned to

see the next problem coming. AI offers the chance to do that at scale.

1720

The Bubble Act

After the South Sea Bubble, the Bubble Act restricted joint-stock companies and chartered two marine insurers, the Royal Exchange and London Assurance. Rules arrived after the crash.

Today: shape AI rules early, before a crisis forces blunt ones.

THE LESSON

Rules written after a crisis are often blunt. Better to shape them early.

1905

The Armstrong Investigation

A New York legislative inquiry into life insurers exposed misuse of policyholders' money and led to sweeping reforms of US life insurance.

Today: AI inventories and complaints data make conduct visible.

THE LESSON

Transparency and data about insurers' conduct protect policyholders.

2016

Solvency II

The EU's risk-based regime for insurers came into force in 2016, linking capital to risk and requiring regular, detailed reporting.

Today: digital returns and SupTech make risk-based supervision live.

THE LESSON

Risk-based supervision needs good, timely data.

Today

AI and SupTech

Supervisors can now use data and AI to monitor markets continuously, and must also supervise the AI that insurers use.

Today: supervise insurers' AI, and use AI to supervise.

THE LESSON

Regulators can finally get ahead of the problem, if they invest in data and skills.

Source: Histories of the South Sea Bubble and the Bubble Act (1720); New York State Armstrong Committee (1905); EU Solvency II Directive (in force 2016).

Each era's regulators were limited by the information they had. AI changes that limit, for regulators willing to build the data and skills to use it.

Where supervision breaks, and how AI helps

Supervision follows a cycle, from licensing to enforcement and market development. At each step, information arrives late or is too much for staff to read.

The exhibit shows where the cycle typically breaks and what AI does.

EXHIBIT 2

The supervisory cycle: eight steps from licensing to market development

	WHERE IT BREAKS TODAY	WHAT AI DOES
Licensing New insurers and products	Long paper applications	Applications checked and summarised automatically
Returns Regular data from insurers	Spreadsheets reviewed by hand	Digital returns validated on arrival
Off-site analysis Solvency, conduct and trends	Few ratios, reviewed late	Early-warning indicators across the market
On-site inspection Visits to insurers	Inspections planned by rotation	Inspections targeted by risk
Conduct and complaints Fair treatment of customers	Complaints read one by one	Themes found across all complaints
Enforcement Action when rules are broken	Slow, case-by-case work	Evidence gathered and organised faster
Policy making New rules and guidance	Rules drafted without market data	Rules tested against real data
Market development Growing inclusion and innovation	Innovation seen as risk only	Sandboxes and data on new models

Source: KaribuTech AI analysis.

The goal is not more data for its own sake. It is seeing the problem early enough to prevent the harm.

Six moves for the insurance regulator

Regulators who handle AI well do six things differently. Each is about better information,

proportionate rules and the right skills, not about buying a technology.

MOVE 1

Learn by doing: labs and sandboxes

Supervisors cannot regulate what they do not understand. An innovation office, an AI lab and a regulatory sandbox let supervisors see new models working, with real customers and clear limits.

What leaders do

- Open an innovation office as a single door for new insurers and partners.
- Run a sandbox for AI, parametric and mobile products, with clear exit criteria.
- Let supervisors test AI tools on their own work, safely.

ASK YOUR TEAM

What have we learned from the last three innovations we approved, or refused?

01

MOVE 2

Move from paper returns to live data

Paper and spreadsheet returns are slow and error-prone. Common data standards and secure digital submission let supervisors receive cleaner data sooner and spend time on analysis, not re-typing.

What leaders do

- Publish data standards for returns, claims and complaints.
- Collect returns through a secure portal or direct system connections.
- Validate data automatically the moment it arrives.

ASK YOUR TEAM

How many days after each quarter do we first see an insurer's numbers?

02

MOVE 3

Use AI to supervise (SupTech)

With good data, AI can flag insurers whose ratios are moving the wrong way, read thousands of complaints for common themes and help plan inspections by risk.

What leaders do

- Build early-warning indicators across solvency, liquidity and conduct.
- Use AI to find themes in complaints and social media.
- Plan inspections by risk, not by rotation.

ASK YOUR TEAM

Which insurer is most likely to fail or mistreat customers next year, and how do we know?

03

MOVE 4

Set proportionate AI rules

Rules should match the risk. A chatbot answering policy questions is not the same as a model that declines claims. Proportionate, principles-based rules protect customers without blocking useful innovation.

What leaders do

- Follow the IAIS principles: governance, robustness, transparency, fairness and redress.
- Tier requirements by the impact of the AI on customers.
- Ask insurers for an AI inventory and a named accountable executive.

ASK YOUR TEAM

Do our rules treat a drafting assistant and a claims-decline model the same way?

04

MOVE 5

05

Build new skills and structures

Supervising AI and using SupTech need data engineers, data scientists and supervisors who understand models. Most regulators need a data unit and a senior data leader.

What leaders do

- Appoint a chief data officer and create a small data and AI unit.
- Train every supervisor in AI basics and model review.
- Second staff to and from industry and other regulators.

ASK YOUR TEAM

How many of our supervisors could review an insurer's pricing model today?

MOVE 6

06

Partner across borders and sectors

AI in insurance crosses boundaries: telcos, banks, platforms and cloud providers, often in several countries. No single regulator sees the whole picture.

What leaders do

- Agree information sharing with central banks, telecom and data-protection regulators.
- Work with regional peers on common standards and shared tools.
- Engage industry early on new rules through consultation and pilots.

ASK YOUR TEAM

Which other authority sees risks in our market that we do not?

AROUND THE WORLD

Already happening around the world

Supervisors are already learning by doing. The UK's Financial Conduct Authority opened an AI Lab in 2024 and launched AI Live Testing in 2025. Singapore's MAS set up its FinTech and Innovation Group in 2015. South Africa's Intergovernmental Fintech Working Group opened an Innovation Hub in 2020, and India's IRDAI launched an insurance regulatory sandbox in 2019.

FCA

AI Lab (2024) and AI Live Testing (2025)

MAS

FinTech and Innovation Group (2015)

IFWG and IRDAI

Innovation Hub (2020); sandbox (2019)


Proportionate supervision of insurers' AI

Supervisors cannot review every model. A proportionate approach sets light expectations for low-risk AI and focuses attention on AI that affects

customers' price, cover or claims. This mirrors the tiered approach in the EU AI Act and the IAIS principles.

EXHIBIT 3

Three levels of supervisory attention, matched to the risk of the AI



LIGHTER TOUCH →

← CLOSER SUPERVISION

Source: KaribuTech AI analysis.

What every insurer should have

- An AI inventory listing every model and its level.
- An accountable executive for AI.
- A governance policy approved by the board.
- Complaint routes for AI decisions.

What supervisors should check

- **Governance:** who owns and approves AI.
- **Fairness:** outcomes across customer groups.
- **Explanations:** can decisions be explained?
- **Third parties:** vendor models and data.

SupTech: how AI helps supervisors

SupTech is technology used by supervisors to do their own work better. The most useful early uses are simple: checking returns, spotting trends and reading large

volumes of text. The table shows where AI helps across the supervisor's work.

EXHIBIT 4

Seven SupTech uses, from quick wins to advanced

USE	WHAT AI DOES	BENEFIT	DIFFICULTY
Returns validation	Checks every submitted return for errors and unusual changes	Cleaner data, days not weeks	Low
Early-warning indicators	Tracks solvency, liquidity and claims ratios across all insurers	Problems seen sooner	Medium
Complaints analytics	Finds themes across complaints, calls and social media	Conduct issues found early	Low
Pricing fairness checks	Tests premiums and declines for unfair patterns	Fairer outcomes for customers	High
Inspection planning	Ranks insurers by risk to target on-site visits	Scarce staff used well	Medium
Money-laundering screening	Finds unusual payments and ownership patterns	Financial crime risk reduced	Medium
Climate exposure	Maps insurers' exposure to floods, droughts and storms	Systemic risk understood	High

Source: KaribuTech AI analysis.

Three rules for SupTech

01

Start with data quality

No analytics will fix late or wrong returns.

02

Keep supervisors in charge

AI flags; supervisors judge and act.

03

Share what works

Reuse tools with regional peers.

Generative AI and AI agents for supervisors

Generative AI reads and writes; AI agents plan and carry out tasks. For supervisors, whose work is full of documents, that means help reading filings, comparing

policies, drafting letters and answering routine questions from industry and the public.

EXHIBIT 5

Six jobs AI can take on in the supervisory team today

Read filings Summarises licence applications, product filings and board papers for supervisors ASSIST	Compare policies Compares policy wordings and flags unfair or unclear terms RECOMMEND	Draft letters Drafts supervisory letters and findings for supervisors to review RECOMMEND
Answer routine queries Answers common questions from insurers and the public from published guidance ACT WITHIN LIMITS	Track global rules Summarises new standards and rules from other regulators ASSIST	Prepare inspections Pulls together an insurer's data, history and issues before a visit ASSIST

Source: KaribuTech AI analysis. Labels show the level of AI autonomy we recommend.

What stays with supervisors

- Every supervisory judgement and decision.
- Enforcement and sanctions.
- Licensing approvals and refusals.
- Public statements and policy.

Guardrails that make it safe

- **Secure, approved tools only:** no confidential data in public AI.
- **Every draft reviewed** before it is sent.
- **Clear records** of how AI was used.
- **Tested** for accuracy on real cases first.

WHERE TO START

Start with summarising filings and tracking global rules. They save time, carry little risk and build confidence with the tools.

The AI rulebook for insurers is taking shape

Rules for AI in insurance are converging around the same principles. Local regulators do not need to copy any one regime, but aligning with international

standards makes it easier for insurers and investors to operate across borders.

EXHIBIT 6

Key milestones for AI regulation in insurance

WHEN	WHAT	WHY IT MATTERS FOR REGULATORS
Dec 2023	NAIC model bulletin on insurers' use of AI (US)	A principles-based model adopted by many US states
Jul 2024	African Union Continental AI Strategy	A shared direction for national AI policies across Africa
Aug 2024	EU AI Act enters into force	A risk-tiered law; life and health pricing is high-risk
Mar 2025	Kenya National AI Strategy 2025–2030	An example of national AI planning in our markets
Jul 2025	IAIS Application Paper on the supervision of AI	The global insurance supervisors' guidance, built on the Insurance Core Principles
Dec 2027	EU high-risk AI rules apply	Deferred from August 2026 by the Digital Omnibus

Source: NAIC (2023); African Union (2024); EU AI Act and Digital Omnibus (2024, 2026); Government of Kenya (2025); IAIS (2025).

The IAIS principles, in plain language

01 Govern Clear ownership, board oversight and a record of AI in use.	02 Explain and be fair Decisions that can be explained and do not treat groups unfairly.	03 Be robust and offer redress Secure, reliable models and a way for customers to complain.
---	--	---

Build once to the common principles, and map each new national rule to them.

Regulating for inclusion: letting AI reach the uninsured

In our markets, most people have no insurance. AI, mobile money and new partners can change that, but only if rules allow simple products and new channels

while protecting customers. Regulators shape how fast inclusion happens.

EXHIBIT 7

Five areas where regulators enable inclusion, and what to protect

AREA	WHAT REGULATORS CAN ENABLE	WHAT TO PROTECT
Microinsurance	Simple products with lighter capital and reporting rules	Clear terms; claims actually paid
Mobile and embedded sales	Cover sold by telcos, banks and platforms	Customer understanding; fair pricing
Digital onboarding	Remote identity checks and e-signatures	Data protection; fraud controls
Parametric cover	Automatic payouts on weather or index triggers	Basis risk explained to customers
Data sharing	Shared data for pricing and fraud, with consent	Privacy; competition; security

Source: KaribuTech AI analysis.

Three tools that help

01

Sandboxes

Test new products with real customers under limits.

02

Tiered licences

Lighter rules for small, simple products.

03

Market data

Publish penetration and claims-paid data by product.

The 100-day quick-wins plan for regulators

Do not start with a large programme. Pick one supervisory process, deliver a working improvement in

100 days, measure it, then build on it. Each cycle is run by a small team from supervision, policy, IT and data.

EXHIBIT 8

One supervisory improvement in 100 days: design, build, run, measure

DAYS 1-30

Discover and design

- Map the process and its data
- Agree data standards and checks
- Agree security and privacy rules

DAYS 31-60

Build

- Build the portal, checks or analysis
- Test on past data
- Train the supervisors who will use it

DAYS 61-100

Run and measure

- Run beside today's process
- Compare speed and findings
- Report to the commissioner and board

Source: KaribuTech AI quick-wins approach.

Three quick wins we recommend

QUICK WIN 1

Industry AI survey

- Ask insurers for AI inventories
- Tier uses by customer impact
- Identify high-risk uses
- Publish anonymised findings

MEASURE

Response rate, high-risk uses found

QUICK WIN 2

Automated returns checks

- Digital submission of one return
- Automatic validation
- Unusual changes flagged
- Queries sent to insurers

MEASURE

Days to first review, errors found

QUICK WIN 3

Complaints analytics

- All complaints in one place
- AI finds common themes
- Insurers ranked by conduct risk
- Findings used to plan visits

MEASURE

Themes found, time to act

The public value case for AI in supervision

For a regulator, the return on AI is measured in public value: problems found earlier, customers protected, supervisors' time used better and a market that grows

safely. The example below is illustrative, for a supervisor of 40 insurers.

EXHIBIT 9

Illustrative example: a supervisor of 40 insurers with 30 supervisory staff

Measure	Before	After AI	Public value
Days to first review of returns	45 days	5 days	Problems seen sooner
Staff time on data checking	40% of time	15% of time	≈ 7 staff freed for analysis
Complaints read for themes	A sample	All of them	Conduct issues found early
Inspections targeted by risk	By rotation	By risk ranking	Visits where harm is likeliest
Overall			Earlier action, safer growth

Source: Illustrative KaribuTech AI example, not a forecast. Staff time: 25% of a 30-person team moved from data checking to analysis and inspection.

Where the value comes from

- **Earlier action:** returns reviewed within days.
- **Better targeting:** inspections go where risk is highest.
- **Customer protection:** complaint themes spotted early.
- **Market growth:** faster, safer approval of new products.

Where the money goes

Data standards and portal	30–40%
Analytics and AI tools	20–25%
Skills and hiring	20–25%
Security and privacy	10–15%
Engagement with industry	5–10%

Typical split of a first-year budget (illustrative)

THREE RULES FOR A CREDIBLE CASE

- 1 Measure today's review times and findings first
- 2 Run new tools beside old processes before switching
- 3 Report outcomes, not activity, to the board


What you need underneath: the supervisory data platform

The commissioner does not need to be a technologist, but should sponsor the platform supervisors work on. The key idea is one secure supervisory data platform,

fed by insurers and other sources, so every supervisor works from the same, timely information.

EXHIBIT 10

One supervisory data platform, fed by the whole market



Source: KaribuTech AI analysis.

Build, buy or partner

Buy

Secure portals, analytics platforms, text-analysis tools

Build

Your data standards, indicators, risk models and reports

Partner

Regional regulators, development partners, universities

Ask one question of every technology decision: will we see the next problem sooner because of it?

The supervisory team in the AI era

AI changes supervision from checking data to analysing it, and adds a new job: supervising insurers' AI.

Regulators need a mix of experienced supervisors and new data skills, working as one team.

EXHIBIT 11

How the main roles change

ROLE	TODAY	WITH AI
Off-site analyst	Re-keys and checks returns	→ Analyses early warnings and trends
On-site examiner	Visits insurers by rotation	→ Targets visits by risk; reviews AI governance
Actuarial supervisor	Reviews reserves and capital	→ Also reviews pricing and AI models
Conduct officer	Handles complaints case by case	→ Finds themes and acts on patterns
Policy officer	Drafts rules from principles	→ Tests rules against market data
Commissioner	Oversees licensing and enforcement	→ Leads a data-driven, AI-aware authority

Source: KaribuTech AI analysis.

Who you work with

- Central bank: banking partners and payments.
- Telecom regulator: mobile channels.
- Data protection authority: privacy and consent.
- Competition authority: data sharing and platforms.
- Regional peers and the IAIS: standards and tools.

Skills to build in the team

Data analysis

Model review basics

AI governance

Plain-language communication

Cross-agency working

Risks of using AI in supervision

SupTech brings its own risks. A supervisor that relies on a faulty indicator, or leaks confidential data, can do real harm. These are the pitfalls to manage.

EXHIBIT 12

Five SupTech risks, and the control for each

RISK	WHAT CAN GO WRONG	THE CONTROL
False signals	Indicators that miss real problems or flag healthy insurers	Test indicators against past failures; supervisors judge
Data security	Confidential insurer and customer data leaked or attacked	Strong access controls; encryption; audits
Over-reliance	Supervisors stop questioning the dashboard	Keep on-site work and challenge; review what models miss
Unfair treatment	Some insurers or customer groups treated inconsistently	Document methods; review outcomes; allow appeal
Vendor lock-in	Dependence on one supplier for critical tools	Open standards; exit plans; own the data

Source: KaribuTech AI analysis; IAIS, Application Paper on the supervision of AI (2025).


Eight questions to ask your team

- 01 How many days after each quarter do we see the numbers?
- 02 Which insurers use AI to decide claims or prices?
- 03 Could we explain how our early-warning scores work?
- 04 Where is our confidential data stored, and who can see it?
- 05 How many supervisors can review a pricing model?
- 06 Which complaint themes have we missed this year?
- 07 Do our rules help or block inclusive products?
- 08 Which other authorities do we share data with?

The supervisory scorecard and 12-month roadmap

EXHIBIT 13


Six numbers to run a modern supervisor, plus four for the AI era



Source: KaribuTech AI analysis. Gold: keep within a target band; green: higher is better; red: lower is better.

EXHIBIT 14

Four quarters to a data-driven, AI-ready supervisor



Source: KaribuTech AI analysis.

WHERE TO START ON MONDAY

Send every licensed insurer a short survey asking which AI tools they use and for what. The answers will show where supervision should focus first.

About the authors



Kennedy Chengeta, PhD

Co-founder, KaribuTech AI

Enterprise Architect and AI Leader · Camunda Chapter Lead, Africa

Leads AI, automation and integration programmes for insurers and banks across Southern Africa.



Justice Muguda

Chief Technology Officer, KaribuTech AI


Cloud, integration and cybersecurity

Leads KaribuTech AI's technology and engineering, from cloud platforms to secure connections between systems.

With contributions from the KaribuTech AI Engineering Team. To discuss this report, or to run a 100-day supervisory AI diagnostic, contact KaribuTech AI at kaributech.co.za.

References

1. IAIS (2025), Application Paper on the supervision of artificial intelligence, July 2025.
2. European Union (2024), Regulation (EU) 2024/1689, the Artificial Intelligence Act; Digital Omnibus (2026).
3. NAIC (2023), Model Bulletin: Use of Artificial Intelligence Systems by Insurers.
4. African Union (2024), Continental Artificial Intelligence Strategy.
5. Government of Kenya (2025), Kenya Artificial Intelligence Strategy 2025–2030.
6. Financial Conduct Authority, AI Lab (2024) and AI Live Testing (2025).
7. Monetary Authority of Singapore, FinTech and Innovation Group (2015).
8. Intergovernmental Fintech Working Group, Innovation Hub (2020); IRDAI regulatory sandbox (2019).
9. Histories of the Bubble Act (1720) and the Armstrong Investigation (1905).
10. European Union, Solvency II Directive (in force 2016).
11. Stanford HAI (2025), Artificial Intelligence Index Report 2025.
12. KaribuTech AI, The AI Blueprint for the Insurance CEO (2026).



About KaribuTech AI

KaribuTech AI is a South African technology company specialising in artificial intelligence, machine learning, intelligent automation and integration for insurers and banks across Africa. Our teams design and build intelligent workflows, decision engines, document and computer-vision AI, and cloud-native core platforms, including the **BluFountain Insure** platform: an insurance AI

marketplace with core systems for life, short-term, medical, funeral and investment business. We combine deep process expertise (Camunda, IBM BPM, MuleSoft, webMethods and document automation platforms) with practical AI engineering, so that AI moves out of the lab and into the operations that customers feel.

Disclaimer: this report provides general information and does not constitute legal, actuarial or investment advice. Third-party figures are cited beneath each exhibit and have not been independently verified by KaribuTech AI. Illustrative figures are not forecasts.